

# Boosting Mobile CNN Inference through Semantic Memory

**Yun Li<sup>1,4</sup>, Chen Zhang<sup>2,\*</sup>, Shihao Han<sup>3,4</sup>, Li Lyна Zhang<sup>4</sup>, Baoqun Yin<sup>1,\*</sup>,**  
**Yunxin Liu<sup>5</sup>, Mengwei Xu<sup>6</sup>**

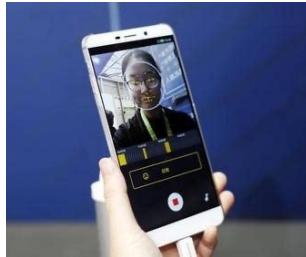
<sup>1</sup>University of Science and Technology of China, <sup>2</sup>Alibaba Group, <sup>3</sup>Rose-Hulman Institute of Technology,

<sup>4</sup>Microsoft Research Asia, <sup>5</sup>Tsinghua University, <sup>6</sup>Beijing University of Posts and Telecommunications

(\* indicates the corresponding author)



# CNNs have catalyzed many emerging mobile vision tasks



Face Recognition



Classification



Action Recognition

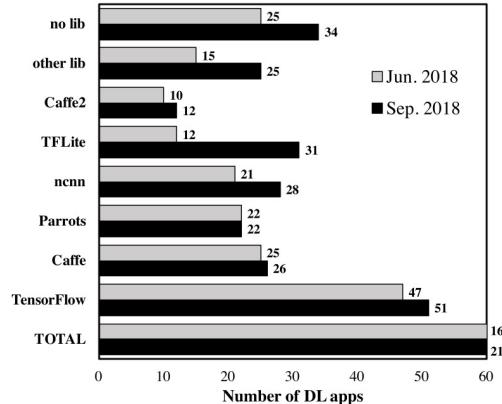
- DNNs have achieved great success in many **continuous mobile vision** applications.
- The mobile/wearable devices need to **perform CNN inference in real time** on these video images.

# Fast inference on mobile devices is urgent

- The CNN executions are costly.
  - high time complexity and energy-consuming
- Offloading to the cloud?
  - tight delay constraint and data privacy concerns
- A notable trend is on-device CNN inference



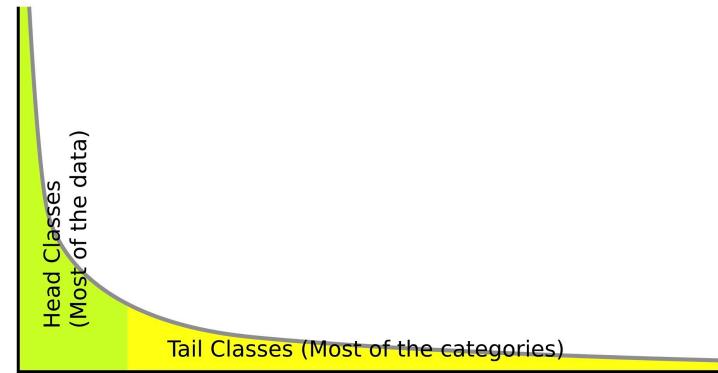
Cloud? Privacy concerns



increased by 27% with 3 month

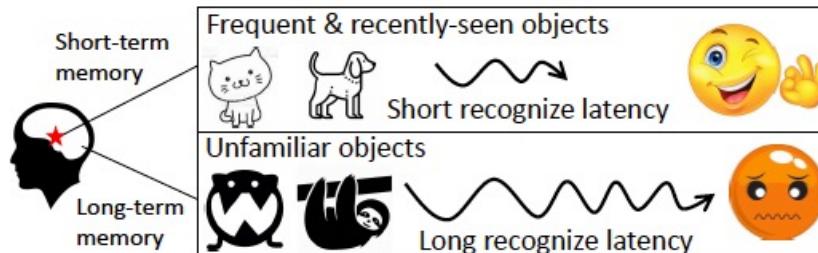
# Two critical observations

- The temporal locality in mobile video streams
  - Recently seen objects are more likely to appear again in the next few frames
- Long-tail distribution
  - The frequency of object occurrence in the mobile video streams typically follows a long-tail distribution



# How does the human brain solve it?

- Human brain leverages temporal redundancy with priming effect
- Priming effect : a psychology phenomenon whereby exposure to one stimulus improves a response to a subsequent stimulus, without conscious guidance or intention.
- Priming effect is related to the long- and short-term memory of human brains

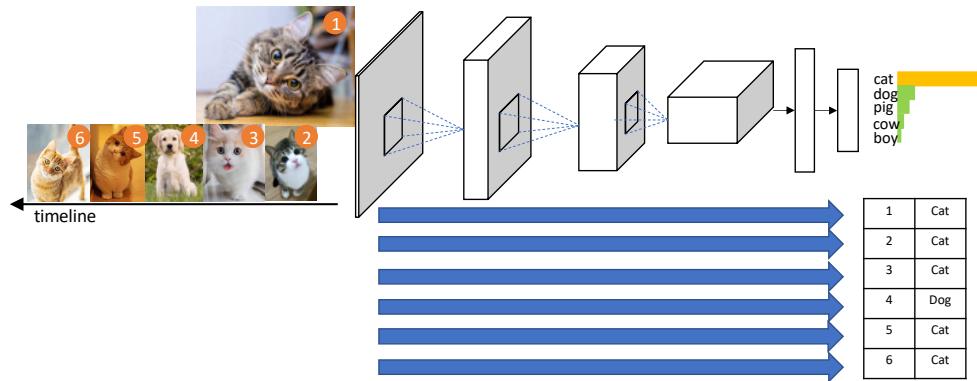


(a) The priming effect

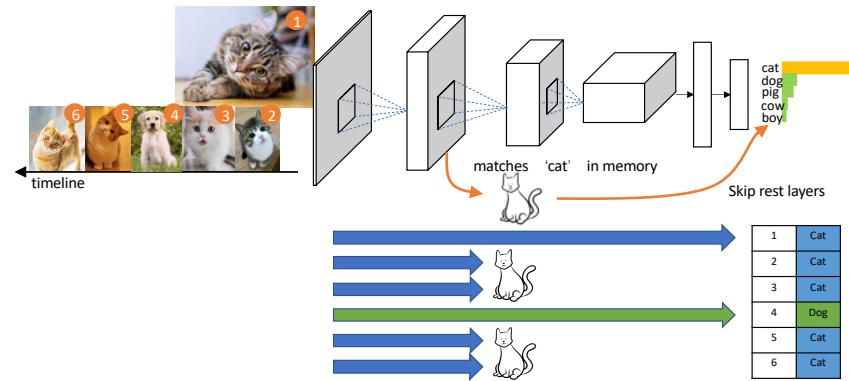


# Motivation

- Infuse the priming effect with CNN inference

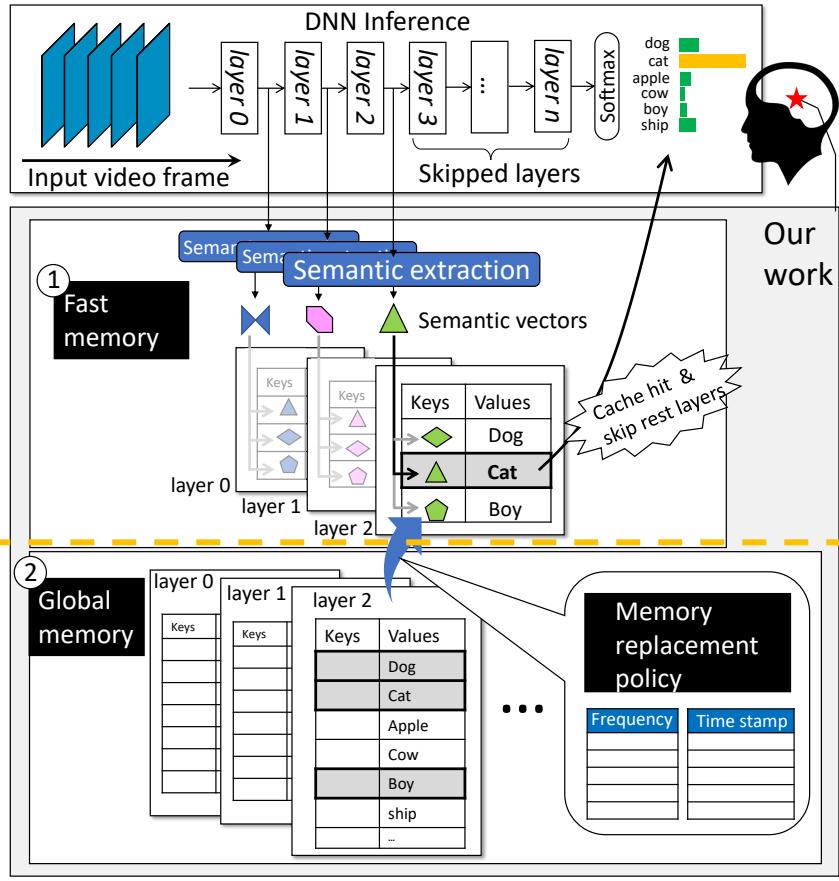


Traditional CNN model: carry out fully inference every time



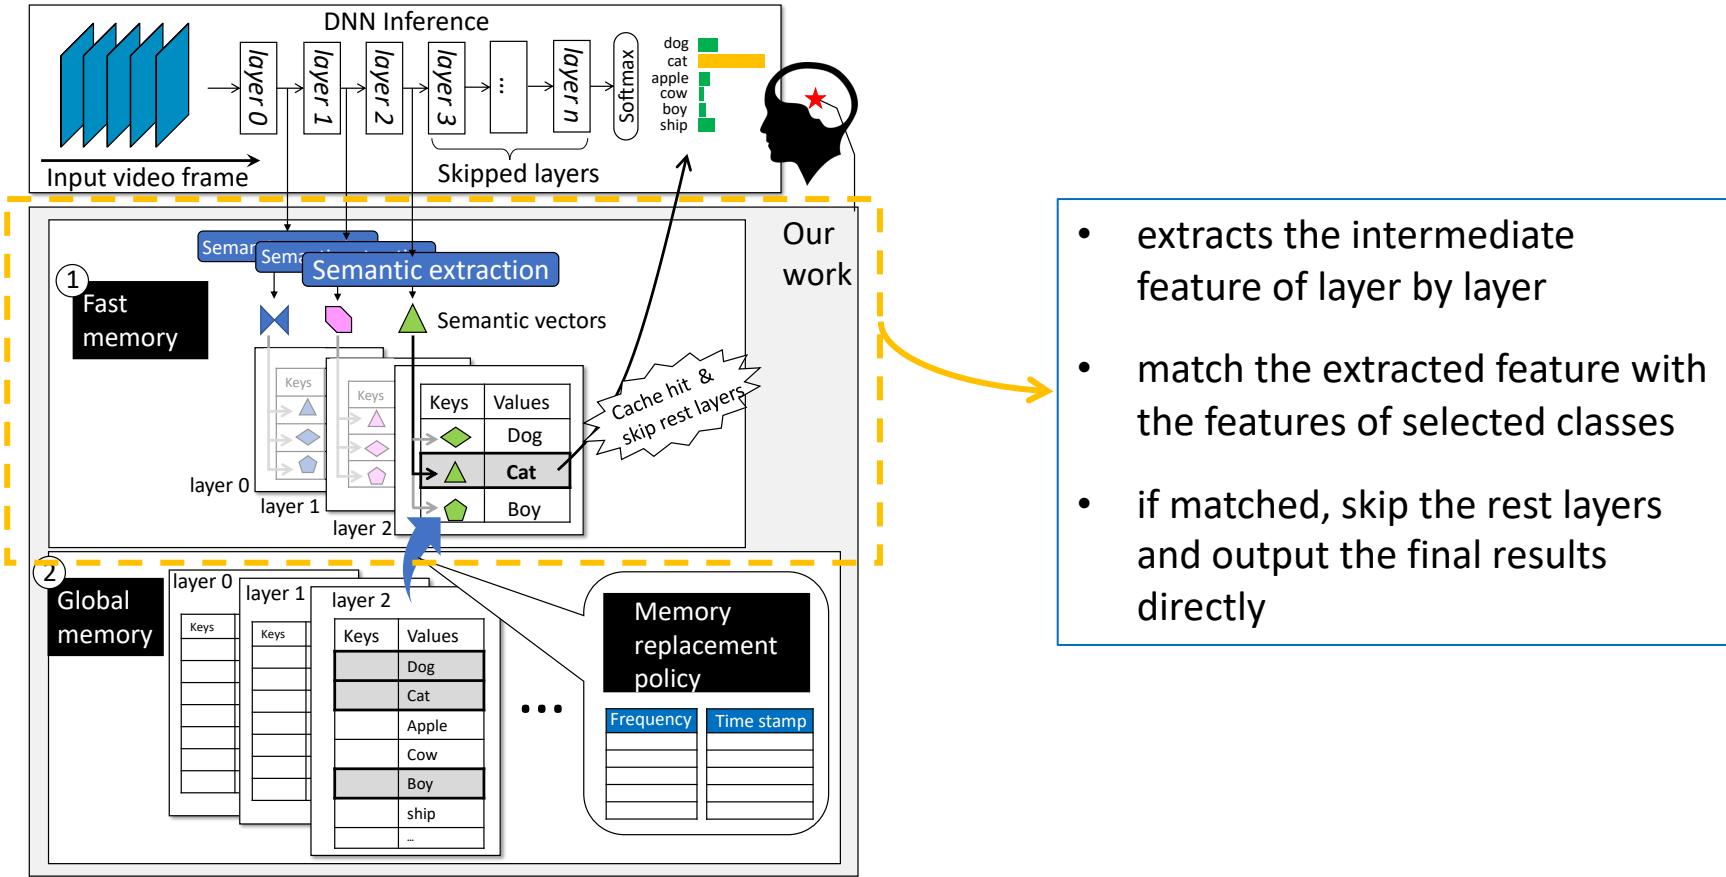
Rarely seen objects: fully inference  
Recently seen objects: early exit

# Our proposal: semantic memory (SMTM)



- store the semantic centers of all classes
- Memory Replacement Policy: cache the frequency and time stamp of recently seen classes
- select few classes with the greatest probability of recurring

# Our proposal: semantic memory (SMTM)



# Challenges

- Efficient memory encoding against CNN models' over-parameterization
  - extremely large volume of intermediate data
  - directly look up feature maps is cumbersome
  - take about 10ms even with GPU acceleration
- Obtaining speedup by high-level vision semantics
  - previous methods: low-level vision information
  - human brain: makes recognition by high-level features
  - traditional execution flow can not reuse semantics
- Battling dynamics on scenario variation
  - the scene change drastically
  - the scene complexity is not known in advance
  - real scenario data  $\neq$  training data, more complicated

# SMTM tech#1: Semantic Memory Encoding

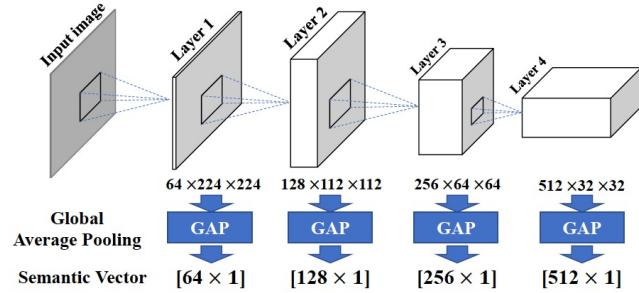


Figure 3: Semantic vectors extraction.

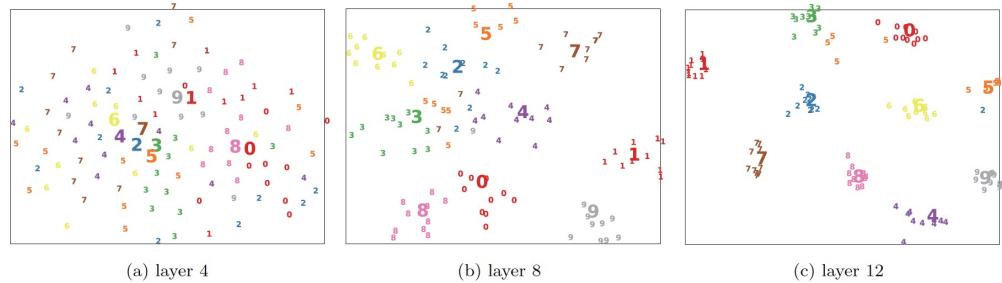


Figure 4: Visualized separability of semantic vectors for different VGG16 layers, showing that going deeper the semantic vectors can be more accurately separated.

- adopt **global average pooling (GAP)** to perform dimensionality reduction
  - much more light-weight
  - an effective indicator: clear separability in hidden layers
- adopt the **cosine distance** to evaluate the distance of different objects

$$\cdot s_j^l = \xi(SV^l, SC_j^l) \in [-1,1], j \in [1, n]$$

$$\cdot sep^l = \frac{s_H^l - s_{SH}^l}{s_{SH}^l}$$

# SMTM tech#2: Early Exit

- The separability in shallow layers is **not as strong or stable as the deeper layers**
- The **cross-layer cumulative similarity**
  - $SA_j^l = \sum_{l_0}^l s_j^{l_0} \times weight_{l_0}, j \in [1, n]$
  - $weight_{l_0} = 2^{l_0-1}, 1 + 2^1 + 2^2 + \dots + 2^{n-1} = 2^n - 1$
- the accumulated confidence (AC)**
  - $SA_H^l$  : the highest similarity accumulation result
  - $SA_{SH}^l$  : the second-highest result
  - $AC^l = \frac{SA_H^l - SA_{SH}^l}{SA_{SH}^l}$
  - If  $AC^l > global\ threshold$ , exit!

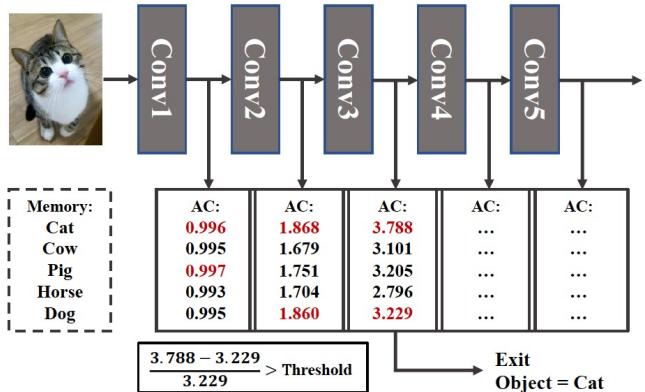
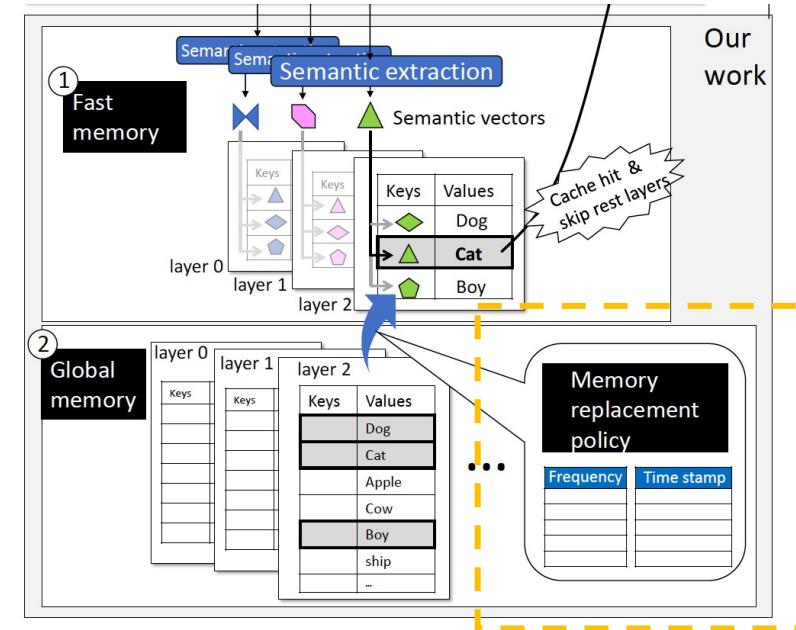


Figure 5: Memory look up by accumulated confidence (AC) metric. Memory: objects in memory.

# SMTM tech#3: Adaptive Priming Memory

1) cache replacement policy   2) adaptive cache size   3) adaptive semantic centers

- Frequency table:
  - keeps a record of the number of times that each object class presented in history
- Time-stamp table
  - keeps a record of the recency of each object class
  - the forgetting mechanism
$$\psi_i = \psi_i \times (0.25)^{\frac{TS_i}{W}}$$
- The replacement policy
  - takes the Top-k highest score
$$Score_i = Score_i \times (0.25)^{\frac{TS_i}{W}}$$
  - cache the Top-k objects in the fast memory



# SMTM tech#3: Adaptive Priming Memory

## 1) cache replacement policy    2) adaptive cache size    3) adaptive semantic centers

### ✓ Adaptive cache size

- probability estimation method
- $P(\theta \in \Psi) = \sum_{i=1}^k \frac{score_i}{\sum_{i=1}^n score_i}$ 
  - Confidence Level (CL) 95%
  - adjust the  $k$ ,  $P(\theta \in \Psi) > CL$
- The experiments show a 21.6% hit ratio improvement

### ✓ Adaptive semantic centers

- warms up using the training data
- update in weighted average manner
  - $\widehat{SC}_{l_0}^j = \frac{SC_{l_0}^j \cdot m_{l_0}^j + SV_{l_0}^j}{m_{l_0}^j + 1}$
- The experiments show a 16.9% accuracy improvement

# SMTM Implementation

- Computing framework:
  - ncnn
- Test Platform:
  - Google Pixel 4XL (Qualcomm Snapdragon 855 Processor)
- Datasets:
  - Action Recognition (UCF-101), Classification (long-tail Cifar-100)
- Five popular CNN models:
  - AlexNet, GoogleNet, ResNet50, MobileNet V2, VGG16
- Five evaluation metrics:
  - latency improvement
  - accuracy loss
  - energy saving
  - memory overhead
  - early exit ratio

# SMTM Evaluation

- latency improvement

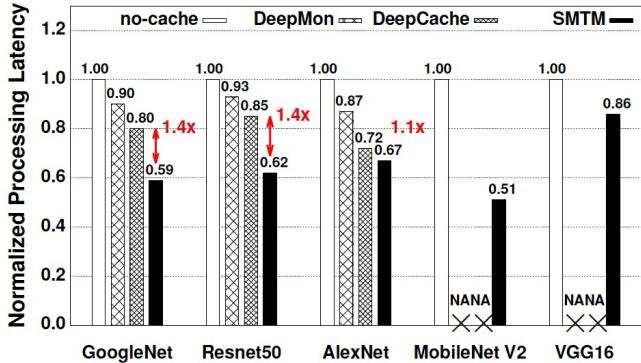


Figure 6: Average processing latency with CPU (w/o SIMD) on action recognition (AlexNet, GoogleNet, ResNet50, MobileNet V2) and classification (VGG16). ‘NA’: ‘not applicable’. SMTM speedup the processing time by  $1.1\times$ - $1.4\times$  comparing to DeepCache [52], and  $1.3\times$ - $1.5\times$  comparing to DeepMon [24]. DeepCache’s and DeepMon’s implementation is not compatible with the two models MobileNet V2 and VGG16, so we are not able to reproduce some results. ‘NA’: ‘not applicable’.

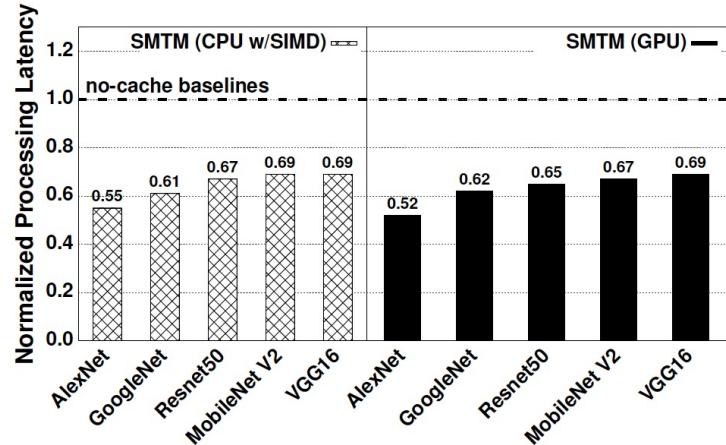


Figure 7: Average processing latency of SMTM with mobile CPU (w/SIMD) and mobile GPU on action recognition (AlexNet, GoogleNet, ResNet50, MobileNet V2) and classification (VGG16).

Mobile CPU and GPU: 30%-50% latency reduction

Compare with SOTA:  $1.1X - 1.5X$

# SMTM Implementation

- Accuracy Drop, Memory Overhead

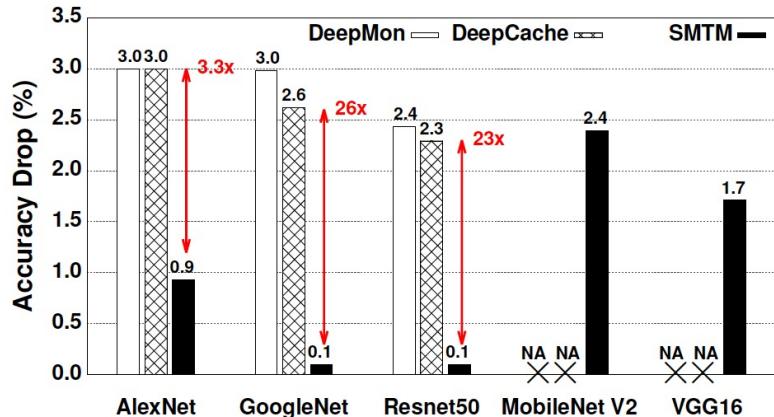


Figure 8: Top-1 accuracy drop of SMTM on action recognition (AlexNet, GoogleNet, ResNet50, MobileNet V2) and classification (VGG16). ‘NA’: ‘not applicable’.

Accuracy drop: 1% on average

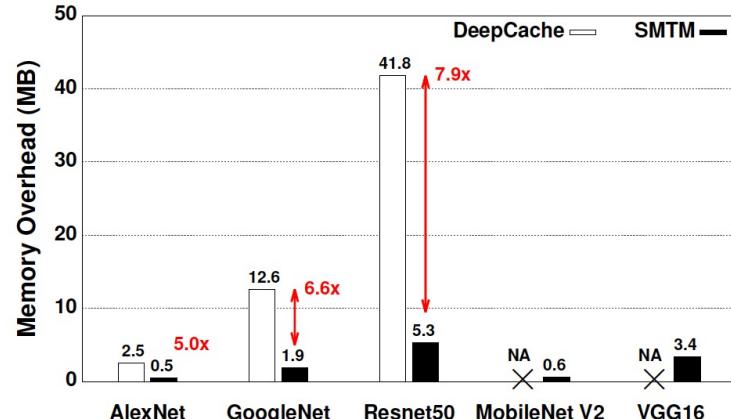


Figure 9: The memory overhead of SMTM on action recognition (AlexNet, GoogleNet, ResNet50, MobileNet V2) and classification (VGG16). ‘NA’: ‘not applicable’.

Memory overhead:

20% of SOTA, less than 5% of original models

# SMTM Implementation

- Energy Saving, Early Exit Ratio, the Effect of the Threshold

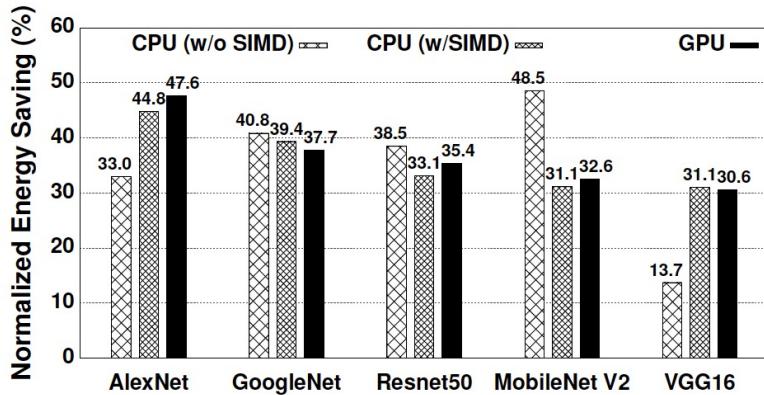
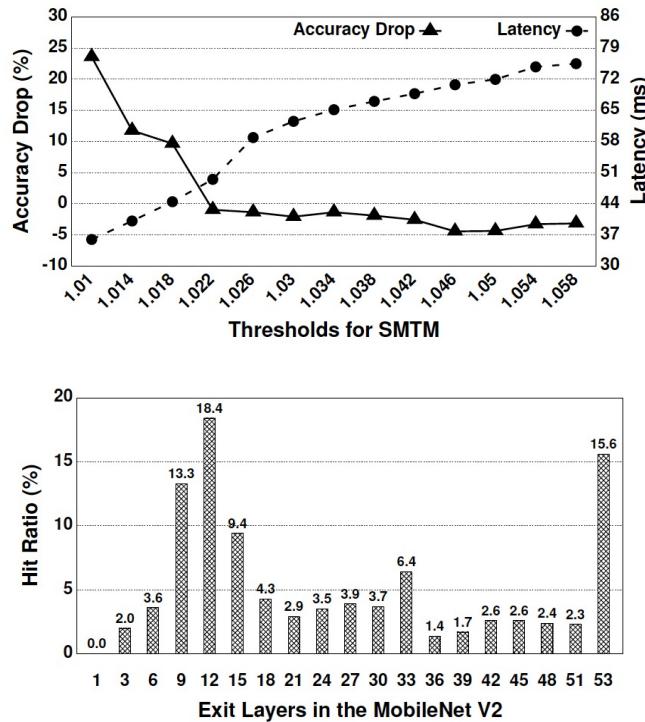


Figure 10: The energy saving ratio with different devices on action recognition (AlexNet, GoogleNet, ResNet50, MobileNet V2) and classification (VGG16).

Energy saving: 36% on average



# SMTM Implementation

- The performance of adaptive semantic memory

	Hit ratio	Latency reduction
SMTM (Constant)	65.39%	25.21%
SMTM (Adaptive)	87.00%	38.46%

Table 1: The impact of adaptive cache size. Tested on ResNet50 model.

21.61% hit ratio improvement

13.25% acceleration

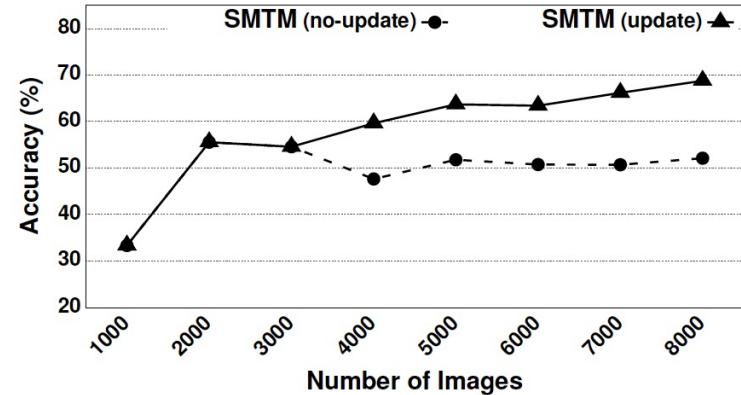


Figure 13: The impact of adaptive semantics center on the prediction accuracy on ResNet50.

16.9% accuracy improvement

# Summary

- SMTM: a novel memory mechanism to accelerate CNN-powered mobile vision by infusing the priming effect with CNN inference.
  - speeds up CNN inference for the frequently and recently-seen objects
  - an accurate yet low-cost memory encoder
  - an early exit method
  - an adaptive priming memory policy
- prototype on commodity engine, evaluate on 5 CNN architectures, 2 datasets, on both mobile CPU/GPU
  - Mobile CPU and GPU: 30%-50% latency reduction
  - Only 5% memory overhead

Thank you for watching!