



PieBridge: Fast and Parameter-Efficient On-Device Training via Proxy Networks


Wangsong Yin¹, Daliang Xu¹, Gang Huang¹, Ying Zhang¹,
Shiyun Wei², Mengwei Xu³, Xuanzhe Liu¹

¹Peking University, ²Zhongguancun Laboratory,

³Beijing University of Posts and Telecommunications

ACM SenSys 2024


On-Device Training NNs



Face detection



Voiceprint verification



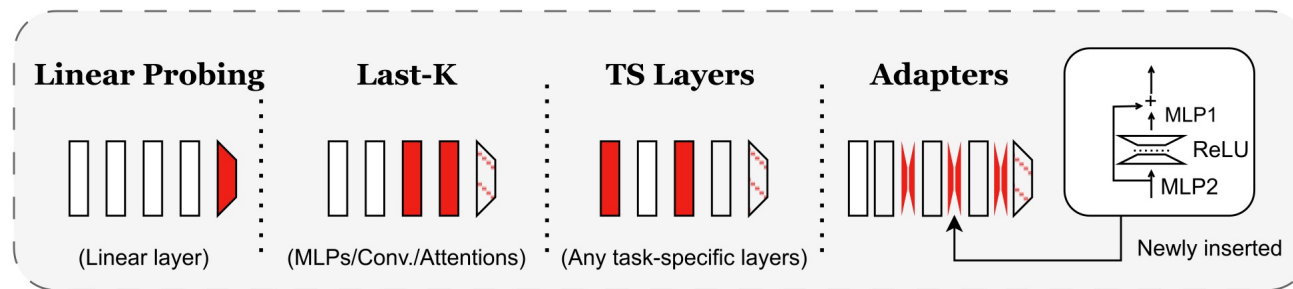
Input habit prediction

Training (fine-tuning) NNs on **devices**
maximizes the protection of **privacy** data!


Popular Paradigm: Parameter-Efficient Training

Parameter-Efficient Training (PET)

Achieving on-par accuracy with training all parameters
by only tuning **a tiny portion** of parameters



↓ *A general abstraction*



Legend:
Blue box: Frozen block Red box: Trainable layer Solid arrow: Forward pass Dotted arrow: Backward pass
Pink box: Trainable block Dashed box: Frozen/trainable blocks Dotted arrow with red X: Backward pass stopping

(Typically static) **trainable/frozen** blocks

Popular Paradigm: Parameter-Efficient Training

Parameter-Efficient Training (PET)


Achieving on-par accuracy with training all parameters
by only tuning **a tiny portion** of parameters

PET Method	Trainable layers	Formula	Trainable Parameters
Linear Probing	The last linear layer	$y = x \cdot W + b$	0.69%
Last-K	The last K layers	N/A	$\geq 8.3\%$
TS Layers	Task-specific layers (Norm. layers here)	N/A	0.87%–2.80%
Adapters	The inserted adapters	$y = x + \sigma(x \cdot W_{down}) \cdot W_{up}$	0.71%

Trainable parameters are very few! (e.g. usually **<1%**)

Key Problem: Param. Efficiency \neq Time Efficiency

Daily idle time on mobile devices^[1]



Training time is still too long!



Model	PET Config.	Upd. Paras.	Frozen Layers		Trainable Layers	
			Time (Sec.)	FLOPs (G)	Time (Sec.)	FLOPs (G)
ViT_base	Linear Probing	0.69%	0.46	67.44	1.88×10^{-3}	0.36
	Adapters (6)	1.75%	0.64	134.86	5.44×10^{-3}	0.54
	Adapters (12)	2.80%	0.84	202.31	1.24×10^{-2}	0.71
ResNet50	Linear Probing	0.81%	0.06	16.53	4.36×10^{-4}	4×10^{-2}
	Last-K (K=3)	14.3%	0.07	15.59	2.56×10^{-3}	2.15

Time breakdown in a PET iteration



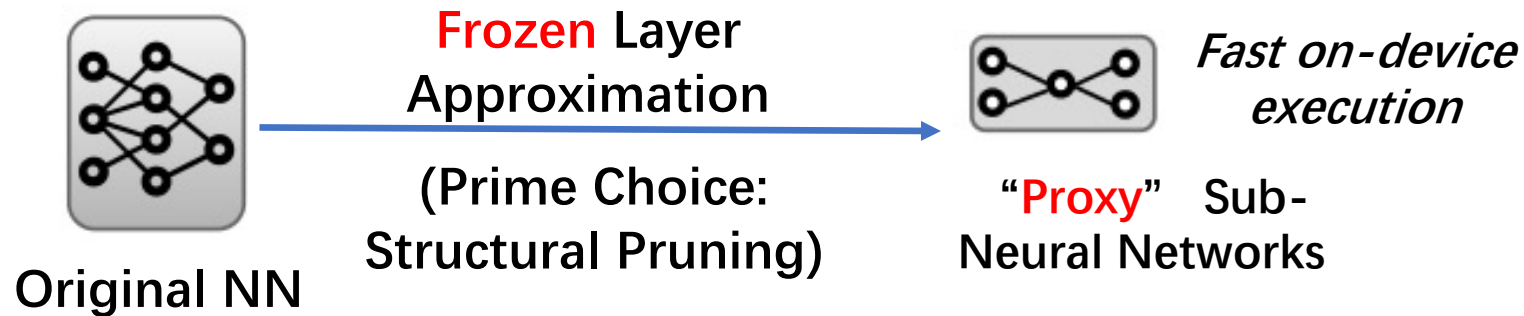
Time consumption:
Frozen \gg trainable blocks!

[1] Mengwei Xu, etc. 2018. Deeptype: On-device deep learning for input personalization service with minimal privacy concern. Proceedings of the ACM on IMWUT (2018), 1–26

Frozen Layer Approximation During Training



Can we *approximate* the frozen blocks for a faster on-device PET?



Convergence accuracy drops significantly !

Training all samples by a proxy sub-NN with 50% retained compute leads to an accuracy drop of up to 29.11%!

Opportunity: Training-Data Diversity


Difficulty




Clean



Blurry



Obstructed



Complicated

Importance




British Shorthair



Tabby Cat 1



Tabby Cat 2



Tabby Cat 3

- Data **Difficulty**

- Samples differ in the difficulty level of extracting high quality features.
- Measure data difficulty by the minimal subnet that captures feature that on-par with the non-approximated network.

$$\begin{aligned} \phi(x) = \max(n \in \{0, 1, \dots, N\}), \\ \text{s.t. } \mathcal{F}_n(x) \simeq \mathcal{F}_0(x). \end{aligned} \quad (1)$$

- Data **Importance**

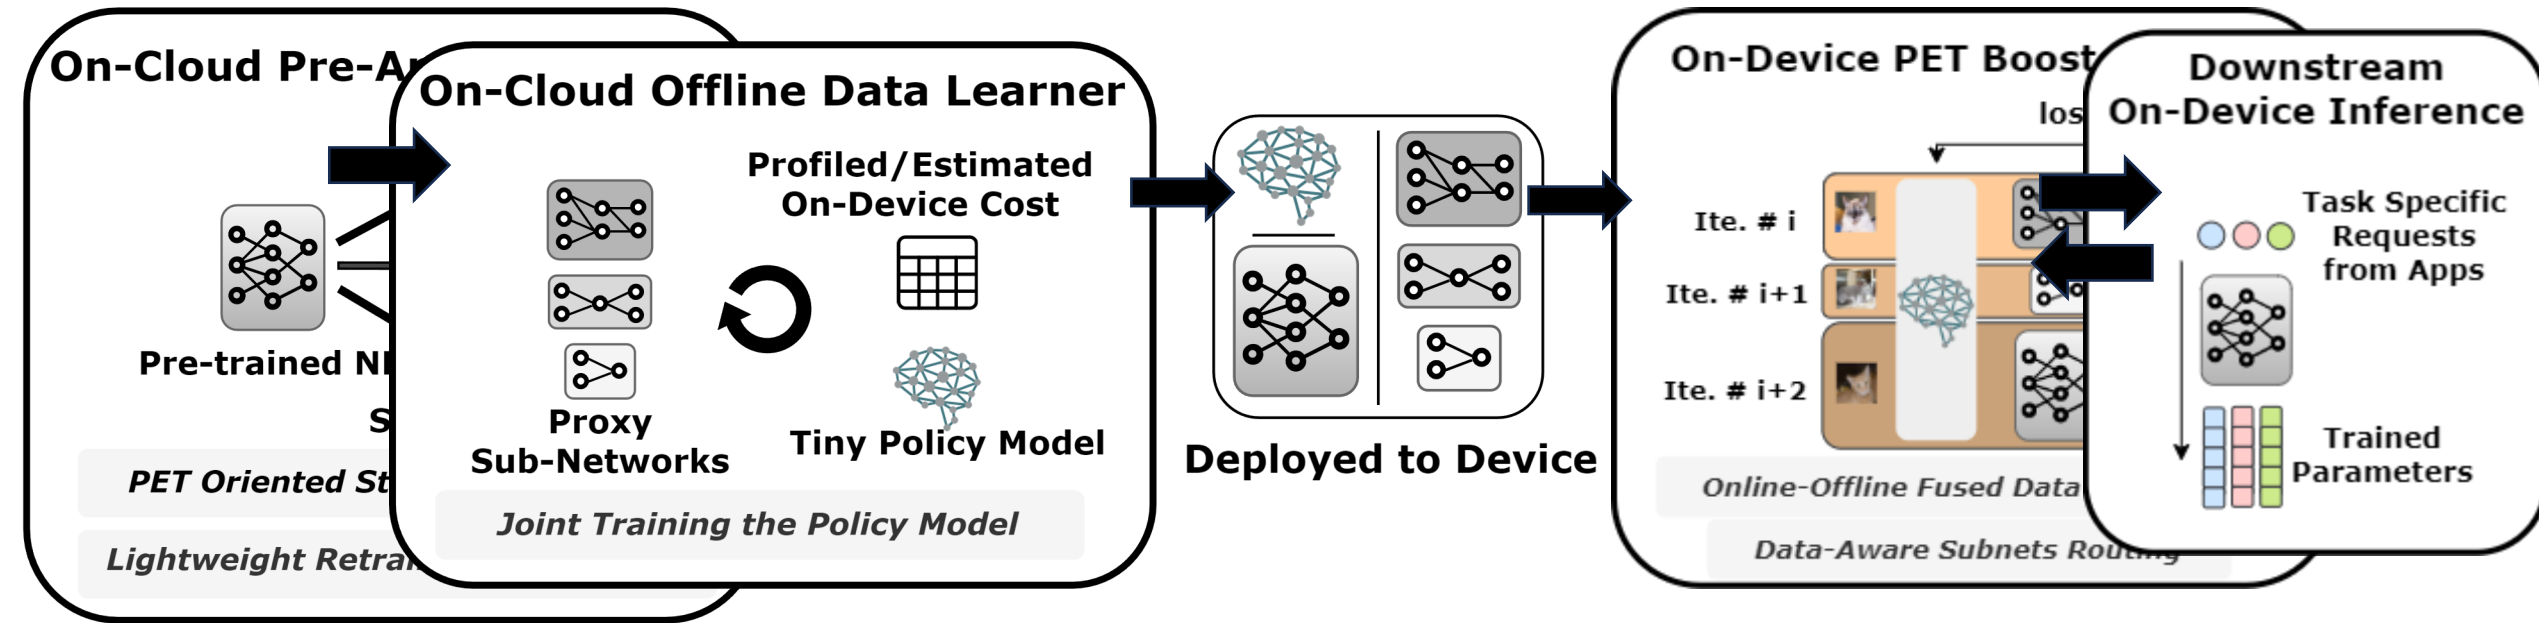
- Samples differ in the contribution to model convergence.
- Measure data importance by loss value.

$$\psi(x, t) = l(x; w_t) = \mathcal{L}[\mathcal{F}(x; w_t), y], \quad (2)$$



Dynamically assign the most suitable approximation for diverse training data!

Our System: PieBridge



Cloud Offline Stage

Generating multi-level subnets
Learning a tiny policy model

Device Online Stage

Perceiving diff./imp. of on-device data
Routing training data to proper subnet


Stage#1: Cloud Offline



How to generate FAST and ACCURATE subnets with fine-grained tradeoffs?

INT8 Quantization

✗ Limited benefit on the same processor



Sparsification

✗ Limited support from on-device NN libs



Early Exiting

✗ Too coarse granularity (a whole NN layer)




Ours: **structural pruning with heuristic**



E2e latency reduction



Fine-grained tradeoff



Identifying and pruning dimensions that are **independent** to trainable blocks



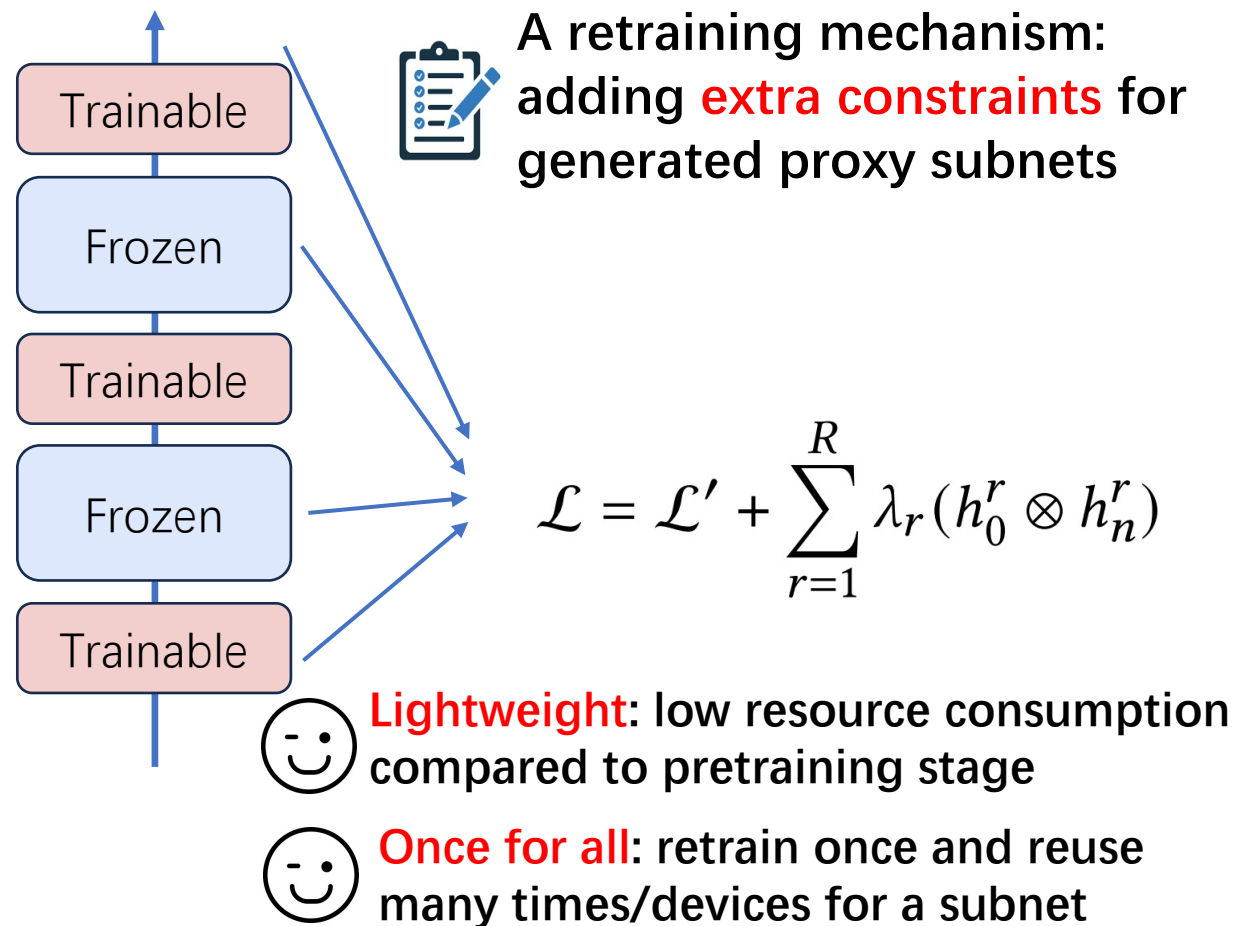
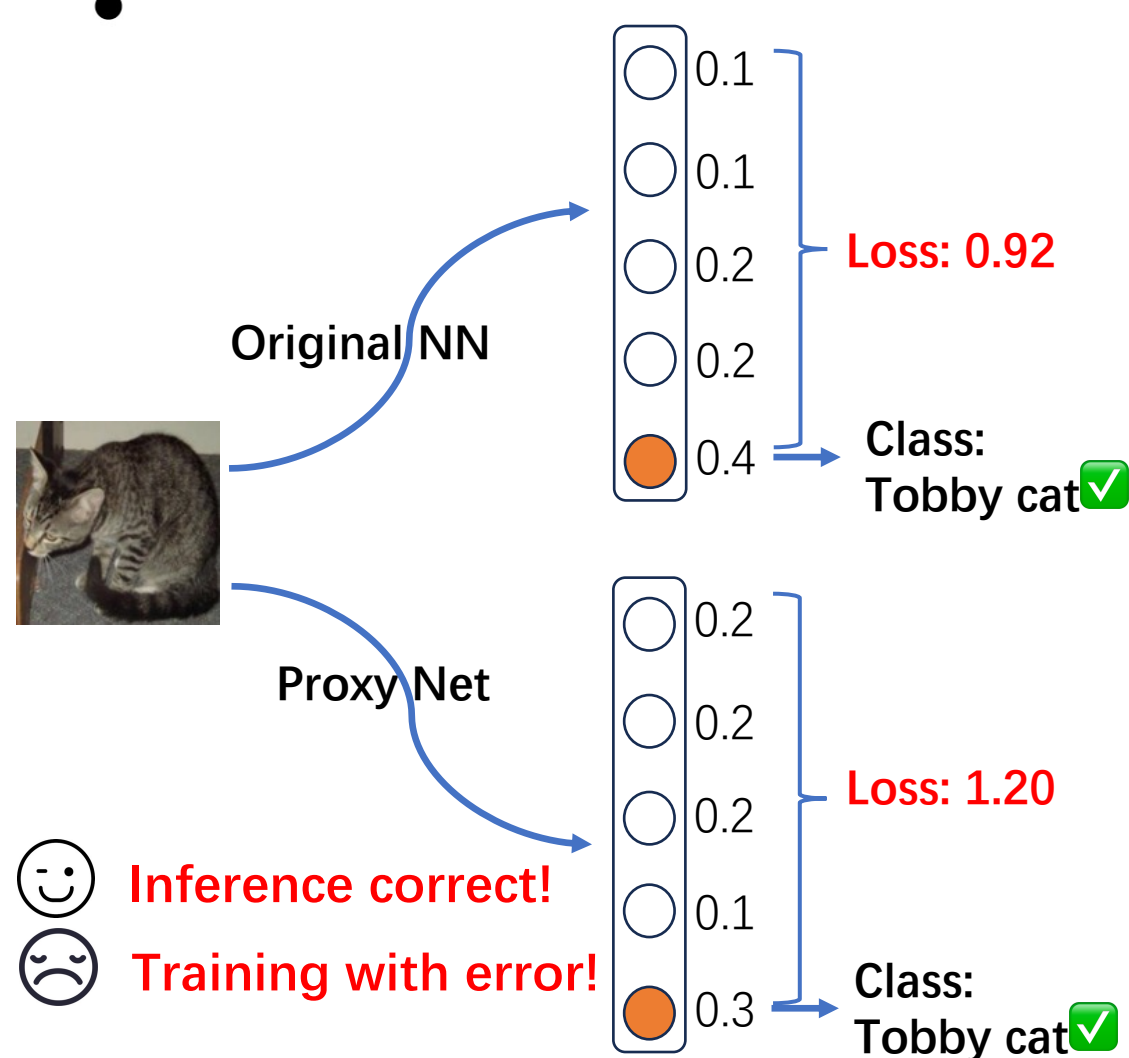
Multiple level proxy nets for a fine-grained approximation



Generating subnets with an **even** pruning step

Stage#1: Cloud Offline

! Is directly generated (pruned) subnet a good approximation? Not yet!




Stage#1: Cloud Offline

- ! Perceiving data diversity online introduces extra
- system level overhead for on-device training...

Needs traversing all proxy subnets


$$\phi(x) = \max_{n \in \{0, 1, \dots, N\}} (1) \\ \text{s.t. } \mathcal{F}_n(x) \simeq \mathcal{F}_0(x).$$

✗ Up to 10x slower training!




What is the corresponding proxy subnet of current sample difficulty?

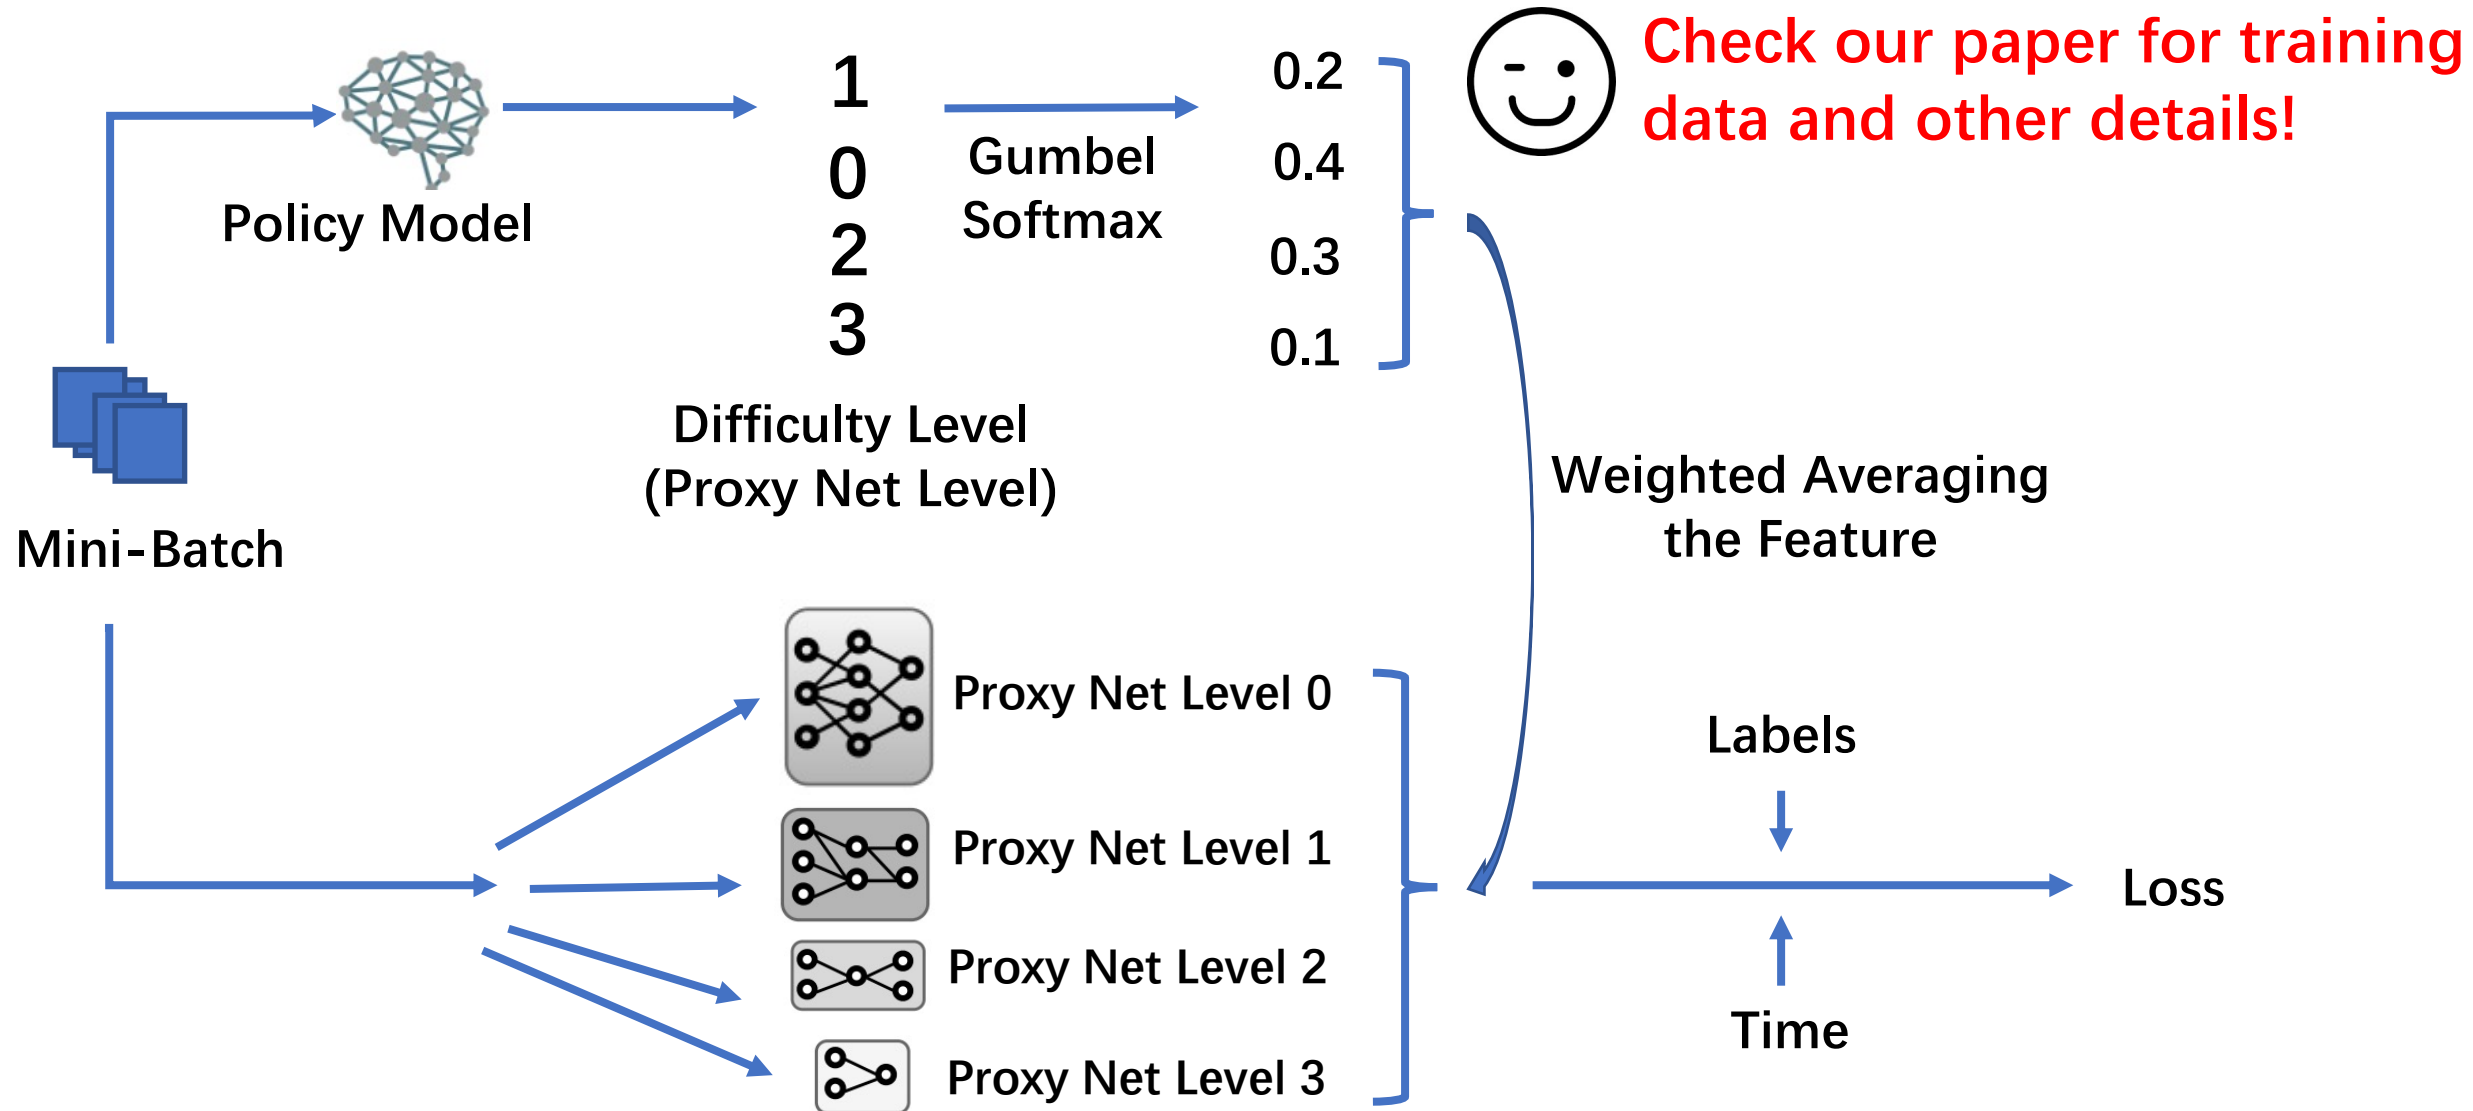
You have to calculate it **during training** to find out...



Can we offline learn data difficulty for a low overhead online prediction ?




Stage#1: Cloud Offline



Stage#2: Device Online




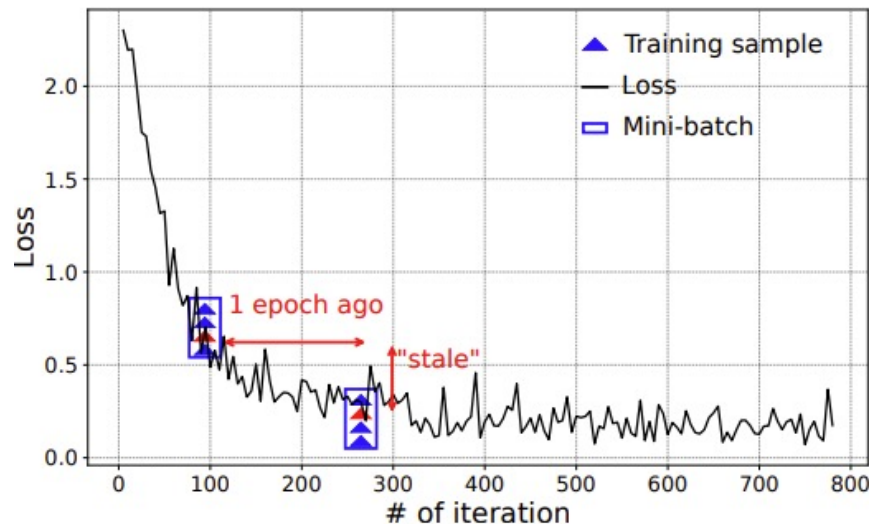
Intra mini-batch scheduling proxy nets: inefficient!



Since the mini-batch size is tiny on wimpy mobile devices (e.g. 4/8), lets schedule the proxy network assignment at **mini-batch granularity**!

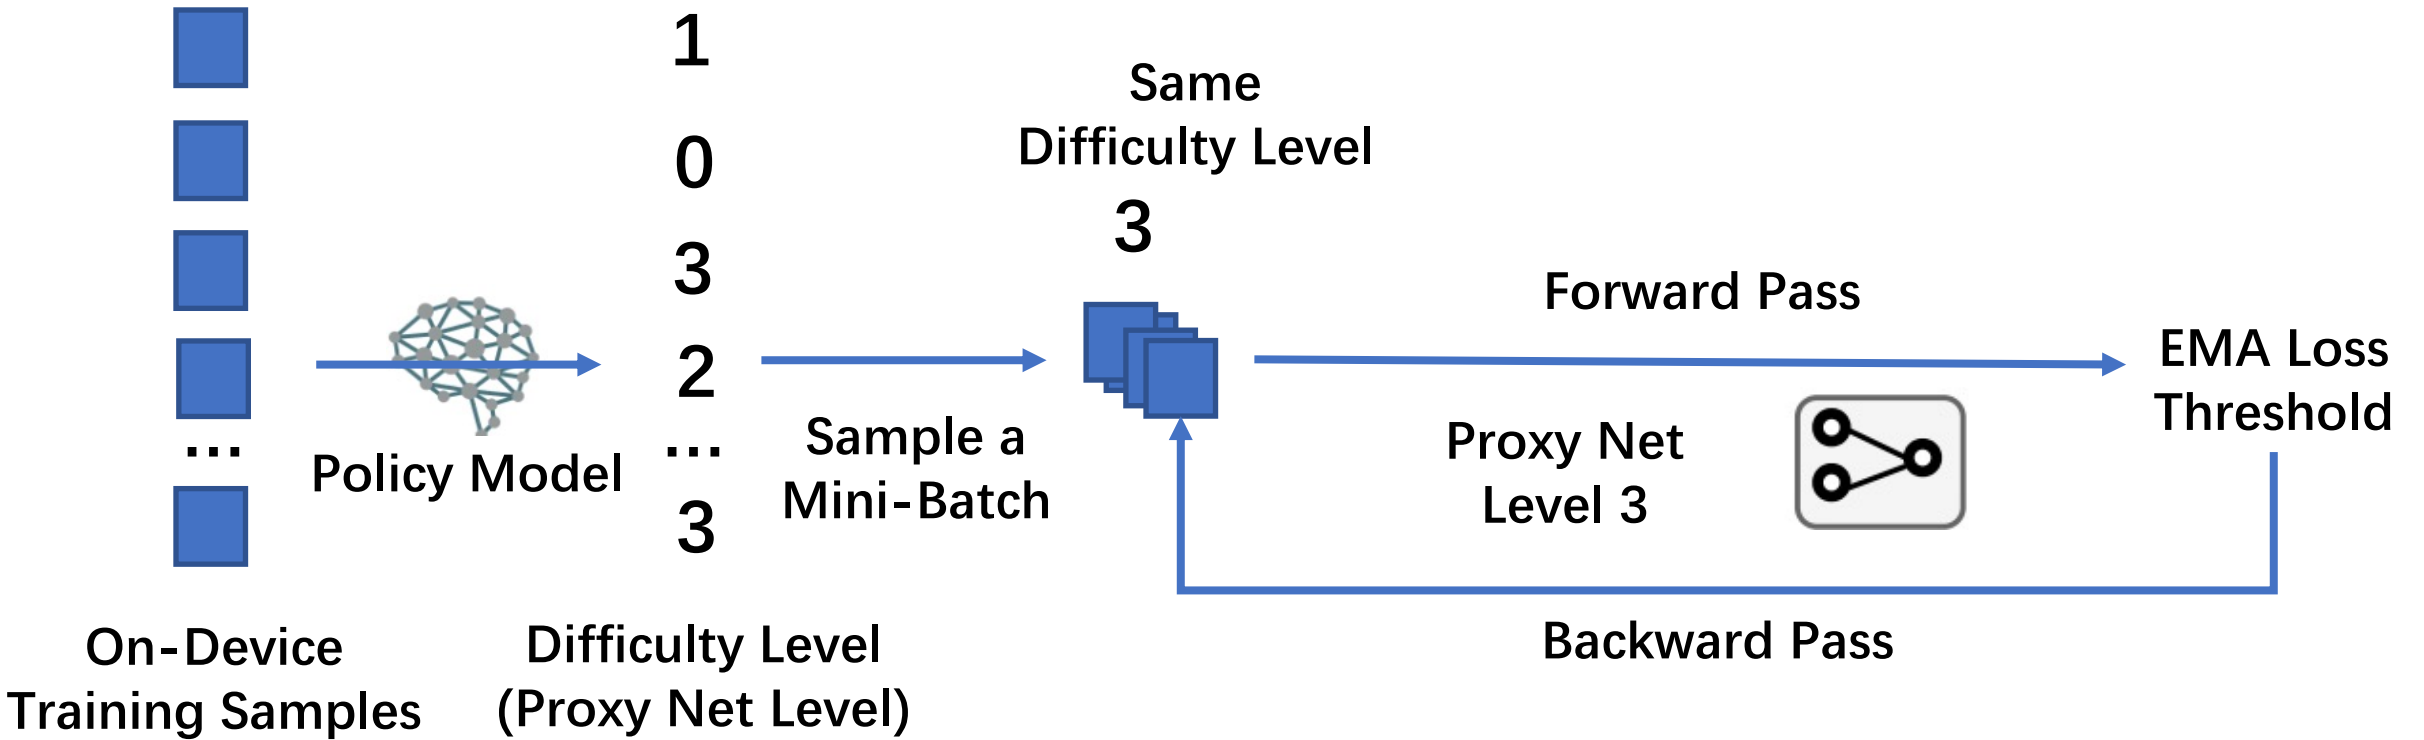


Data importance: determined by current training state!

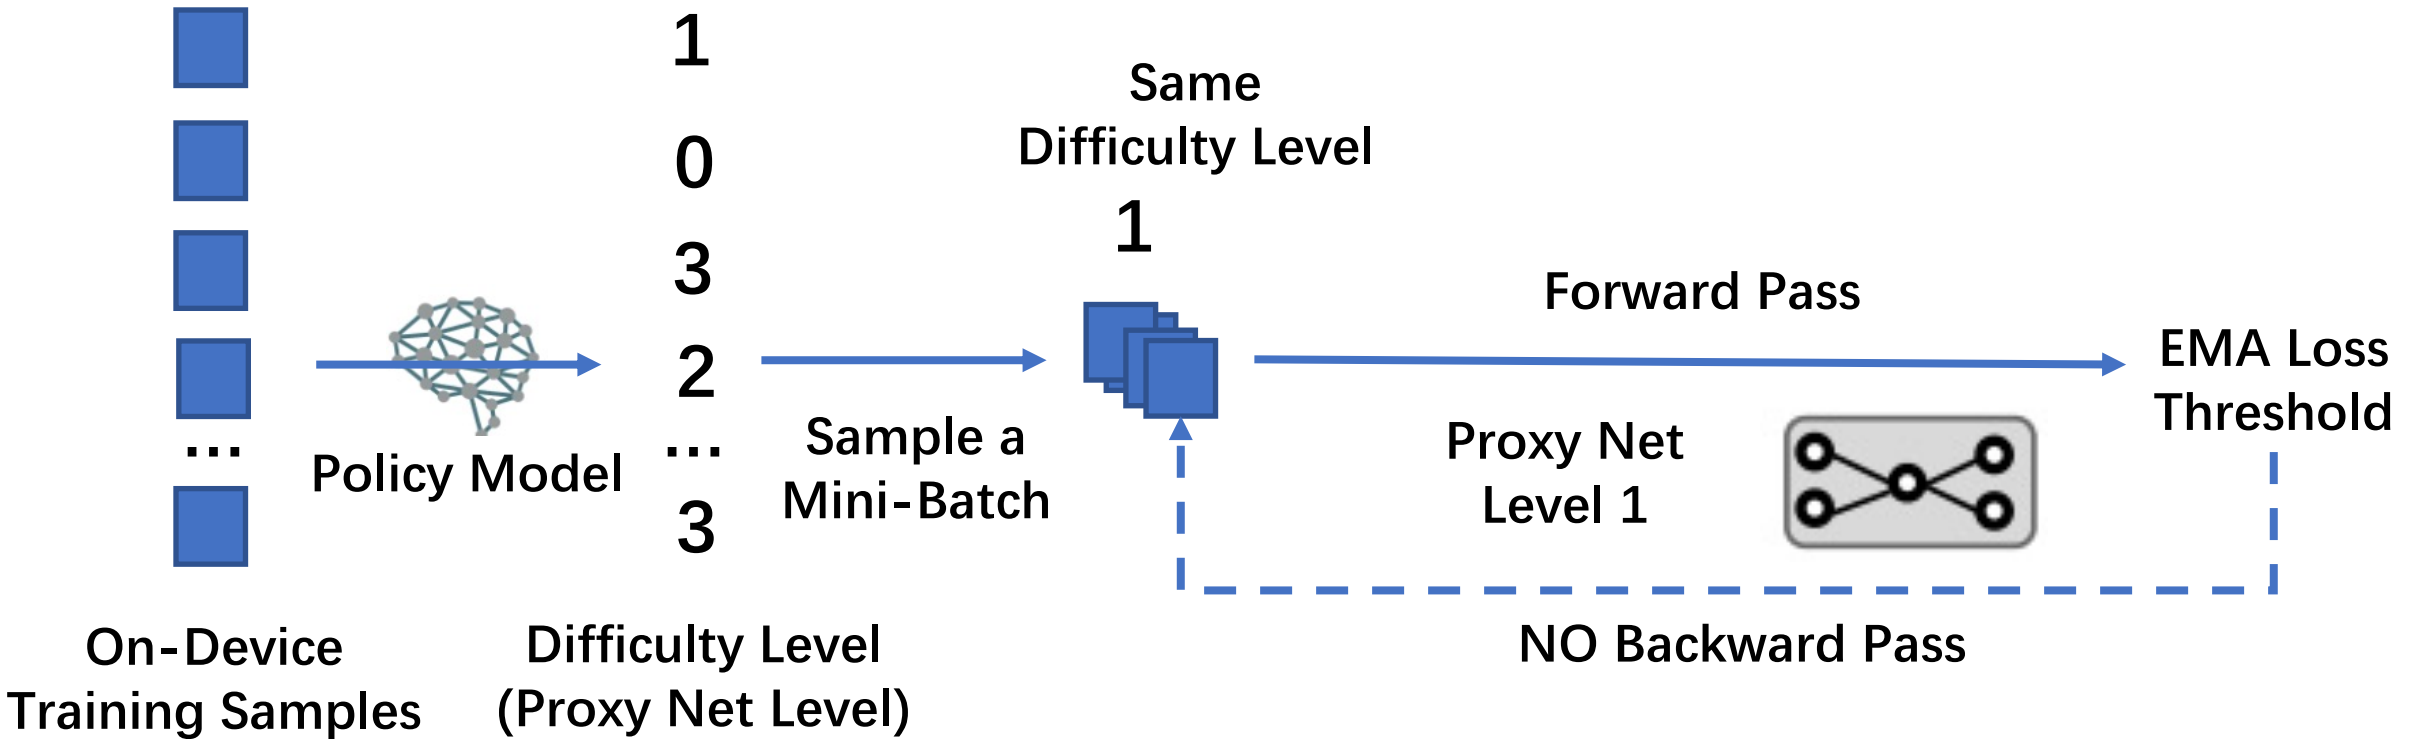


Lets harness it **only for backward stage**!

Stage#2: Device Online



Stage#2: Device Online



Please refer to the Algo. 1 in our paper for more details...

Implementation and Evaluation

- **Datasets**

- Cloud (pre)training
 - ImageNet2012
- On-device training
 - Caltech101
 - Caltech256
 - DogsvsCats
 - DTD

- **Baselines**

- Full parameter training (FT-All)
- Parameter efficient training (PET)
 - ResNet50 : Linear Probing
 - MobileNetV3 : Last-3
 - ViT_base : Adapters(12)
- PruneTrain (PT) [1]
- ElasticTrainer (ET) [2]

- **Models**

- MobileNetV3-L
- ResNet50
- ViT_base

Name	Processor	Software env.
Jetson TX2 [12]	Dual-Core NVIDIA Denver 2 64-Bit CPU, 256-core NVIDIA Pascal™ GPU.	Ubuntu 18.04 LTS PyTorch 1.7.1.
RPI 4B [14]	Broadcom BCM2711B0 quad-core A72 64-bit @ 1.5GHz CPU.	Raspbian 11, PyTorch 1.7.1.
MI 10 [10]	2.84GHz Cortex-X1, 3× 2.4GHz Cortex A78, 4× 1.8GHz Cortex A55 CPU.	Android 10, MNN 2.0.0, ONNX 1.13.1.
Huawei Mate 30 [7]	2x 2.86 GHz ARM Cortex-A76, 2x 2.09 GHz ARM Cortex-A76, 4x 1.86 GHz ARM Cortex-A55 CPU, Kirin 990 NPU.	


Device hardware/software

[1] Sangkug Lym, etc. 2019. PruneTrain: Fast Neural Network Training by Dynamic Sparse Model Reconfiguration

[2] Kai Huang, Boyuan Yang, and Wei Gao. 2023. ElasticTrainer: Speeding Up On-Device Training with Runtime Elastic Tensor Selection.

Wall-Clock Training Time

- Wall-Clock training time on standalone datasets



- Compared to baselines, PieBridge accelerates on-device training by up to **6.6x**
- Compared to vanilla PET, PieBridge accelerates on-device training by **1.4- -2.5x**

Training Accuracy

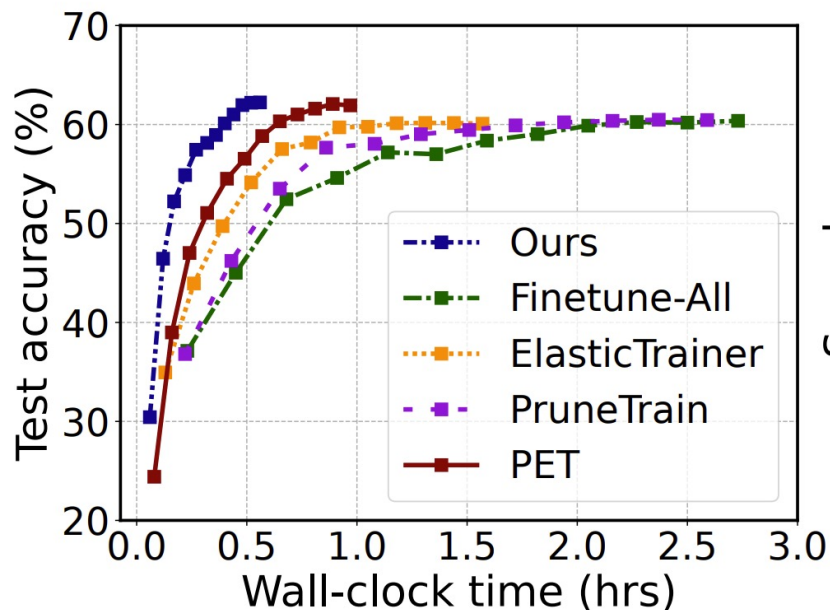
- Best convergence accuracy on standalone datasets

	MobileNetV3-L					ResNet50					ViT_base				
	FT-All	PET	PT	ET	<i>Ours</i>	FT-All	PET	PT	ET	<i>Ours</i>	FT-All	PET	PT	ET	<i>Ours</i>
C-101	87.66	84.27	87.76	89.20	89.82	91.96	91.13	91.23	91.09	92.03	94.53	96.00	93.98	95.64	96.12
C-256	72.00	72.50	73.08	71.72	73.51	83.93	82.38	82.94	80.81	82.89	85.61	84.93	81.39	81.83	84.91
DVC	95.35	95.50	95.15	95.20	94.60	95.55	95.25	95.25	95.55	95.50	95.40	95.96	95.05	95.50	95.00
DTD	53.19	50.27	50.48	50.69	52.87	60.37	62.93	60.48	60.16	62.23	65.05	66.12	62.55	66.54	67.71

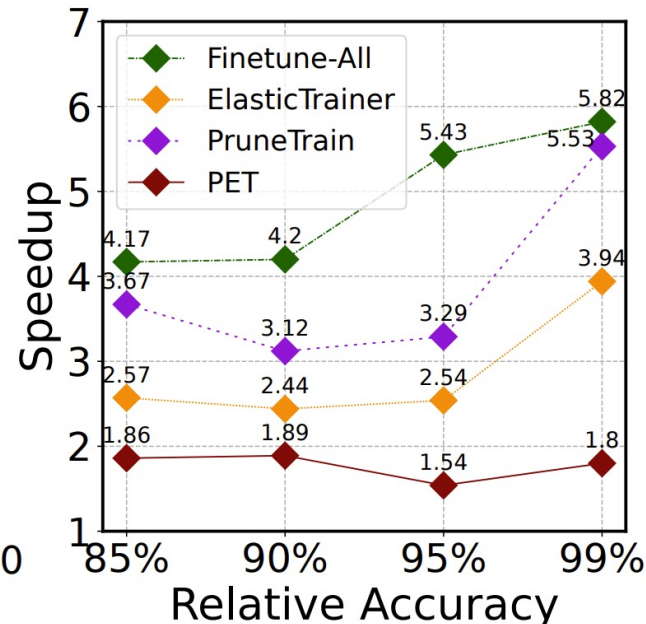
- Compared to baselines, PieBridge has no noticeable influence on accuracy (<2%)

Training Accuracy

- Performance under various accuracy budgets
 - There exists a marginal effect of training
 - Under an 85%/90%/95%/99% accuracy budget
 - PieBridge speeds up PET by **1.86/1.89/1.54/1.80x**
 - PieBridge speeds up FT-All by **4.17/4.20/5.43/5.82x**



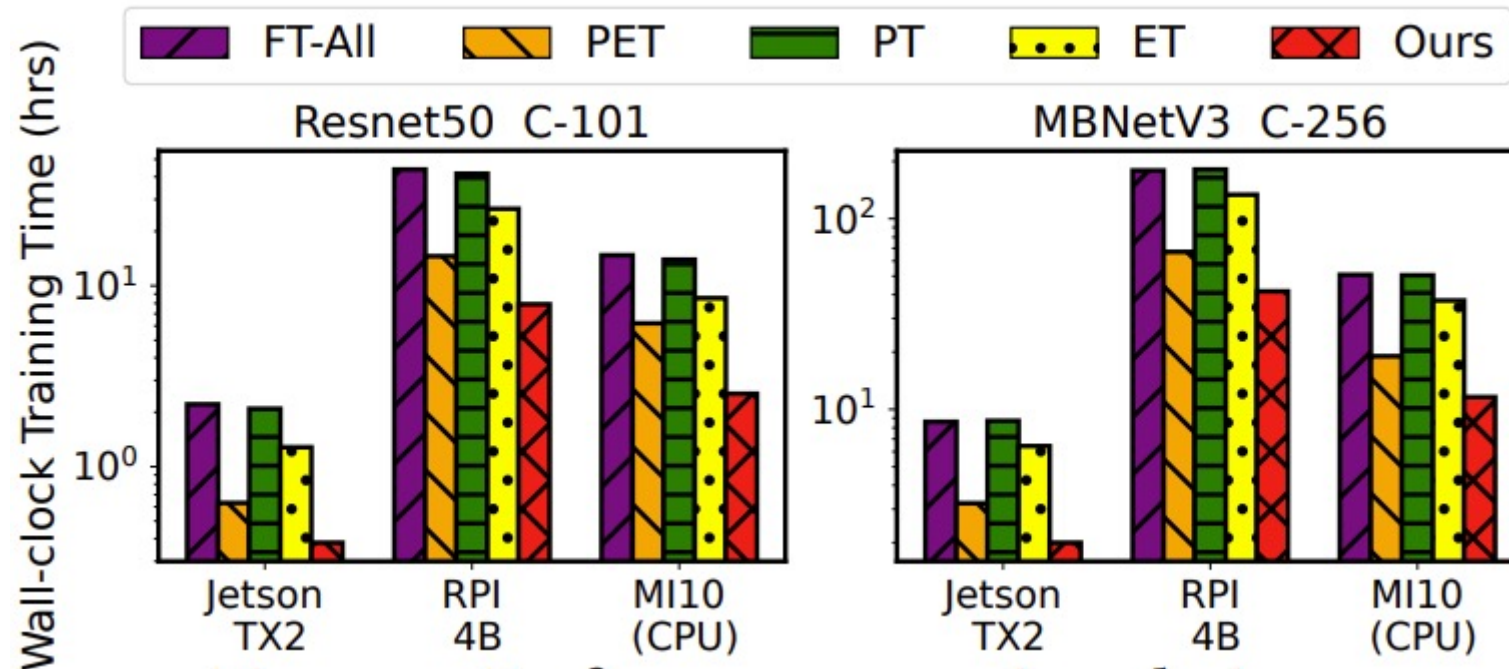
(a) Time to accuracy



(b) Speedup

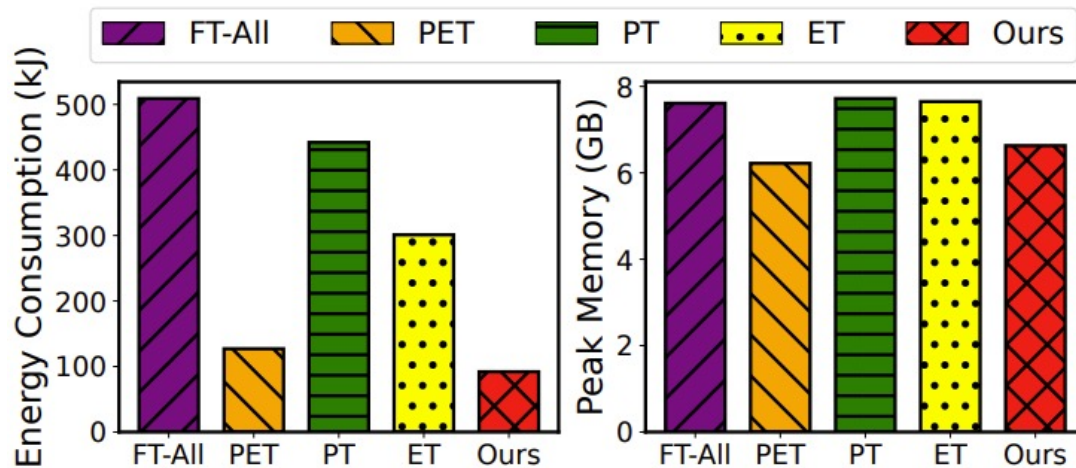
Performance on Various Devices

- Performance on various devices
 - Jetson TX2, Raspberry Pi 4B, MI 10 smartphone
 - Speeds up PET by **1.61/1.63/1.72x**
 - A clear performance gain on both CPU and GPU devices



Others

- Energy consumption
 - **5.53x** lower than FT-All
 - **1.28x** lower than PET
- Memory consumption
 - No significant increase of memory consumption compared to baselines
- Other sensitivity analysis
 - # of proxy nets (pruning step)
 - Policy model training data distribution
 - Few-shot learning performance
 - ...



Please refer to our paper!

Take Aways

- We present PieBridge, an on-device training framework with both **time- and parameter-** efficiency.
- PieBridge innovatively leverages **training data diversity** and **proxy subnetwork approximation** to reduce the overhead of frozen layers during training.
- PieBridge achieves up to **6.6x** training speedup compared to strong baselines, paving the way for personalized on-device AI.



yws@stu.pku.edu.cn

Code:

<https://github.com/yinwangsong/PieBridge>