



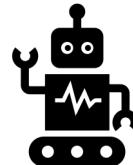
# 端侧大模型和智能体：挑战和尝试

# LLM-powered Mobile Devices

Mengwei Xu (徐梦炜)

Beijing University of Posts and Telecommunications

<https://xumengwei.github.io>



# Aren't they smart..Already?

- Yes, to a certain extent.



## DNN-embedded mobile apps

- Increased by almost **10x** (2018 to 2021)<sup>[1,2]</sup>
- Downloaded **billions of times** in one year
- Include almost every high-popularity app
- Up to **200+ DNNs** in a single app<sup>[3]</sup>

[1] Mengwei Xu, et al. "A First Look at Deep Learning Apps on Smartphones". In WWW 2019

[2] Mario Almeida, et al. "Smart at what cost? Characterising Mobile Deep Neural Networks in the wild". In IMC 2021.

[3] Through offline communication with application developers.

# Aren't they smart..Already?

- Yet, not even close to our expectation.



*“AI is a mirror, reflecting not only our intellect, but our values and fears.”*

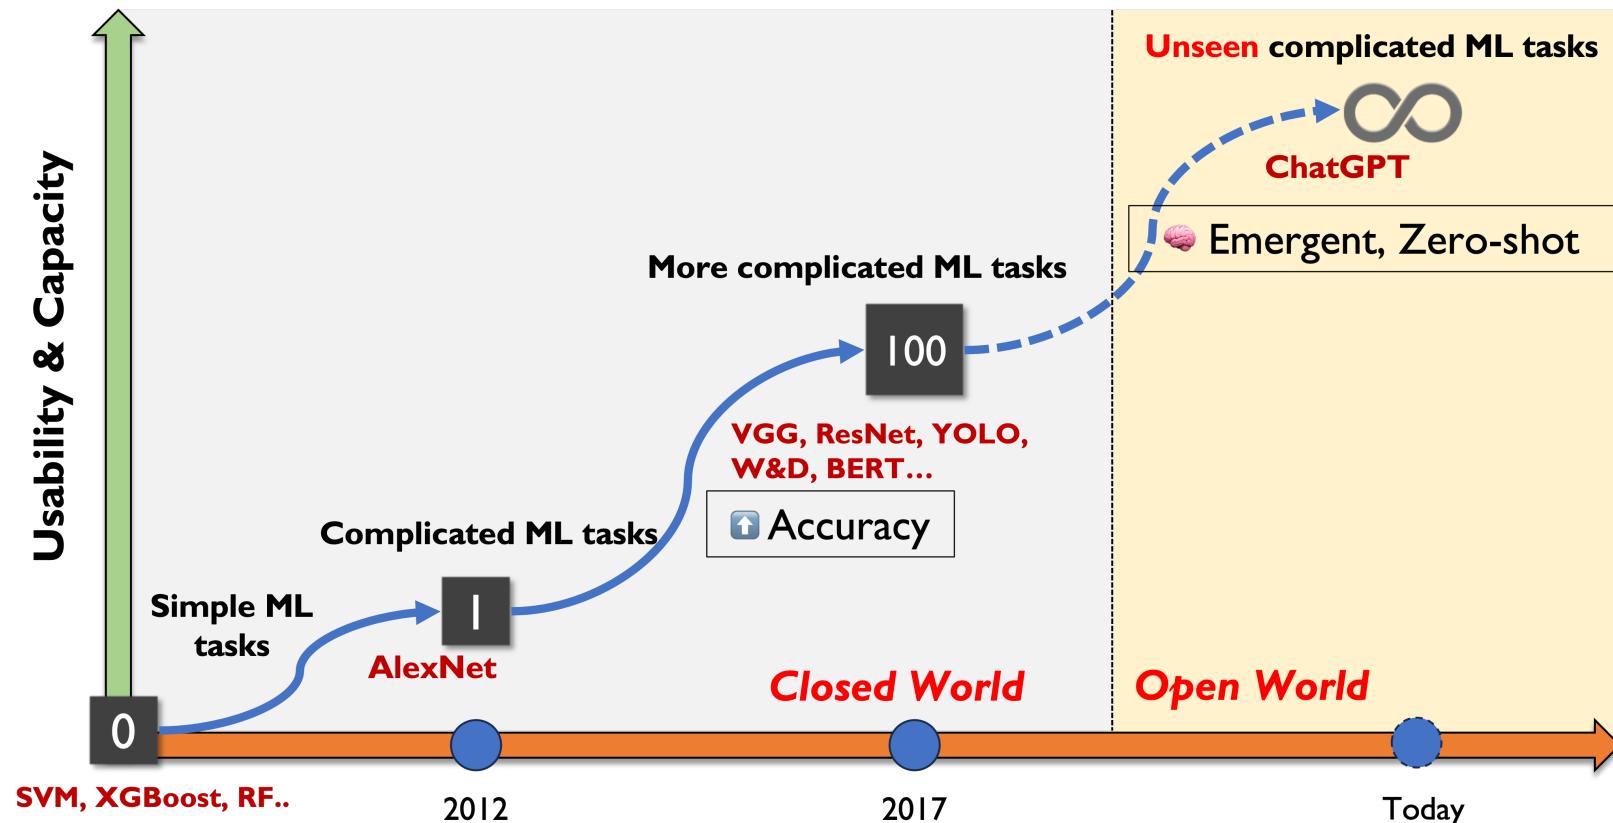
# A cool example of "smart device"



- *Comprehend physically*
- *Proactively sense, plan, and act*
- *Retrieval from Internet or Remote DB*
- *Predict the future (multimodal)*
- *Instruction following*
- *Fast response*

# The opportunity: LLM

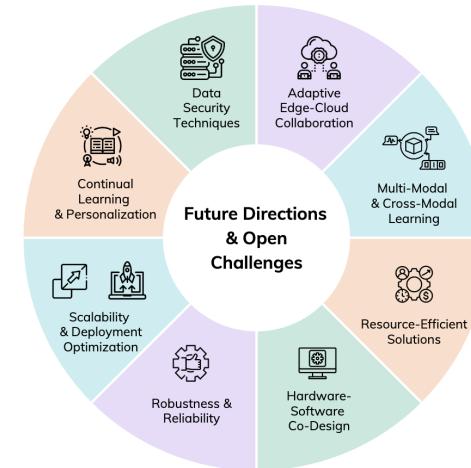
- To bring mobile devices the “next-level” intelligence



- Comprehend human language
- Zero-shot & in-context learning
- Multimodal alignment and input/output
- Reasoning & Planning
- Long context

# On-device LLM is crucial

- On-device LLMs handle language tasks in a way that is ..
  - ✓ **cost-efficient** (important, obviously)
  - ✓ **more available** (even w/o network)
  - ✓ **faster** (not always)
  - ✓ **privacy-preserving** (very important, LLMs can leverage almost every bits of local data)
- LLMs on devices does not obviate mega-scale LLMs on clouds!
  - Creating music/poetry, solving math problems, etc.



[1] Jiajun Xu, et al. "On-Device Language Models: A Comprehensive Review". In preprint'24.

# On-device LLM is crucial

- We already have a mobile device that can function with high intelligence!



*A mobile device that can comprehend, reason, and plan **without a cloud!***

# So, what's unique to mobile LLM?

(compared to traditional DNN-powered apps)

<b>Workload:</b>	fragmented tasks	→	a unified agent
<b>OS:</b>	model-agnostic	→	LLM-native
<b>Hardware:</b>	heterogeneous H/W	→	DSA-dominated

# So, what's unique to mobile LLM?

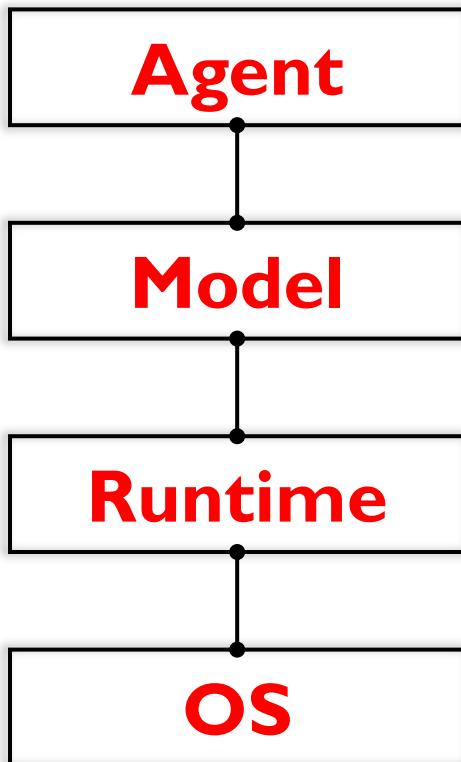
(compared to traditional DNN-powered apps)

<b>Workload:</b>	fragmented tasks	→	a unified agent
<b>OS:</b>	model-agnostic	→	LLM-native
<b>Hardware:</b>	heterogeneous H/W	→	DSA-dominated

*The source of research/industrial opportunities*

# Call for full-stack design

- Our response: **agent-model-runtime-OS** co-design



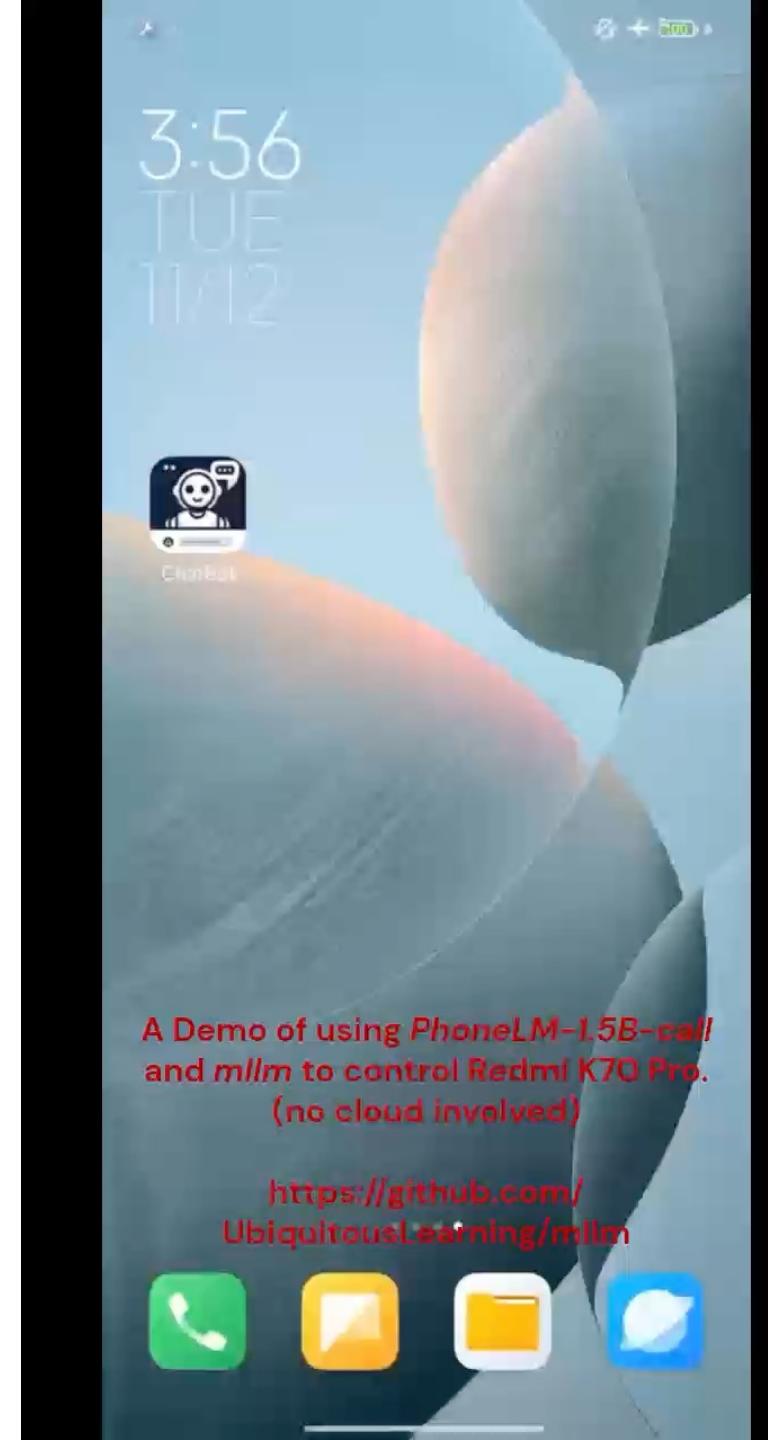
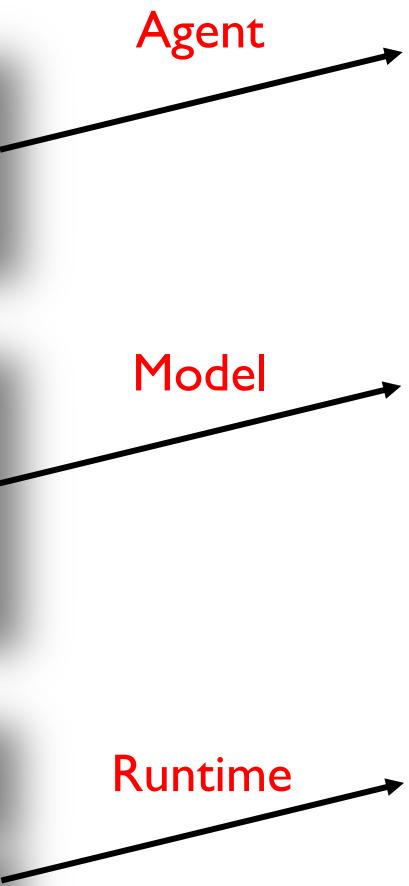
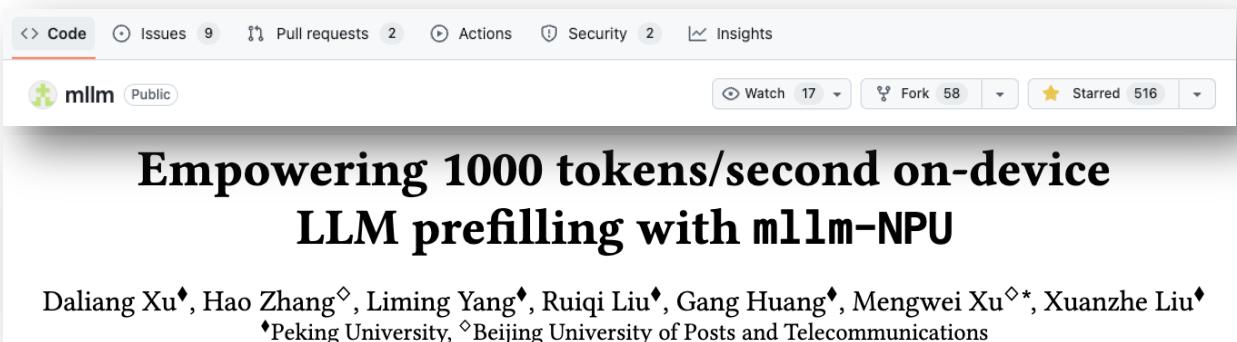
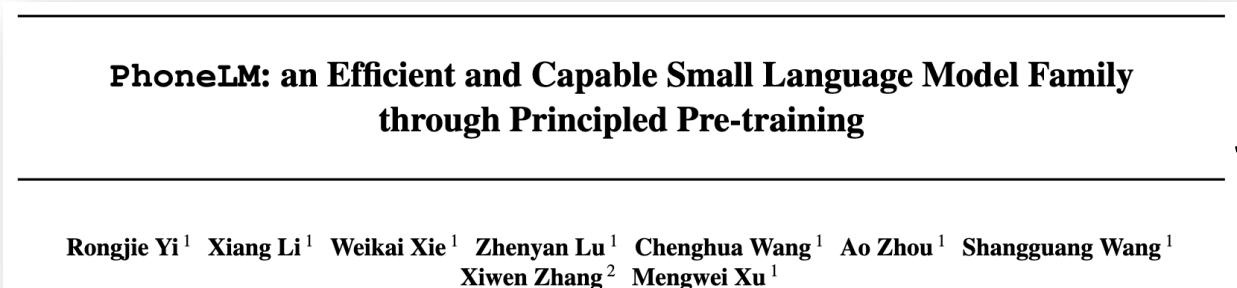
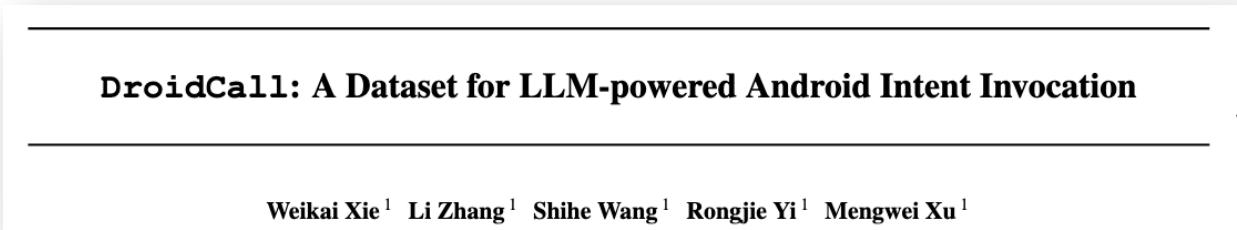
Device control and GUI agents **testbed** [LlamaTouch, UIST'24], **datasets** [DroidCall, preprint'24][SHORTCUTSBENCH, ICLR'25], and **privacy enhancements** [SILENCE, NeurIPS'24]

A training-from-scratch, fully-reproducible **SLM family** [PhoneLM, preprint'24], Any-to-any modality **mobile foundation model** [M4, MobiCom'24] , and **Federated LLM** techniques [FwdLLM, ATC'24][AdaFL, MobiCom'23] [FeS, MobiCom'23]

Acceleration through **NPU** [llm.npu, ASPLOS'25], **SpecDecoding** [LLMCad, TMC'24] , **Sparsity** [EdgeMoE, TMC'25], **Early Exiting** [Recall, preprint'24], etc

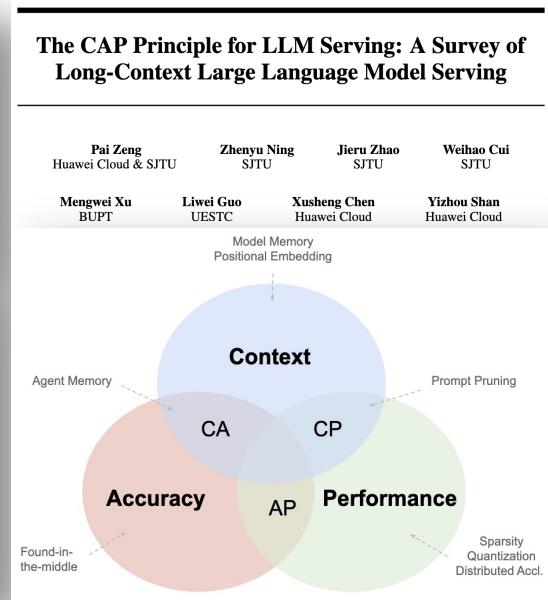
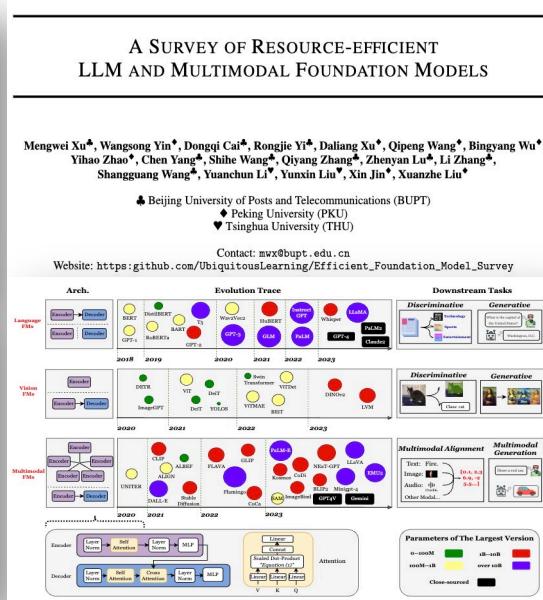
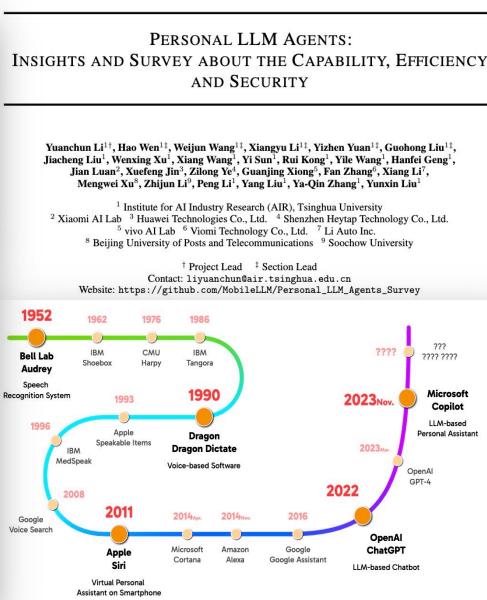
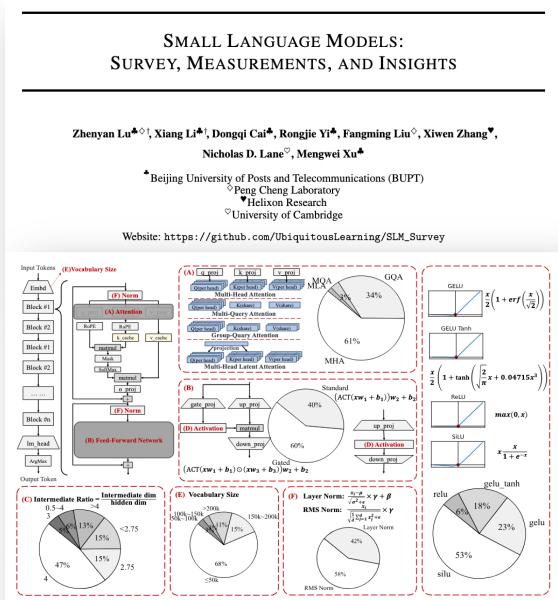
**LLMaaS Context Management** [LLMS, preprint'24] and **QoS** [ELMS, preprint'24]

# An e2e demo



# Call for full-stack design

- Our response: **agent-model-runtime-OS** co-design



- [1] "Small Language Models: Survey, Measurements, and Insights", Zhenyan Lu, et al.
- [2] "Personal LLM Agents: Insights and Survey about the Capability, Efficiency and Security", Yuanchun Li, et al.
- [3] "A Survey of Resource-efficient LLM and Multimodal Foundation Models", Mengwei Xu, et al.
- [4] "The CAP Principle for LLM Serving: A Survey of Long-Context Large Language Model Serving", Pai Zeng, et al.

# So, what's unique to mobile LLM?

(compared to traditional DNN-powered apps)

**Workload:** fragmented tasks → a unified agent

OS: model-agnostic → LLM-native

Hardware: heterogeneous H/W → DSA-dominated

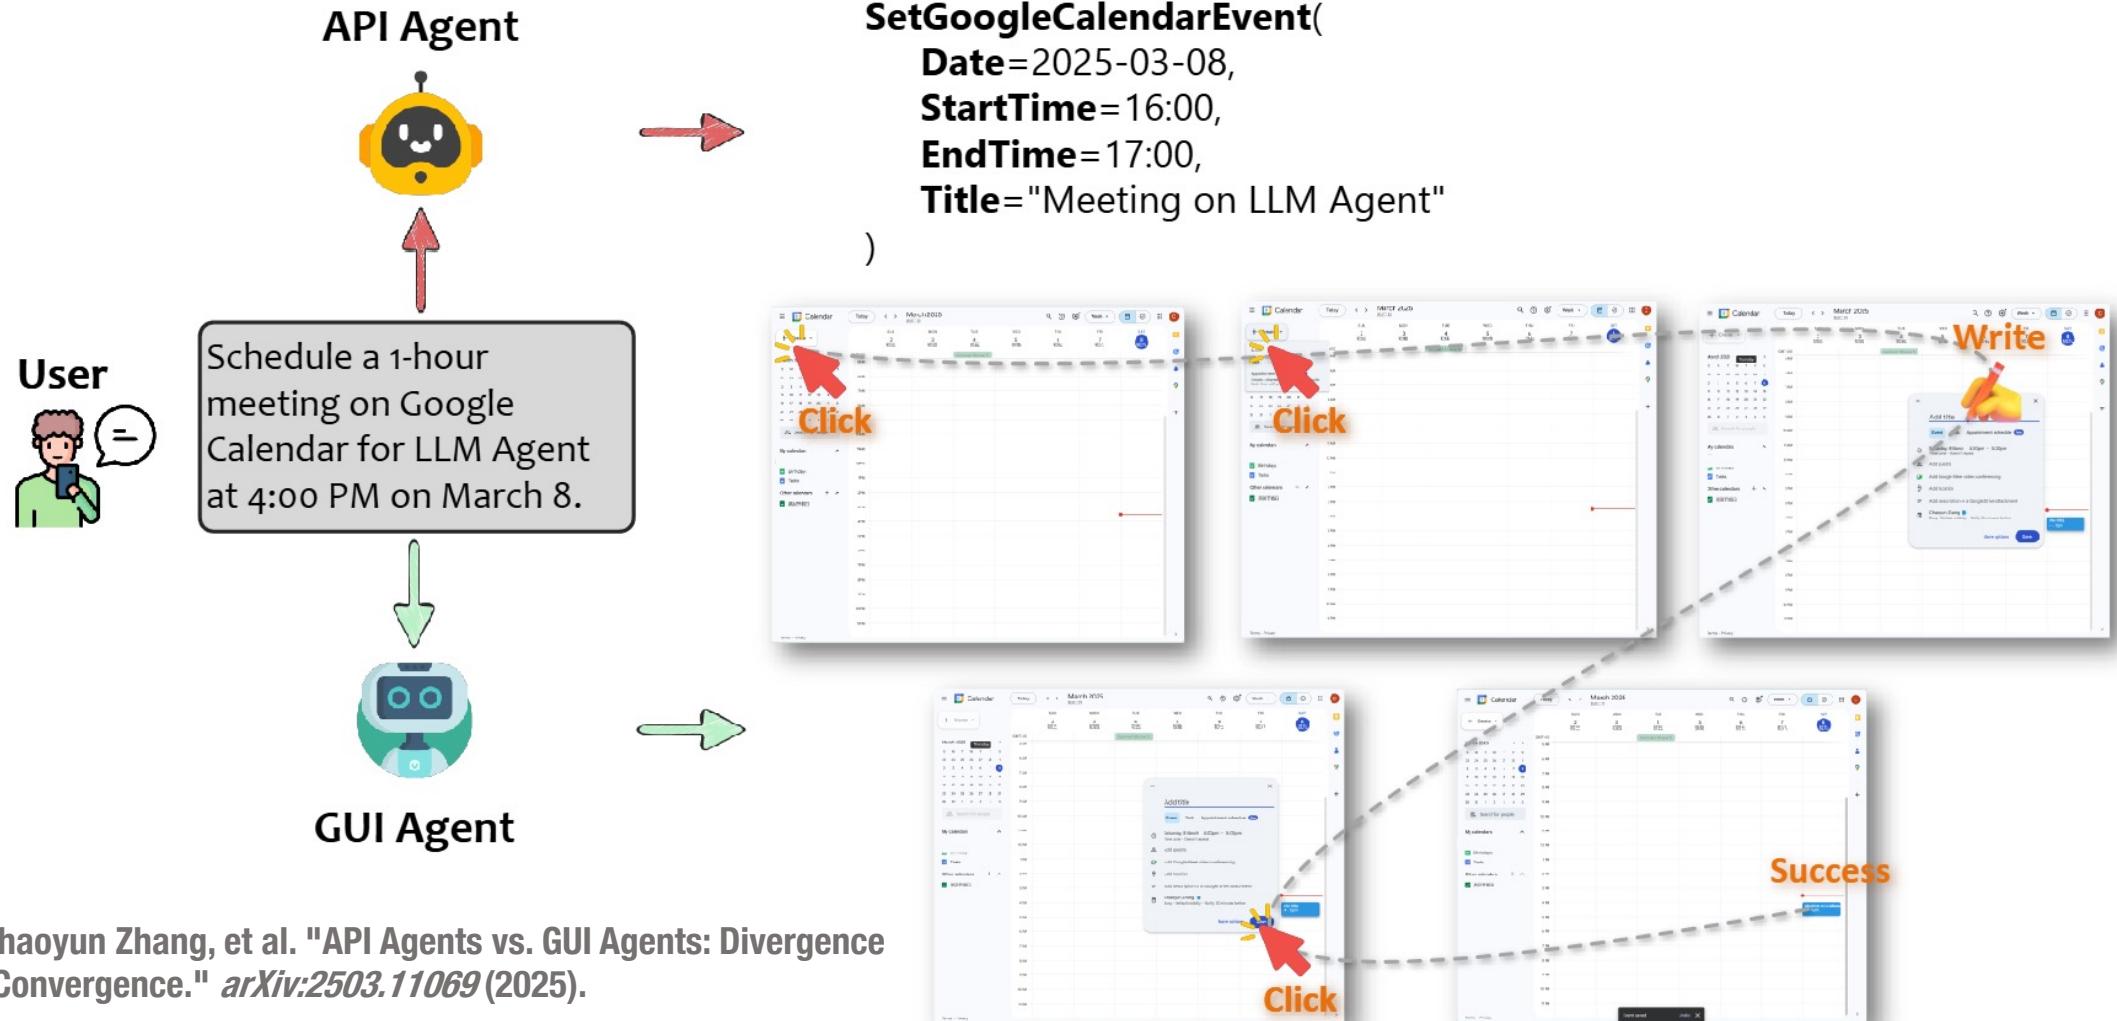
*How to build a capable, generalized, and personalized mobile agent?*

# Our vision of an agent

Making electronic devices (smartphones, robots, cars, IoTs, satellites) more accessible to **anyone** (those with cognitive difficulties) at **anytime** (when driving)

- *Comprehend physically*
- *Proactively sense, plan, and action*
- *Retrieval from Internet or Remote DB*
- *Predict the future (multimodal)*
- *Instruction following*
- *Fast response*

# General approaches: API (MCP) vs. GUI



[1] Chaoyun Zhang, et al. "API Agents vs. GUI Agents: Divergence and Convergence." *arXiv:2503.11069* (2025).

# GUI Agent: Status Quo

Offline

SFT/RL/In-context Prompting



Instruction:  
Display  
calendar in  
week view

Response:  
(0.30, 0.24)

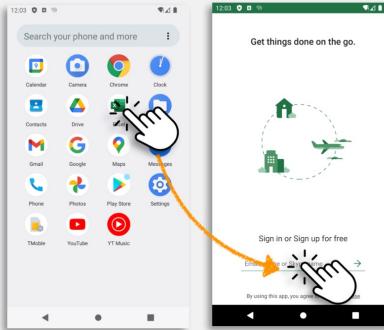
Grounding  
Training



Question:  
Who created  
the album, and  
when?

Response:  
Green Day in  
2019

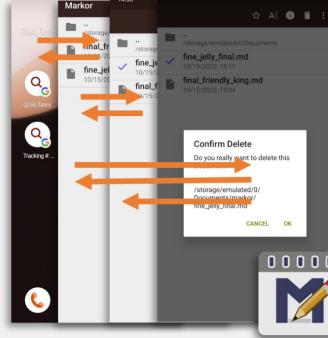
Screen Q&A



Instruction:  
Open Excel and  
Login

Response:  
Click (0.38,  
0.7)

GUI Navigation



App Knowledge

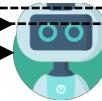
VLMs

- GPT-4o
- Claude 3.7 Sonnet
- Qwen-VL

GUI Datasets

- Rico
- AndroidInTheWild
- MobileViews

Online



GUI Agent

Agentic Workflow

Input

Task  
Instruction

Pixel-based  
Screenshot



GUI  
Represent  
ation

Working  
Memory

Text-based  
View Hierarchy

```
<?xml version='1.0'  
encoding='UTF-8'?>  
<node index=0  
class=FrameLayout>  
  <node index=0  
  class=TextView  
  text='21st Country'  
  />  
  <node index=1  
  class=TextView  
  text='Album 2009'  
  />  
  </node>  
<node index=1  
  class=FrameLayout>  
</node>
```

Set-of-Mark  
Prompting



Tool Using



Multi-agent  
Collaboration

Reflection and  
Backtracking



Test-time Scaling



Environment Perception

Interactive Testbed  
(Smartphone, Android Emulator)



- AndroidWorld
- LlamaTouch

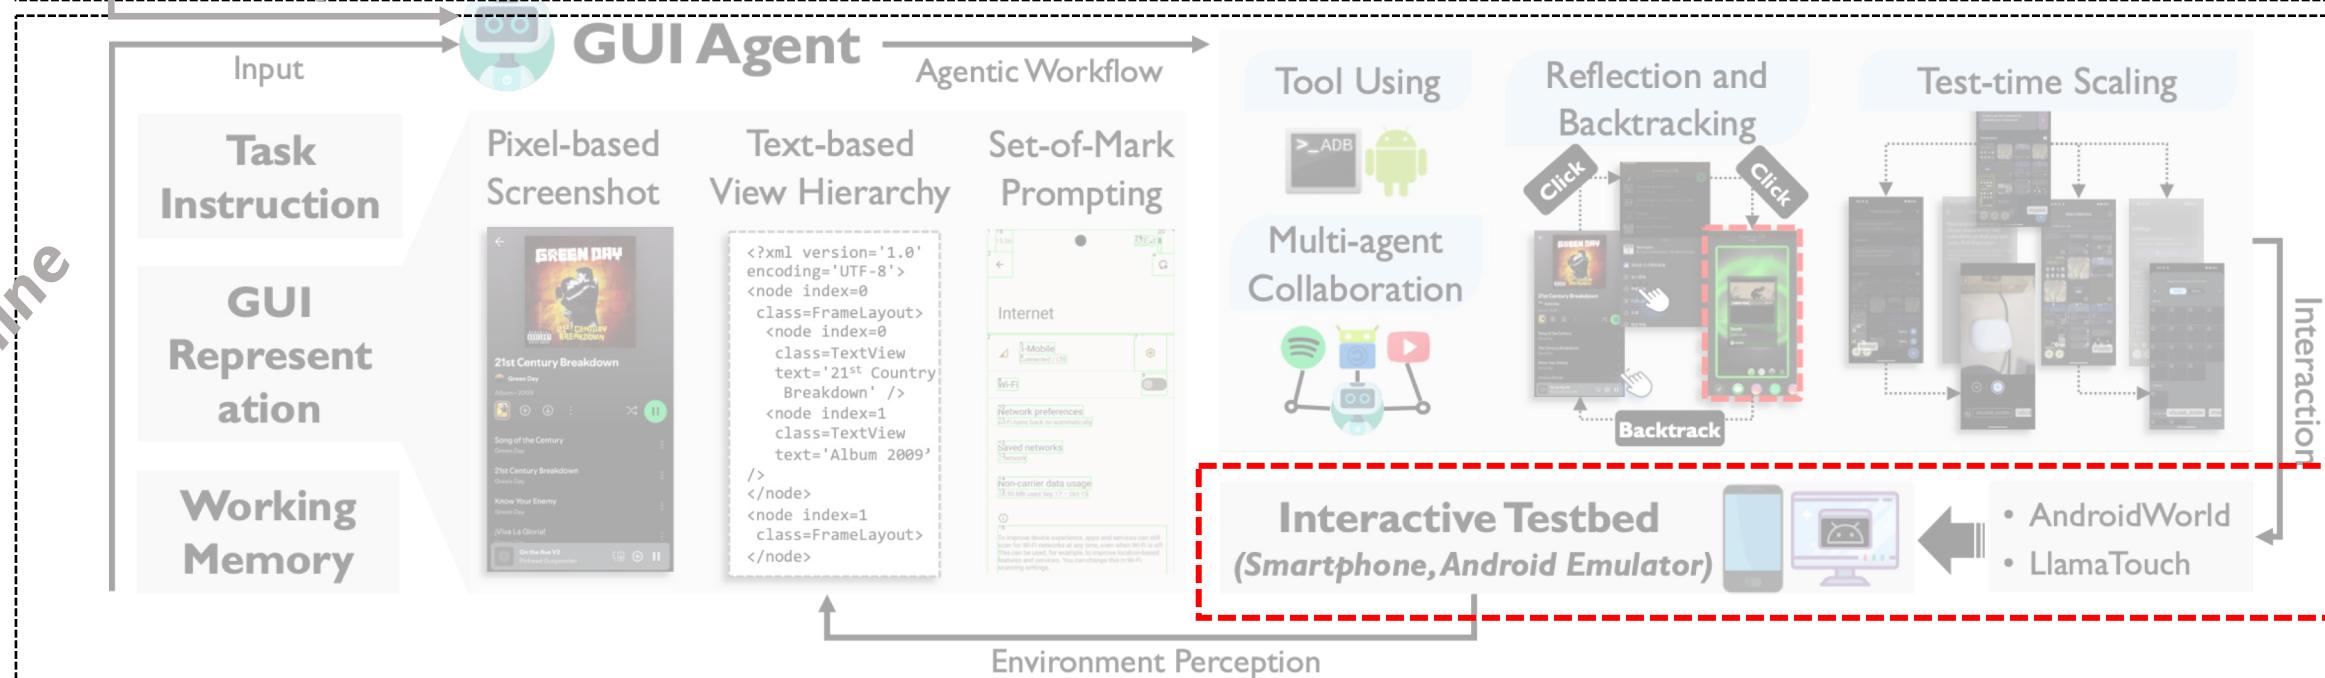
Interaction

# GUI Agent: Status Quo

Offline



Online

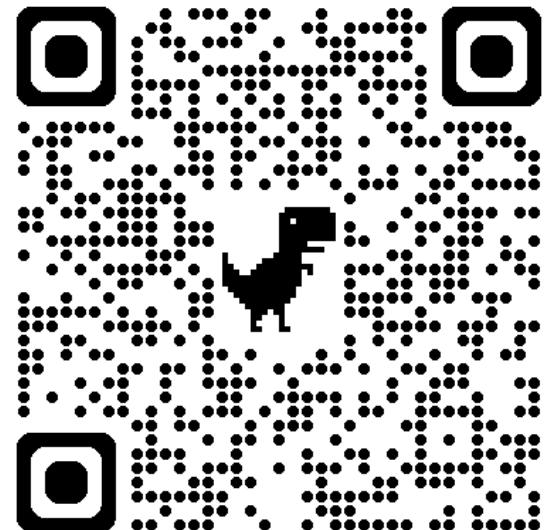


# *Collecting mobile GUI datasets, with good quality and quantity*

[arxiv'24] MobileViews:A Large-Scale Mobile GUI Dataset

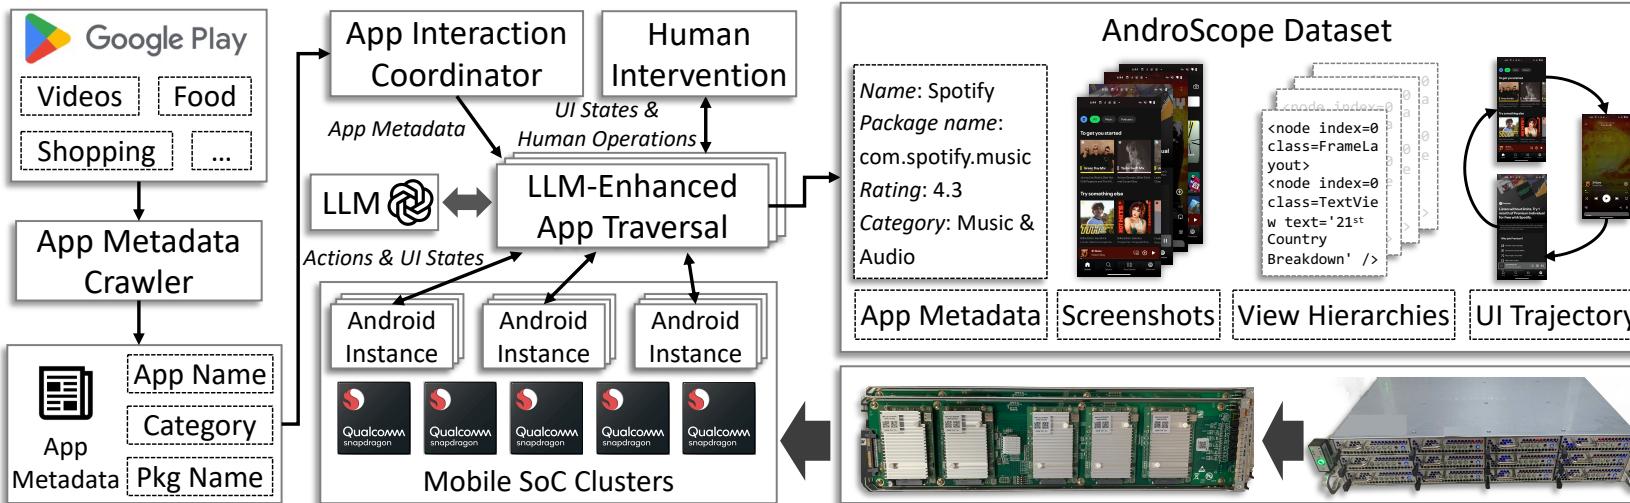
Data at

<https://huggingface.co/datasets/mllmTeam/MobileViews>



# MobileView: largest (>1M) open mobile GUI dataset

- LLM-enhanced app traversal + SoC Clusters



- Handling log-in actions through LLM
- 2x 2U SoC clusters, 120 Snapdragon 865 in total, further virtualized

- Largest coverage
  - >1M UIs
  - >20K apps
- Bilingual apps
- Both UI and VH
- Action traces

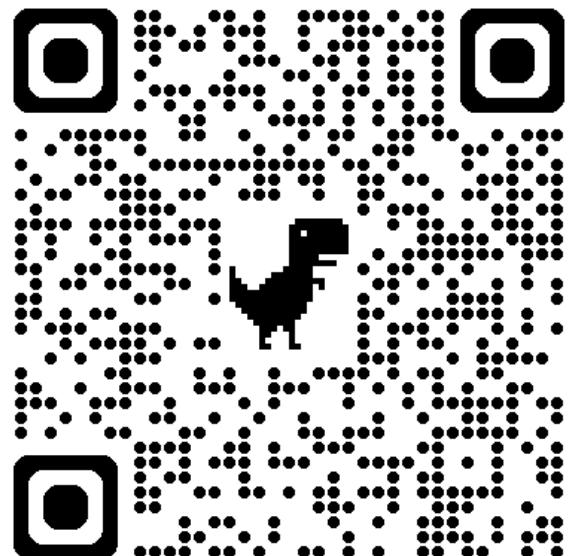
Mobile Screen Dataset	Scale		Content Diversity					Data Collection	
	Apps	Unique Screens	App Metadata	Screenshot-VH Pair	UI Trajectory	Chinese UI	Different Resolution	Automation	Hardware
Rico (Deka et al., 2017)	9,772	63,370	✓	✓	✓	✗	✗	✗	Physical Devices
PixelHelp (Li et al., 2020a)	4	187	✗	✓	✓	✗	✗	✗	Emulators
Screen2words (Wang et al., 2021)	6,269	22,417	✓	✓	✗	✗	✗	✗	N/A
ScreenQA (Baechler et al., 2024)	N/A	35,352	✓	✓	✗	✗	✗	✗	N/A
META-GUI (Sun et al., 2022)	11	24,825	✗	✓	✓	✗	✗	✗	Physical Devices
DroidTask (Wen et al., 2024)	13	362	✗	✓	✓	✗	✗	✗	Emulators
AITW (Rawles et al., 2023)	357	2,282,533	✗	✗	✓	✓	✗	✗	Emulators
LlamaTouch (Zhang et al., 2024b)	57	3,281	✗	✓	✓	✗	✗	✗	Emulators
<b>GUIScope</b>	<b>30,037</b>	<b>1,213,866</b>	✓	✓	✓	✓	✓	✓	SoC clusters

[1] Longxi Gao, et al. "MobileViews: A Large-Scale Mobile GUI Dataset". In preprint'24.

# *Benchmarking mobile GUI agents, properly and efficiently*

[UIST'24] LlamaTouch: A Faithful and Scalable Testbed for  
Mobile UI Task Automation

Code at <https://github.com/LlamaTouch/LlamaTouch>



# LlamaTouch: mobile GUI testbed

- Existing approaches: human/LLM eval.; step-wise action match
- Our approach: **critical states** matching

Match Type	State Type	Primitive	Keyword	Use Case
Fuzzy match	UI state	Screen info	fuzzy<1>	Check if the contents on two screens are approximately identical.
		Textbox	fuzzy<n>	Check if the content of the target textbox is semantically similar to the content of the original textbox<n> in the ground-truth UI.
		Activity	activity	A coarse-grained approach to determining if two UIs represent the same functional screen in an application.
Exact match	UI component	UI component	exact<n>, exclude<n>	Check if the UI component is exactly identical to the UI component<n>, or does not occur, in the ground-truth UI.
		System state	(Un)installation	installed<app>, uninstalled<app>
	Action	Action	click<n>, type<input_text>	Check if the target application named ‘app’ has been successfully installed/uninstalled.
				Check if two actions and their parameters are identical.

**Two types of matching primitives:**  
**Exact match and fuzzy match**

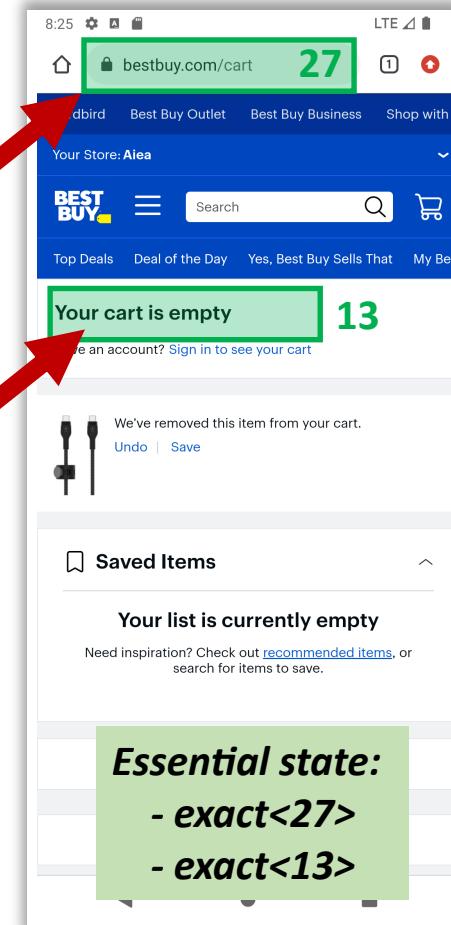
Table 7: Accuracy (Acc. %) of different evaluation approaches among all successful tasks in human validation.

Mobile Agent	Step-wise action match	LCS action match	LlamaTouch	Human
	Acc.	Acc.	Acc.	# succ.
AutoUI	0.00	0.00	77.78	9
AutoDroid	0.00	0.00	73.91	69
AppAgent	0.00	3.03	93.94	33
CoCo-Agent	0.00	0.00	70.00	10
Average	0.00	0.76	78.91	30

**High eval. acc**

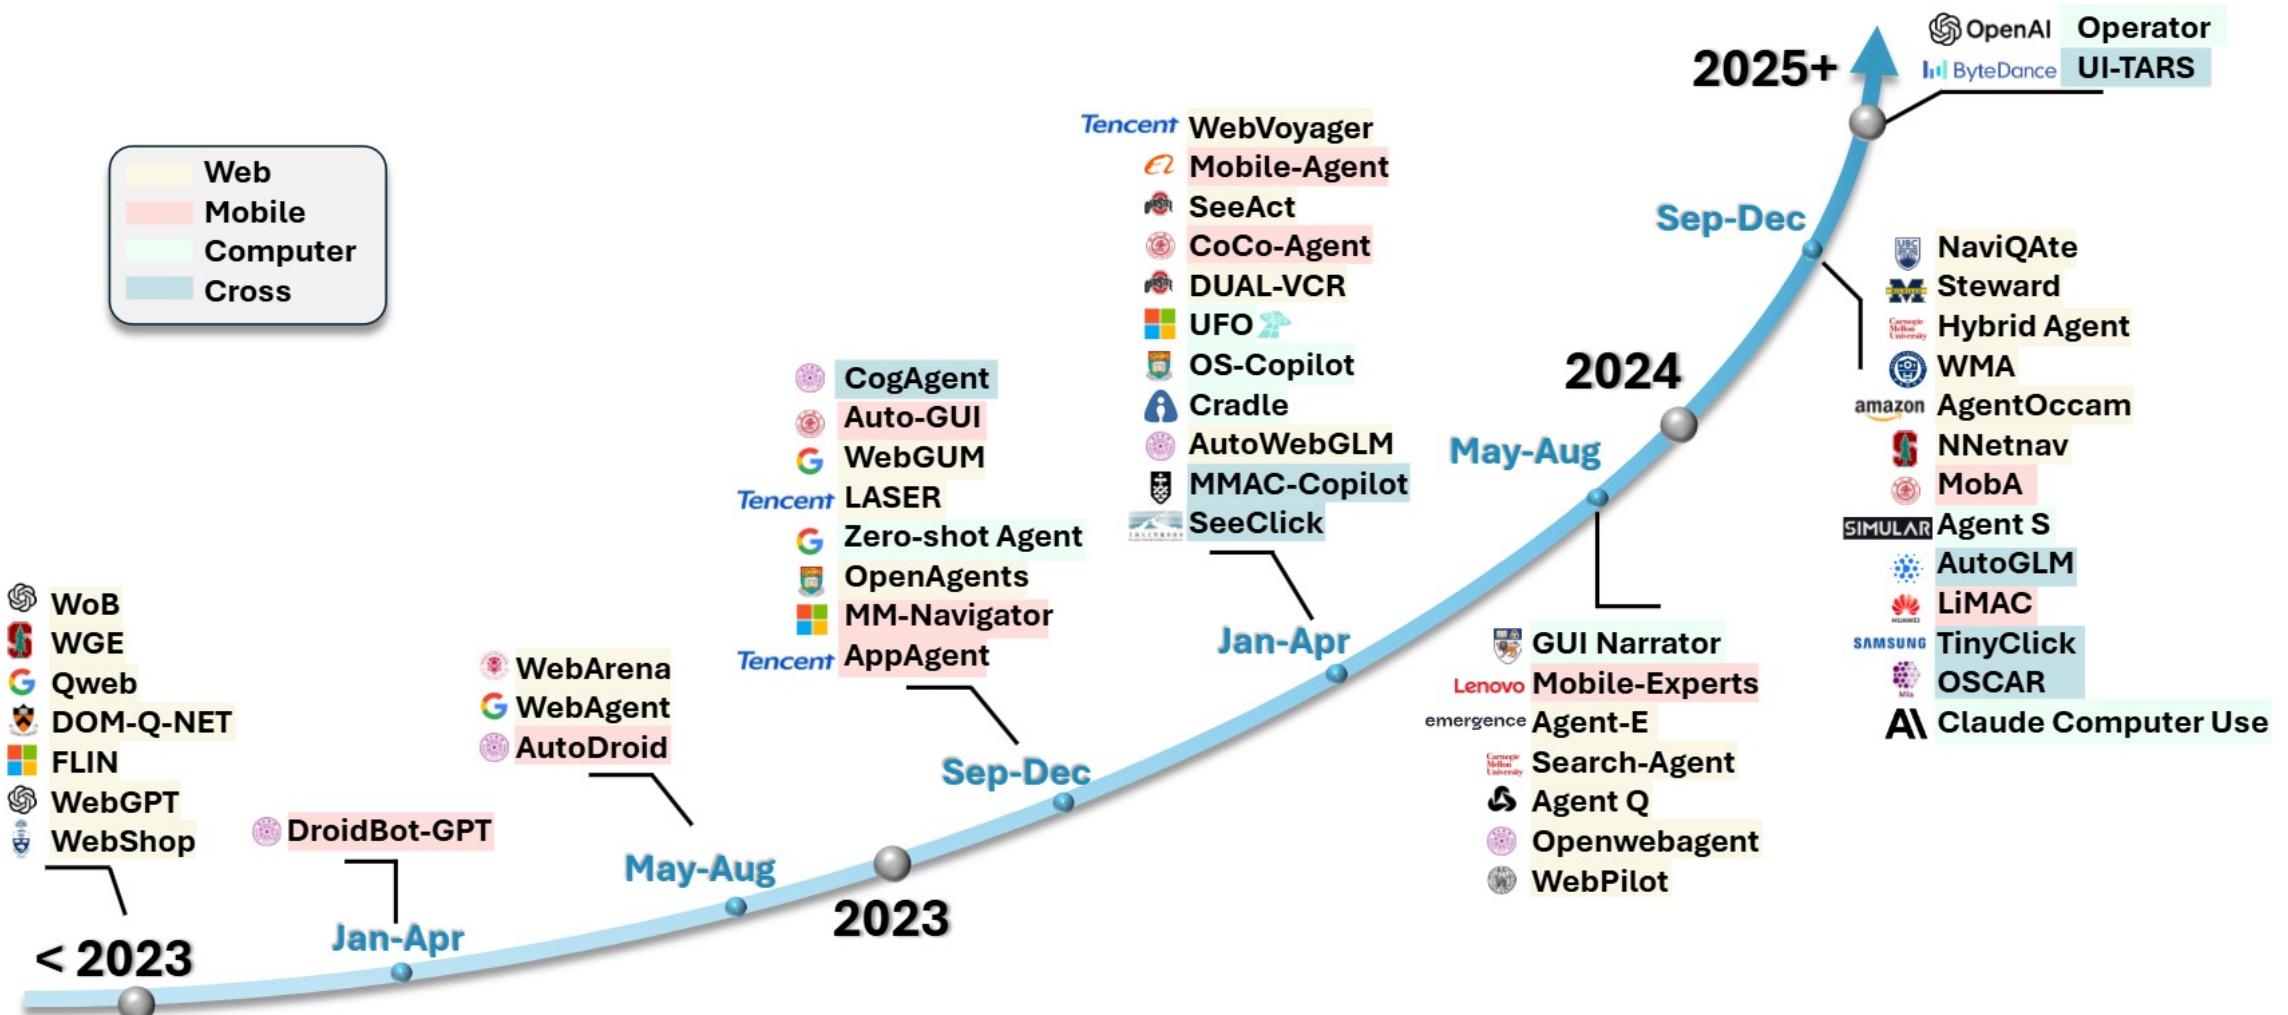
**Task: Empty my shopping cart on BestBuy**

- ❑ **Exact match** on the URL:  
“bestbuy.com/cart”
- ❑ **Exact match** on the UI text  
“Your cart is empty”
- ❑ All others (actions, UI components) are omitted.



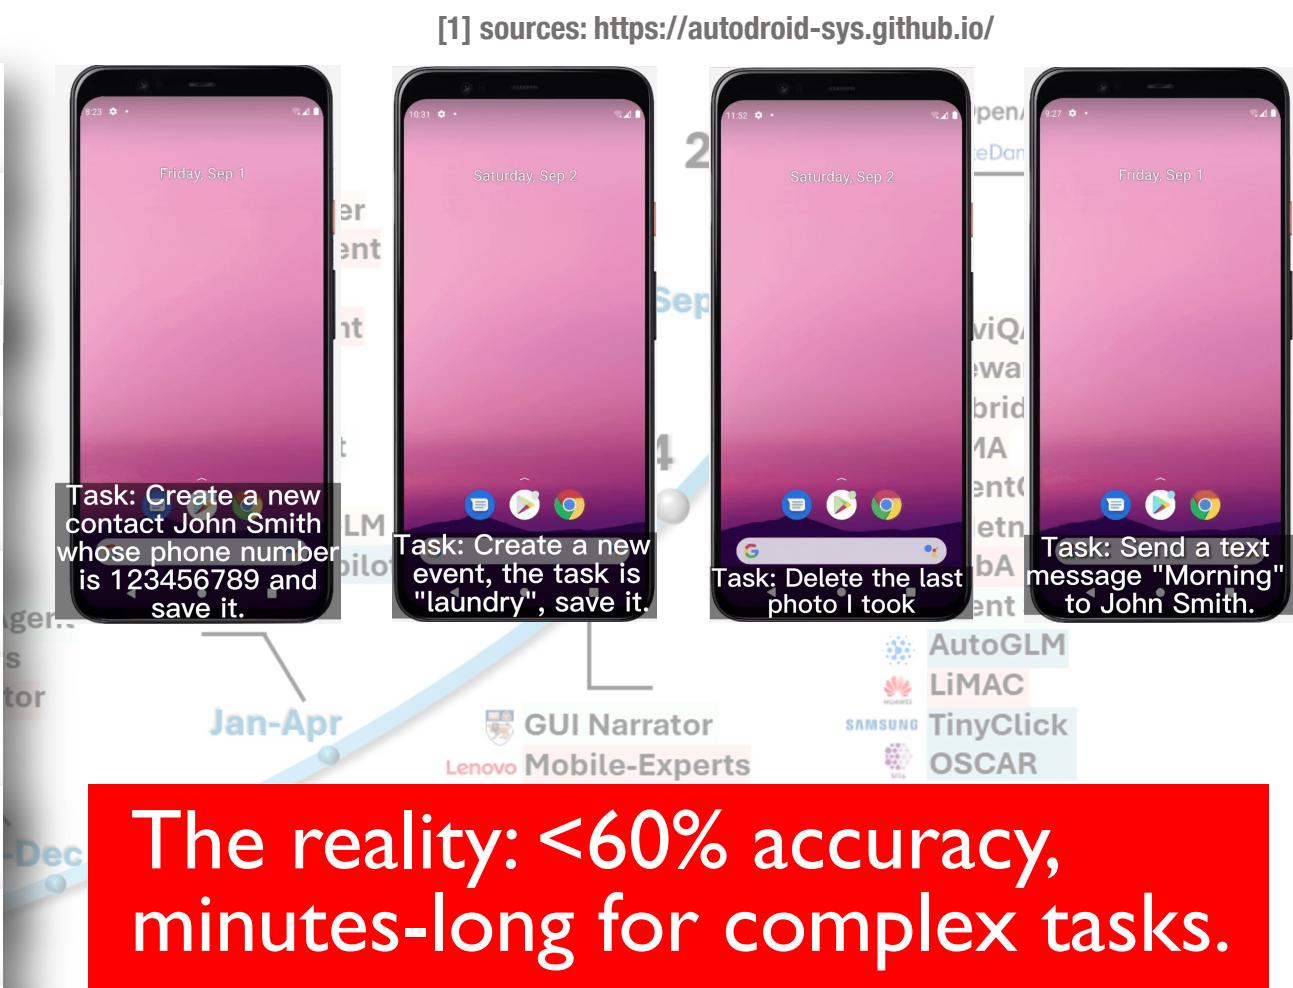
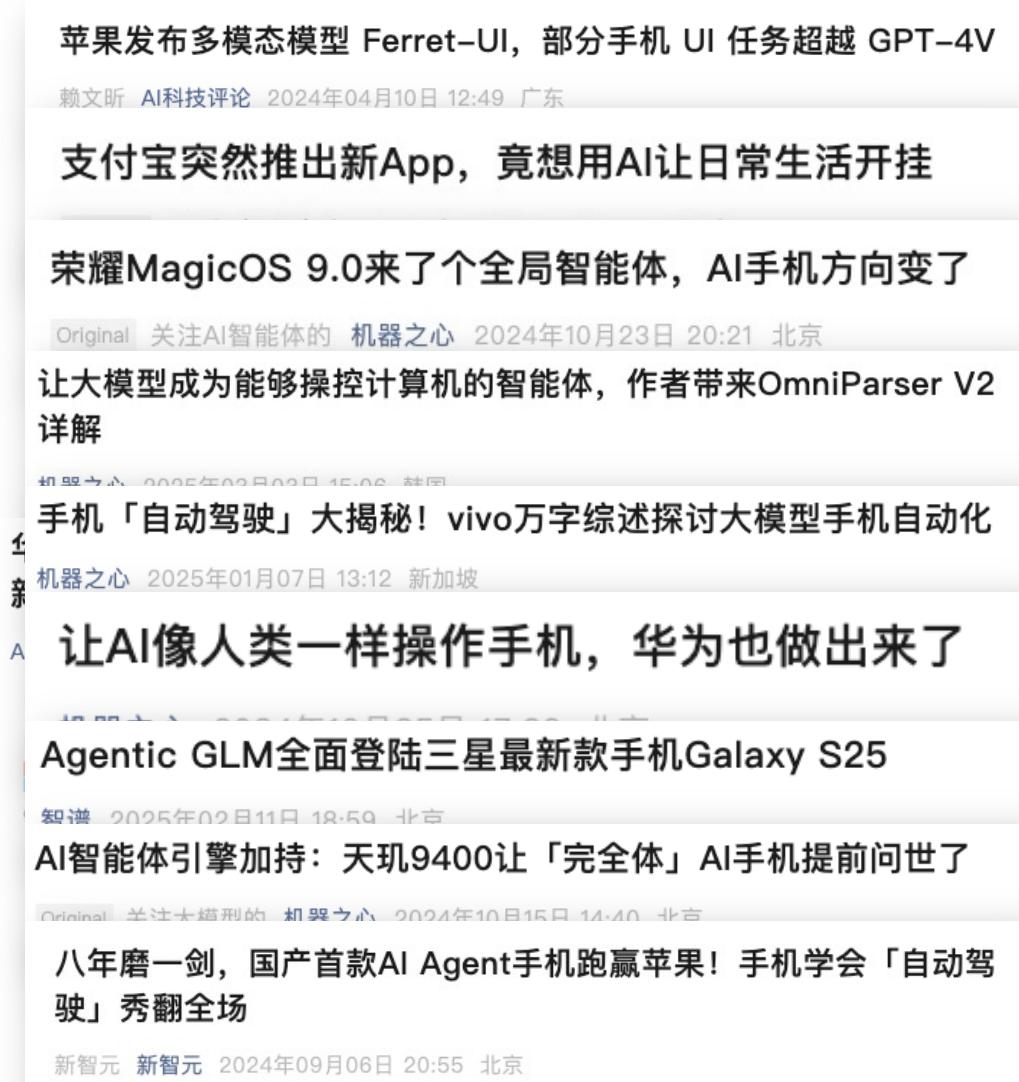
[1] Li Zhang, et al. “LlamaTouch: A Faithful and Scalable Testbed for Mobile UI Task Automation”. In UIST’24.

# GUI Agent: Status Quo



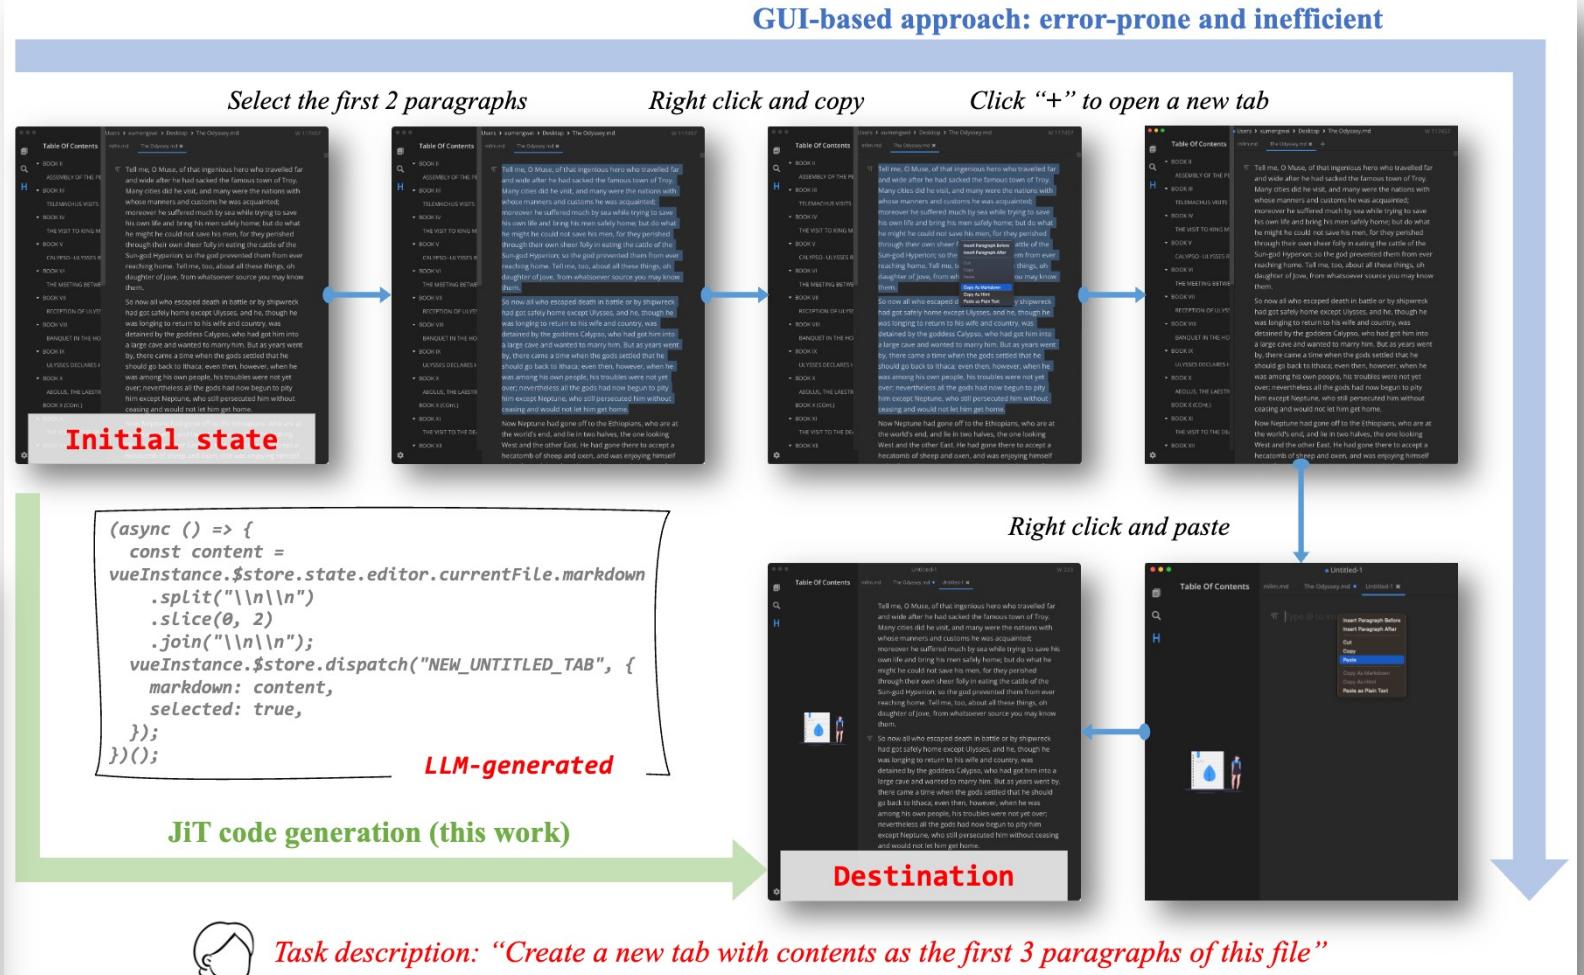
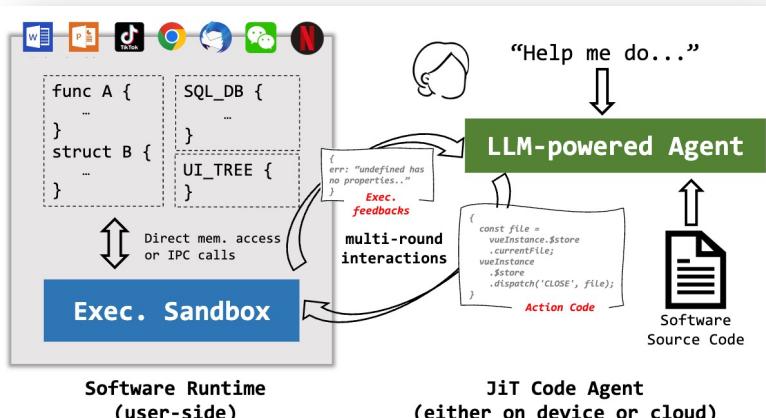
[1] Chaoyun Zhang, et al. "Large Language Model-Brained GUI Agents:A Survey." *arXiv:2411.18279* (2024).

# GUI Agent: Status Quo



# A Third Path? Codegen is all you need

Just-in-time code generation and in-app execution



[1] Mengwei Xu. "Every Software as an Agent: Blueprint and Case Study" arXiv:2502.04747(2025).

# So, what's unique to mobile LLM?

(compared to traditional DNN-powered apps)

**Workload:** fragmented tasks → a unified agent

**OS:** model-agnostic → **LLM-native**

**Hardware:** heterogeneous H/W → DSA-dominated

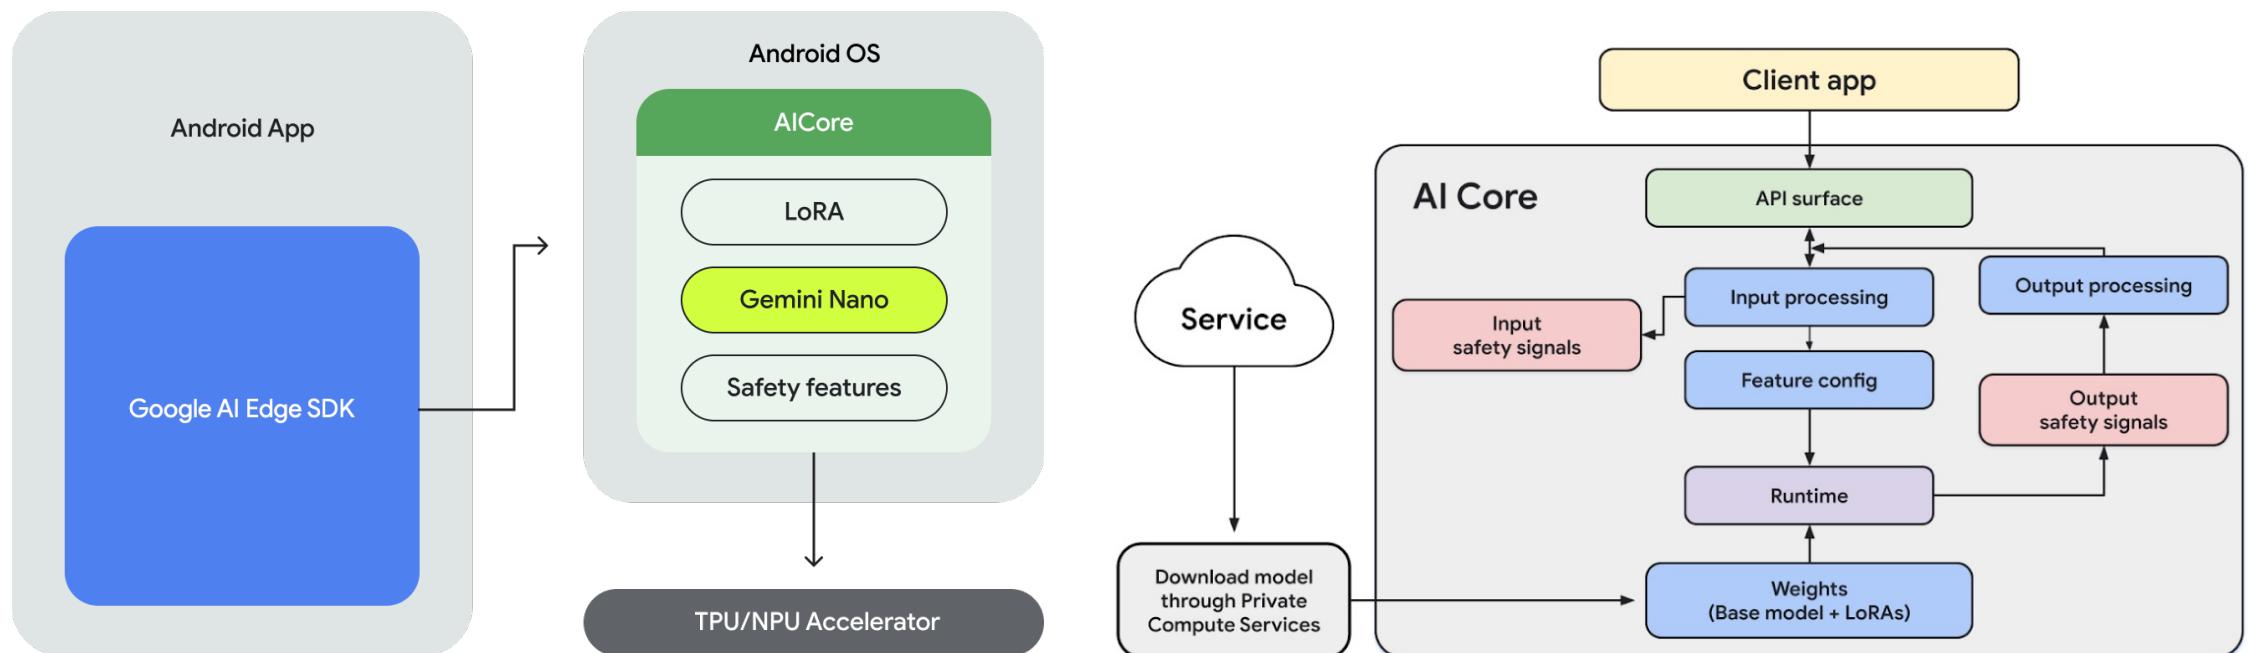
*How should OS better serve/manage device-wise LLM requests?*

# LLM as a system service

- LLM integrated into OS as a **system service** (LLMaaS)
  - Scales to infinite number of tasks
  - Hardware-design-friendly
  - OS gains full visibility into LLM requests
- Opening new research opportunities and challenges
  - *Efficiency*: how to schedule, batch, and cache-reuse system-wise LLM requests?  
How to manage the LLM context states across apps?
  - *Security*: how to protect app-owned LoRa? How to isolate cross-app requests?
  - *Usability*: how to upgrade LLM? How to design LLMaaS interface?
  - Etc..

# A pioneering case: Android

- Android introduced an LLM system service (a.k.a. AICore)
  - Since end of 2024, but still in preview



[1] Source: <https://developer.android.com/ai/gemini-nano>

# *Towards a capable and generalizable LLMaaS*

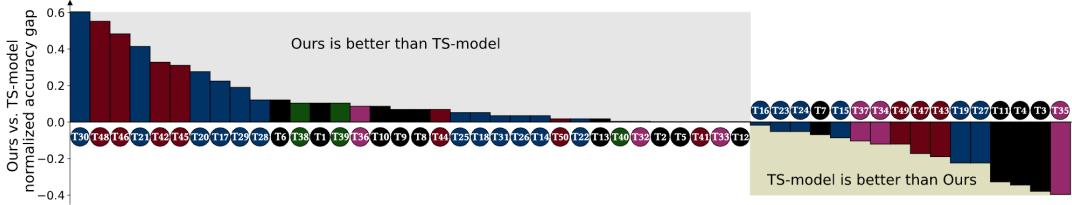
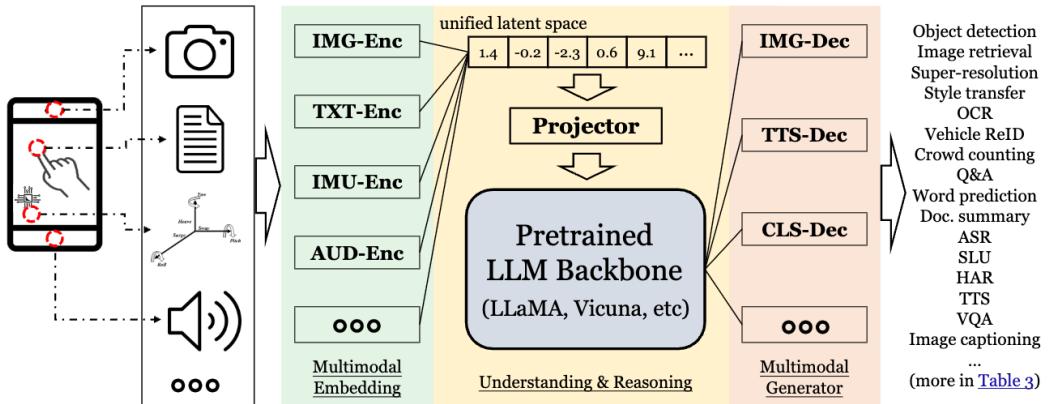
[MobiCom'24] Mobile Foundation Model as Firmware

Code at <https://github.com/UbiqitousLearning/MobileFM>



# M4: a one-size-fits-all mobile MLLM

- Can one model (as an OS service) solve all mobile AI tasks?



**M4 outperforms prior arts on most tasks**

Category	Tasks	Mobile Application	Dataset	Specific-Models	Results	Metrics
NLP	Input word prediction $T_1$	Input method (GBoard)	PTB	RNN [23]	0.17*	Accuracy
	Question answering $T_2$ , $T_3$	Private assistant (Siri)	SQuAD v2.0	RoBERTa [35]	0.79*	F1
	Machine translation $T_4$	Translator (Google Translate)	wmt22-en-de	AraELECTRA [36], Transformer [32]	0.87	F1
	Emoji prediction $T_5$	Input method (GBoard)	tweet_eval	RoBERTa [22]	0.34*	BLEU
	Emotion prediction $T_6$	Conversational analytics (Clarabridge)	go_emotion	RoBERTa [29]	0.57*	Accuracy
	Sentiment analysis $T_7$	Conversational analytics (Clarabridge)	tweet_eval	RoBERTa [27]	0.77*	Accuracy
	Text classification $T_8$ , $T_9$	Spam SMS filtering (Trucaller)	ag_news	BERT [37]	0.93*	Accuracy
	Grammatical error correction $T_{10}$	Writing assistant (Grammarly)	SST2	DistilBERT [38]	0.91*	Accuracy
	Text summary $T_{11}$	Reading assistant (ChatPDF)	JFLEG	FLAN-T [30]	0.68*	BLEU
	Code document generation $T_{12}$	Code editor (Javadoc)	CNN Daily Mail	BART [5]	0.43*	ROUGE1
	Code generation $T_{13}$	Code editor (Copilot)	CodeSearchNet	CodeT5-base [20]	0.33*	ROUGE1
		Shellcode_IAS2	CodeBERT [21]	0.92	BLEU	
	Object detection $T_{14}$ , $T_{15}$	Augmented Reality (Google Lens)	COCO	Libra-rcnn [24]	0.43*	mAP
	Image retrieval $T_{16}$	Image searcher (Google Photos)	LVIS	X-Paste [33]	0.51	AP
	Super-resolution $T_{17}$	Video/Image super-resolution (VSCO)	Resnet50-arcface [31]	Real-ESRGAN [19]	0.90*	Recall
CV	Style transfer $T_{18}$	Painting & Beautifying (Meitu)	COCO, Wikiart	StyleGAN-nada [4]	0.23	CLIP score
	Semantic segmentation $T_{19}$ , $T_{20}$	Smart camera (Segmentix)	ADE20K-150	DeepLabv3plus [25]	0.43*	mIoU
		PASCAL VOC	DeepLabv3plus [26]	0.79*	mIoU	
	Optical character recognition $T_{21}$	Intelligent document automation	Rendered SST2	CTIP [34]	0.71	Accuracy
	Image captioning ...					
	ASR					
	SLU					
	HAR					
	TTS					
	VQA					
Im	...					
	(more in Table 3)					
	Tr					
	Ve					
	Gender recognition $T_{22}$	Smart camera (Face++)	Adience	MiVOLO-D1 [2]	0.96	Accuracy
	Location recognition $T_{23}$	Navigation search (Google Maps)	Country211	CLIP [34]	0.46	Accuracy
	Pose estimation $T_{24}$	AI fitness coach (Keep)	AP-10K	ViTPose [134]	0.69	AP
	Video classification $T_{25}$	Video player (YouTube)	kinetics400	SlowFast [28]	0.79	Accuracy
	Crowd Counting $T_{26}$	Smart camera (Fitness Tracking)	UCF-QNRF	CSS-CCNN [12]	437	MAE
	Image matting $T_{27}$	Virtual backgrounds (Zoom)	RefMatte-RW100	MDETR [79]	0.06	MSE
Audio	Automatic speech recognition $T_{28}$	Private assistant (Siri)	LibriSpeech	CTC+attention [14]	3.16%*	WER
	Spoken language understanding $T_{29}$ , $T_{30}$	Private assistant (Siri)	FSC	Transformer [18]	0.37%	WER
	Emotion recognition $T_{31}$	Emoji recommendation (WeChat)	SLURP	CRDNN [3]	0.82*	Accuracy
	Audio classification $T_{32}$	Music discovery (Shazam)	ECAPA-TDNN [15]	ACDNC [1]	0.64*	Accuracy
	Keyword spotting $T_{33}$	Private assistant (Siri)	ESC-50	Speech command Cnn-trad-fpool3 [17]	0.87	Accuracy
Sensing	Human activity recognition $T_{34}$ , $T_{35}$ , $T_{36}$	AI fitness coach (Keep)	Using Smartphones	Ts-TCC [51]	0.90	Accuracy
			HHAR	LIMU-BERT [16]	0.84	Accuracy
			MotionSense	LIMU-BERT [16]	0.91	Accuracy
	Text-to-speech $T_{37}$	Voice broadcast (WeChat reading)	LJSpeech	Transformer [13]	3.26	MCD
Multimodal	Audio captioning $T_{38}$ , $T_{39}$	Hearing-impaired accessibility (Ava)	Clotho	Transformer [10]	0.52*	BLEU
	Image captioning $T_{40}$ , $T_{41}$	Visual-impaired accessibility (Supersence)	AudioSet	Transformer [10]	0.64	BLEU
	Text-to-image retrieval $T_{42}$ , $T_{43}$	Image search (Google Photos)	MSCOCO/14	LSTM [7]	0.73*	BLEU
	Text-to-image generation $T_{44}$	Art creation (Verb Art)	Flickr8k	LSTM [110]	0.58	BLEU
	Audio/Text-to-image generation $T_{45}$	Visual question answering (Answerables)	Flickr30k	NAPReg [69]	0.39	Recall
	Visual question answering $T_{46}$ , $T_{47}$		VGGSound	CLIP [34]	0.69	Recall
			Wav2clip [11]	VQA v2.0	99.89	FID
			MUTAN [41]	MUTAN [41]	0.63	Accuracy
			VizWiz	MUTAN [41]	0.52	Accuracy

Tested on 50 mobile AI tasks

[1] Jinliang Yuan, et al. "Mobile Foundation Model as Firmware". In MobiCom'24.

Mengwei Xu @ BUPT

# *Towards elastic LLMaS*

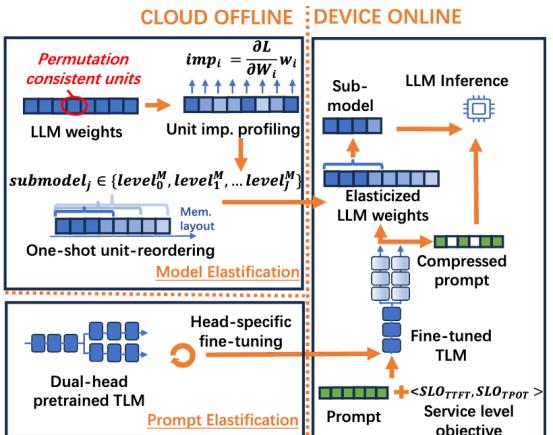
[arxiv'24] ELMS: Elasticized Large Language Models  
On Mobile Devices

# Serving LLM requests with different QoS

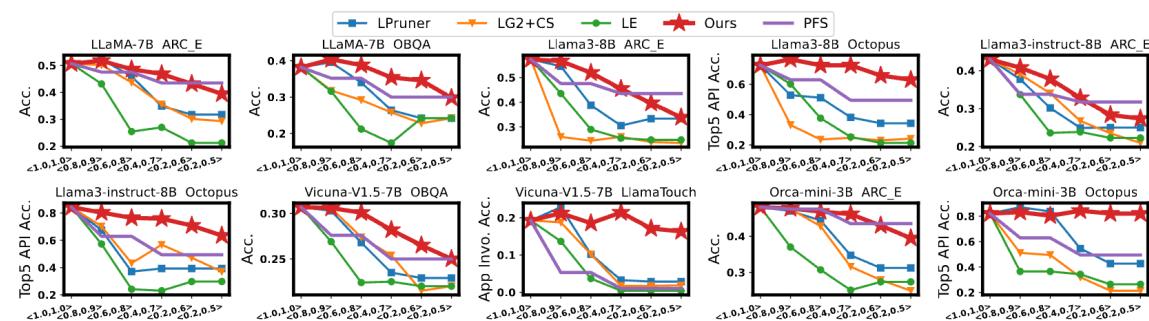
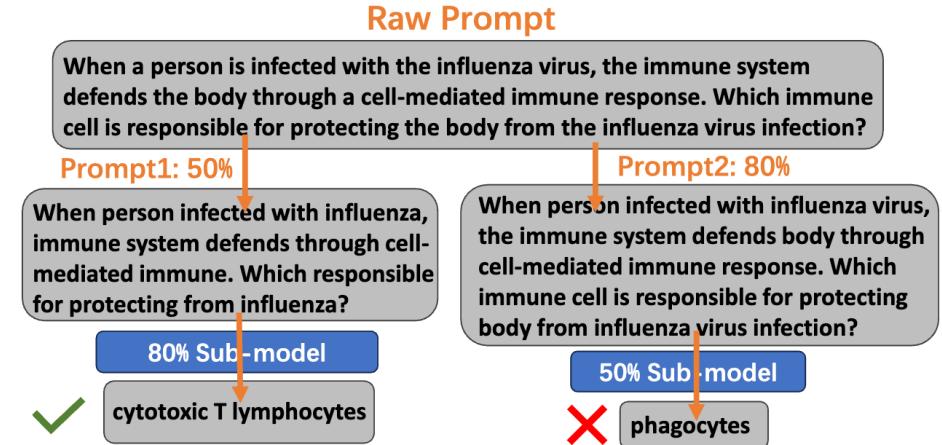
- Key idea: a joint planning of token/model pruning

Mobile LLM App.	Service-Level Objective
Chatbot [9]	Readable TTFT/TPOT
Always-on Voice Assistant [6, 16]	Very-Low TTFT, medium TPOT
Background Screen-Event Recorder [15]	Tolerable TTFT/TPOT
Smart Message Reply [5]	Low TTFT, low TPOT
API-Calling Agent [23]	Low TTFT, acceptable TPOT
UI-Automation Agent [71, 84]	Low TTFT, acceptable TPOT

Different apps demand diversified QoS



A offline-guided, joint planning of token pruning and weights pruning



Significant improvement over static approaches

# So, what's unique to mobile LLM?

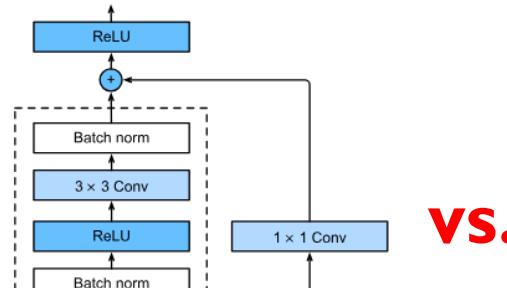
(compared to traditional DNN-powered apps)

<b>Workload:</b>	fragmented tasks	→	a unified agent
<b>OS:</b>	model-agnostic	→	LLM-native
<b>Hardware:</b>	<b>heterogeneous H/W</b>	→	<b>DSA-dominated</b>

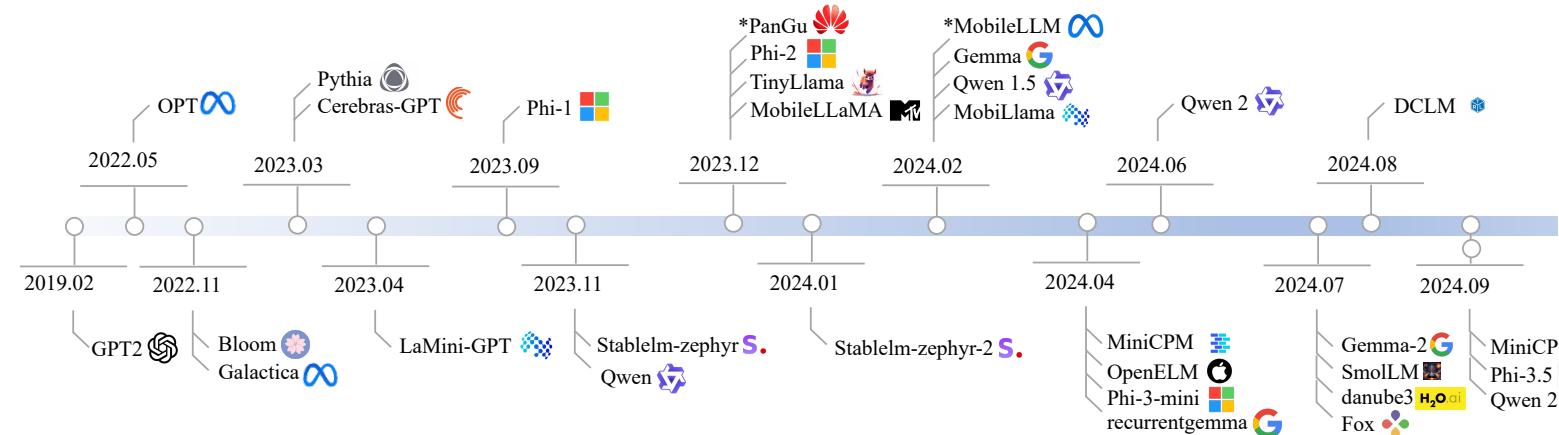
*How to serve LLM requests with low latency and energy efficiency?*

# On-device LLM needs LLM-processor

- On-device **resource scarcity** further exacerbated.



VS.



ResNet, YOLO, LSTM, etc  
(<200M)  
<100ms to process one image  
<100MB memory footprint  
Easy to quantize (integer-only)  
Static shape and cost

>10sec to process one prompt on CPU  
>1GB memory footprint  
Difficult to quantize (FP required)  
Dynamic shape and increased cost with longer prompt

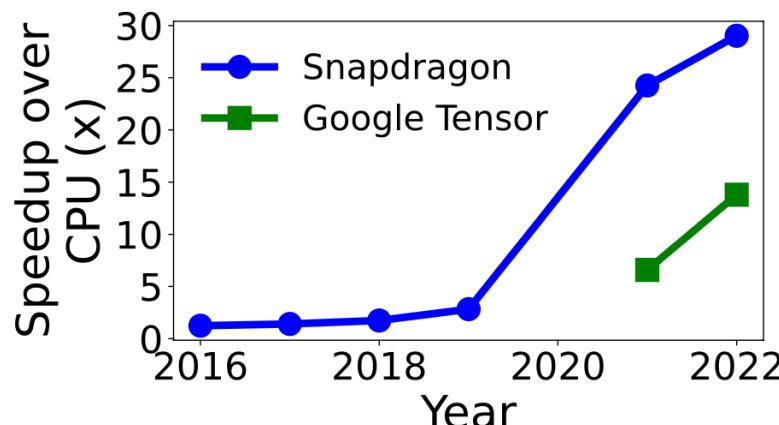
[1] Zhenyan Lu, et al. "Small Language Models: Survey, Measurements, and Insights". In preprint'24.

# On-device LLM needs NPU

- DSA (LLM-processor) is the answer to on-device LLM.

Vendor	Latest NPU	SDK	Open	Group	INT8 Perf.
Qualcomm	Hexagon NPU [15]	QNN [23]	✗	✗	73 TOPS
Google	Edge TPU [17]	Edge TPU API [7]	✗	✗	4 TOPS
MediaTek	MediaTek APU 790 [11]	NeuroPilot [13]	✗	N/A	60 TOPS
Huawei	Ascend NPU [6]	HiAI [9]	✗	✗	16 TOPS

"Open": Open-source?; "Group": Support per-group quantization MatMul? "N/A": No available documents for public; "INT8 Perf.": Int8 performance.



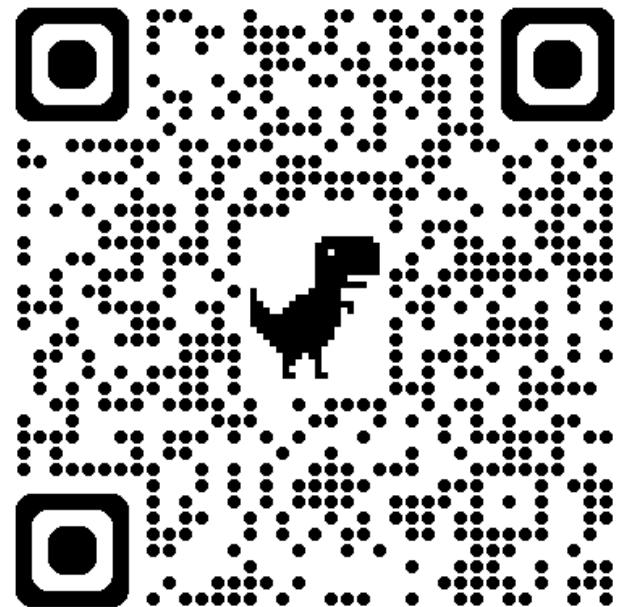
- *The gap between CPU/GPU and NPU increases over time*
  - *Moore's law still stands for NPU*
- *The gap of energy efficiency is even larger*

[1] Jinliang Yuan. “Mobile Foundation Model as Firmware”. In MobiCom’24.

# *Filling the design gap between legacy NPUs and modern LLM inference*

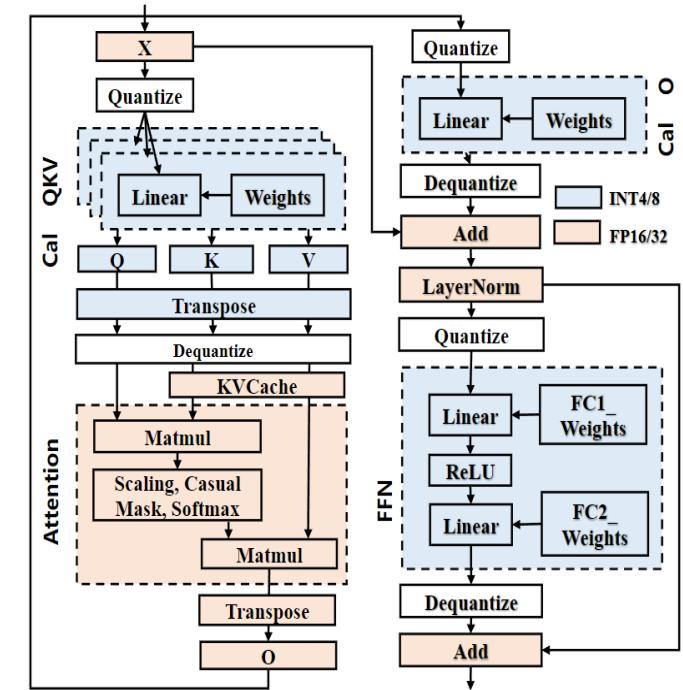
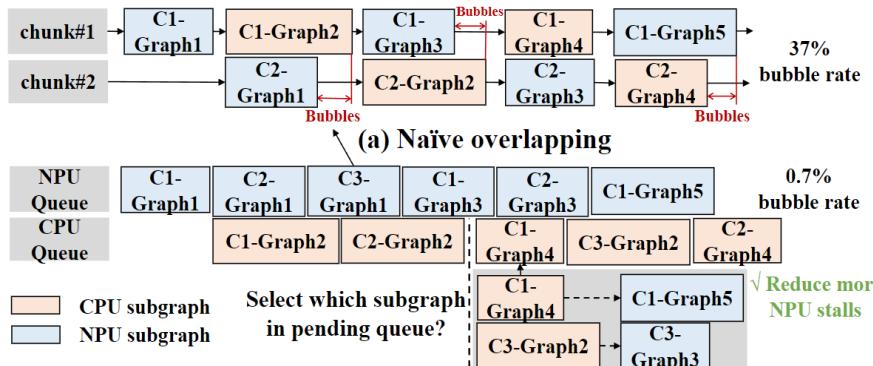
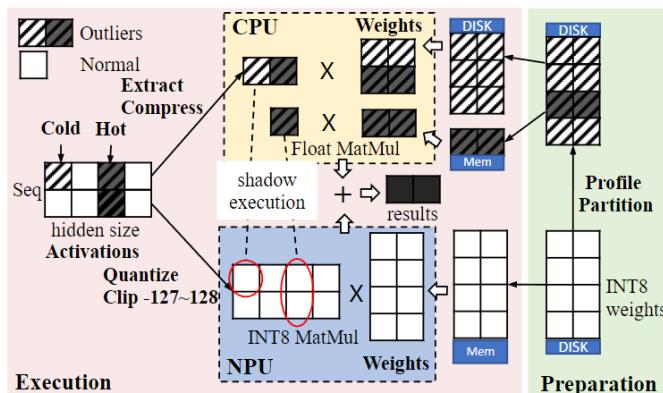
[ASPLOS'25] Fast On-device LLM Inference with NPUs

Code at <https://github.com/UbiqitousLearning/mllm>



# llm.npu: accelerating LLM prefilling with NPU

- Legacy mobile NPU has poor support for
  - (1) Dynamic shape; (2) FP operations; (3) group-level quantization
- llm.npu proposes
  - Chunked prefill with partial sharing
  - Shadow outlier execution across CPU/NPU
  - Out-of-order scheduling among CPU/NPU

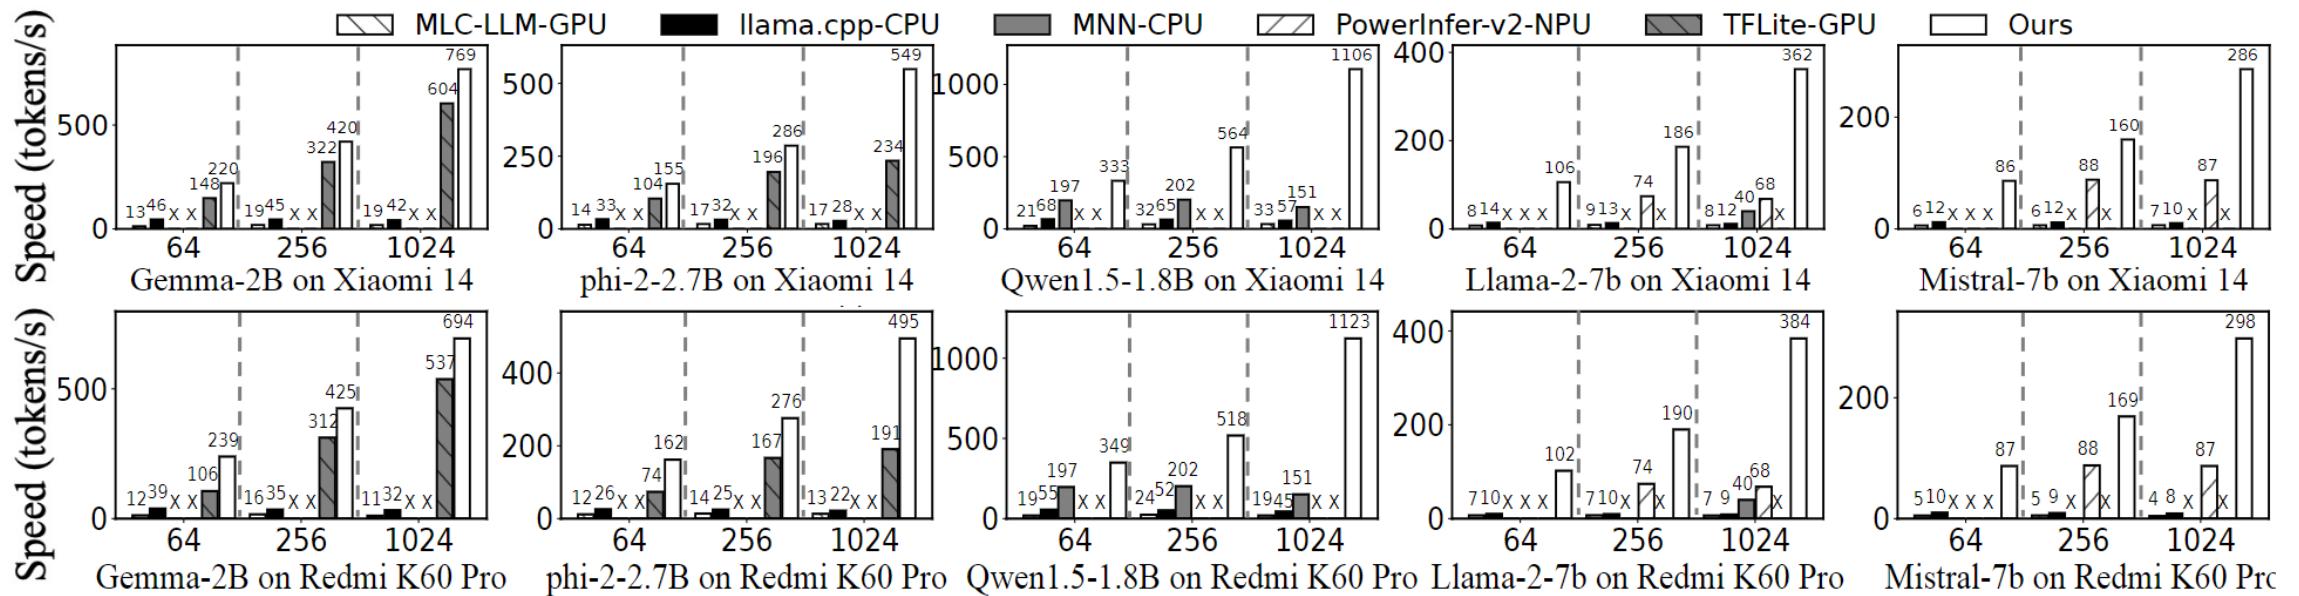


[1] Daliang Xu, et al. "Fast On-device LLM Inference with NPUs". In ASILO'25.

Mengwei Xu @ BUPT

# Highlighted results

**Prefill speed under different prompt lengths on different devices (datasets: Longbench-2wiki-Multi-doc QA)**  
Baselines: MLC-LLM (GPU), llama.cpp (CPU), MNN (CPU), PowerInfer-v2 (NPU), TFLite (GPU)



**7.3×–18.4× faster than baselines on CPU, and 1.3×–43.6× on GPU with prompt length of 1024**  
**Achieves >1000 tokens/second on Qwen1.5-1.8B (for the first time)**

# *Filling the design gap between legacy NPUs and modern LLM training*

[USENIX ATC'24] FwdLLM: Efficient Federated Finetuning of Large  
Language Models with Perturbed Inferences

Code at <https://github.com/UbiquitousLearning/FwdLLM>



# FwdLLM: BP-free LLM finetuning

- Key idea: leveraging forward gradient for LLM finetuning

$$g(\theta) = (\nabla f(\theta) \cdot \mathbf{v}) \mathbf{v}$$

A random, independent perturbation with same size as trainable weights  $\theta$

Forward gradient, an unbiased estimator of  $f(\theta)$ 's gradient

The directional derivative of  $f$  at point  $\theta$  in direction  $v$ . Computing it takes only forward, no need for backpropagation

Compared to BP approach:

- Legacy NPU-compatible
- More memory efficient

- Further optimizations:
  - var.-controlled perturbation pacing
  - discriminative perturbation sampling
- Highlighted results: federated Llama-7B finetuning on devices, with significant speedup and memory saving

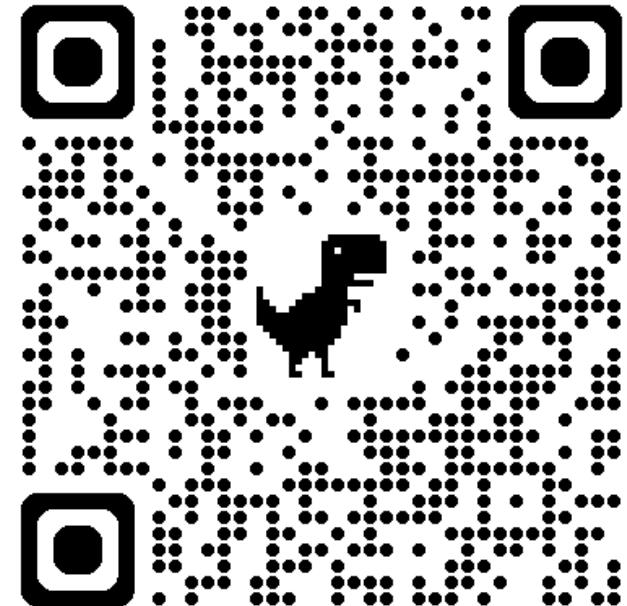
Methods	Mem. (GB)	Centralized Training (A100)			Federated Learning		
		Acc.	Round	Time	Acc.	Round	Time
BP, FP16	39.2	89.7	500	0.1 hrs			
BP, INT8	32.4	88.6	500	0.06 hrs			
BP, INT4	28.5	87.8	500	0.04 hrs			
Ours, FP16	15.6	87.0	240	1.5 hrs			
Ours, INT8	7.9	86.9	260	0.8 hrs			
Ours (CPU), INT4	4.0	85.8	130	0.25 hrs	85.8	130	0.19 hrs
Ours (NPU*), INT4							0.07 hrs

[1] Mengwei Xu, et al. "FwdLLM: Efficient Federated Finetuning of Large Language Models with Perturbed Inferences". In ATC'24.

# *Filling the design gap between legacy NPUs and modern LLM design*

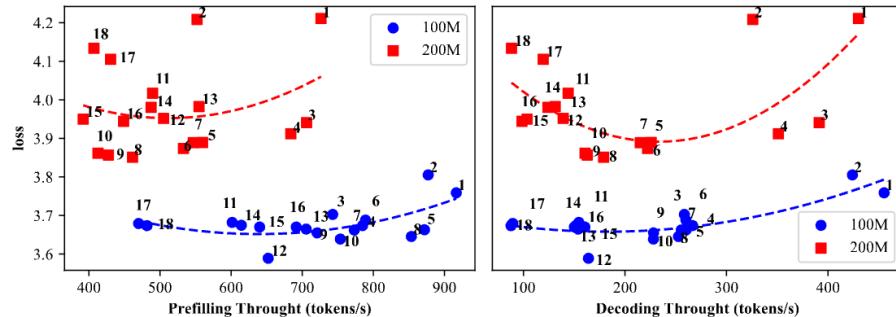
[arxiv'24] PhoneLM: an Efficient and Capable Small Language Model Family  
through Principled Pre-training

Code at <https://github.com/UbiqitousLearning/PhoneLM>  
Models at <https://huggingface.co/mllmTeam/PhoneLM-1.5B>



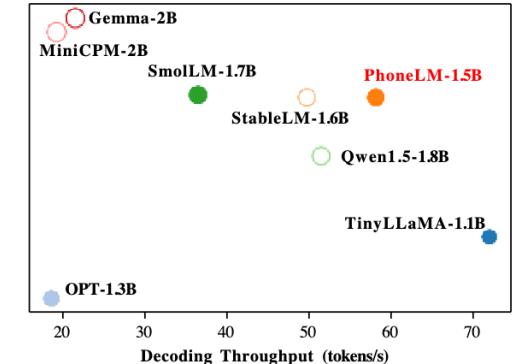
# PhoneLM: efficient SLMs for devices

- An argument: SLM shall adapt to the target device hardware
  - Hardware-specific, ahead-of-pretraining hyperparameter search for runtime resource efficiency



The SLM design (hyper-parameters) has more impacts on the runtime performance than the capability

Name	Size	Date	Training tokens	HellaSwag	WinoGrande	PIQA	SciQ	BoolQ	ARC Easy	ARC Challenge	Average
Pythia (EleutherAI, 2023.03a)	1.4B	23.03	207B	52.0	57.2	71.1	79.2	63.2	53.9	28.3	57.84
OPT (Facebook, 2022.05a)	1.3B	22.05	180B	53.7	59.0	71.0	78.1	57.2	51.3	28.0	56.90
BLOOM (BigScience, 2022.11b)	1.1B	22.11	350B	43.0	54.9	67.2	74.6	59.1	45.4	25.6	52.83
TinyLlama (Unknown, 2023.12)	1.1B	23.12	3B	59.1	58.9	73.0	82.3	58.6	55.7	31.0	59.80
MobileLLaMA (Meituan, 2023.12)	1.4B	23.12	1.3T	56.1	59.4	73.0	81.9	56.7	55.8	30.3	59.03
MobiLlama (MBZUAI, 2024.02)	1B	24.02	1.25T	62.2	59.3	74.8	82.8	60.3	56.4	31.7	61.07
OpenELM (Apple, 2024.04)	1.1B	24.04	1.5T	64.8	61.7	75.6	83.6	63.6	55.4	32.3	62.43
DCLM (Toyota, 2024.08)	1.4B	24.08	4.3T	53.6	66.3	77.0	94.0	71.4	74.8	41.2	68.33
SmolLM (HuggingFace, 2024.07)	1.7B	24.07	1T	49.6	60.9	75.8	93.2	66.0	76.4	43.5	66.49
Qwen 1.5 (Alibaba, 2024.02)	1.8B	24.02	2.4T	60.9	60.5	74.2	89.4	66.5	59.1	34.7	63.61
Galactica (Facebook, 2022.11)	1.3B	22.11	106B	41.0	54.4	63.8	87.7	62.0	58.6	30.5	56.86
StableLM 2 (StabilityAI, 2024.01)	1.6B	24.01	2T	68.8	64.1	75.1	76.9	80.0	60.3	39.2	66.34
Cerebras-GPT (Cerebras, 2023.03a)	1.3B	23.03	371B	38.4	51.9	66.8	73.0	59.3	45.8	25.3	51.50
MiniCPM (OpenBMB, 2024.04)	1B	24.04	1.2T	67.5	63.7	75.1	91.0	70.5	62.9	38.1	66.97
MiniCPM (OpenBMB, 2024.04)	2B	24.04	1.2T	67.2	63.9	76.1	92.5	74.6	69.0	42.7	69.43
Gemma (Google, 2024.02)	2B	24.02	3T	71.4	65.2	78.4	91.4	69.9	72.3	42.0	70.09
Gemma 2 (Google, 2024.07)	2B	24.07	2T	55.0	68.7	78.7	96.0	73.6	80.3	46.9	71.31
<b>PhoneLM</b>	<b>1.5B</b>	<b>24.11</b>	<b>1.5T</b>	<b>66.9</b>	<b>63.0</b>	<b>77.3</b>	<b>88.8</b>	<b>65.5</b>	<b>69.7</b>	<b>39.9</b>	<b>67.31</b>



SOTA tradeoff between capability and efficiency

[1] Rongjie Yi, et al. “PhoneLM: an Efficient and Capable Small Language Model Family through Principled Pre-training”. In preprint’24.

# *Looking into the future..*

,

## Mortal Computation

- If we abandon immortality and accept that the knowledge is inextricable from the precise physical details of a specific piece of hardware, we get two big benefits:
- Huge energy savings
  - We can use very low power analog computation.
- Much cheaper hardware
  - The hardware could be grown cheaply in 3-D instead of being manufactured very precisely in 2-D.
  - This would require lots of new nano-technology or perhaps genetic re-engineering of biological neurons.

We shall probably look for hardware-software co-evolution, e.g., mortal computation by Geoffrey Hinton

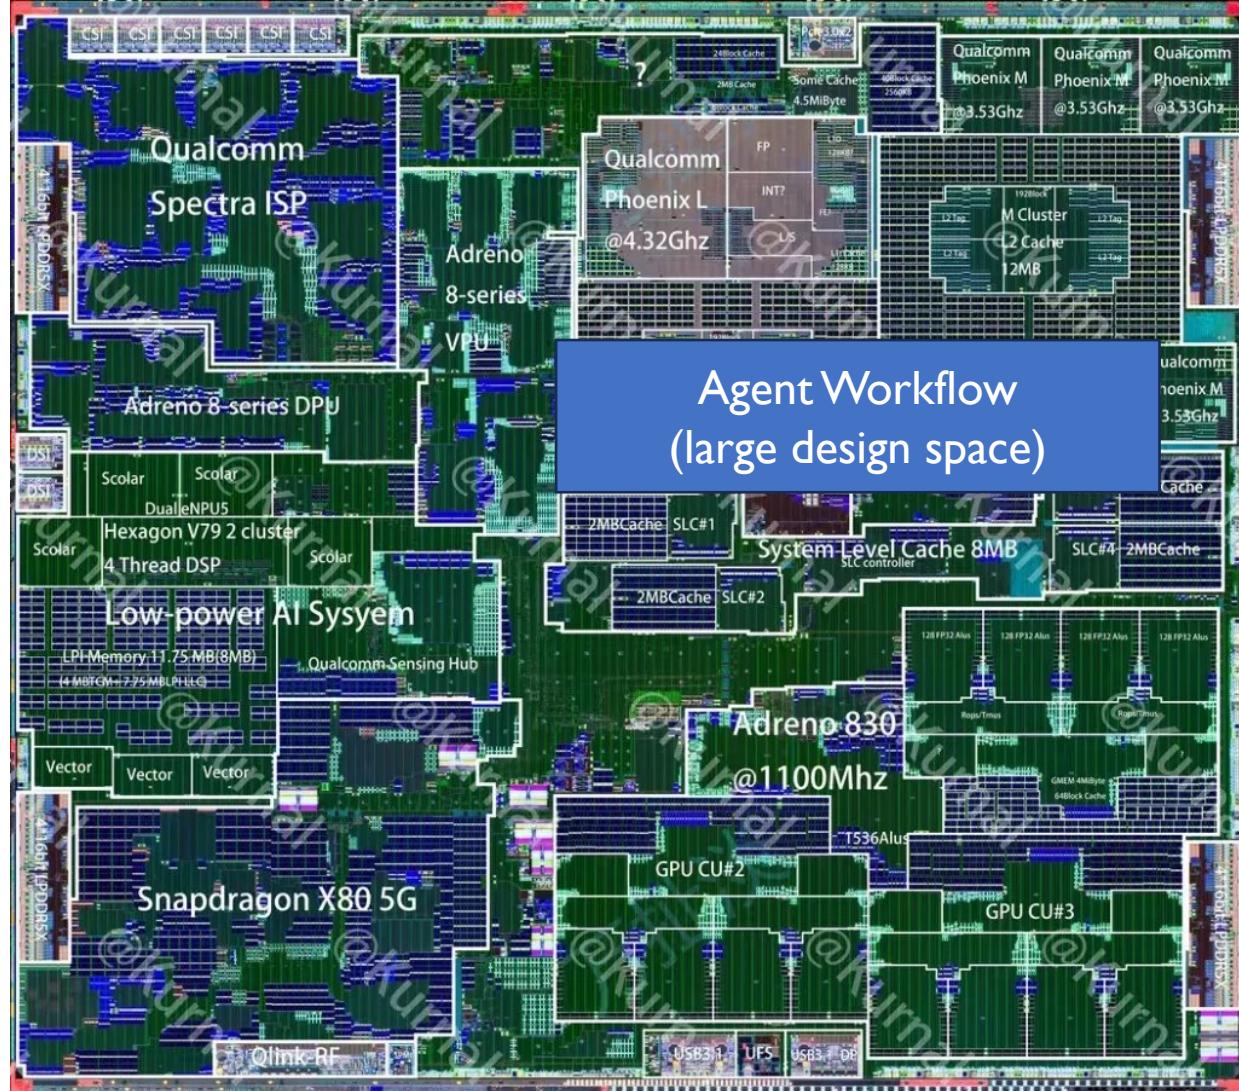
“Two paths to Intelligence” by Geoffrey Hinton

# The Future: full-stack design!



<https://innogyan.in/2024/10/28/die-shot-of-snapdragon-8-elite-reveals-component-space-allocation/>

# The Future: full-stack design!



Agent Workflow  
(large design space)

Prompt/LoRa  
Response

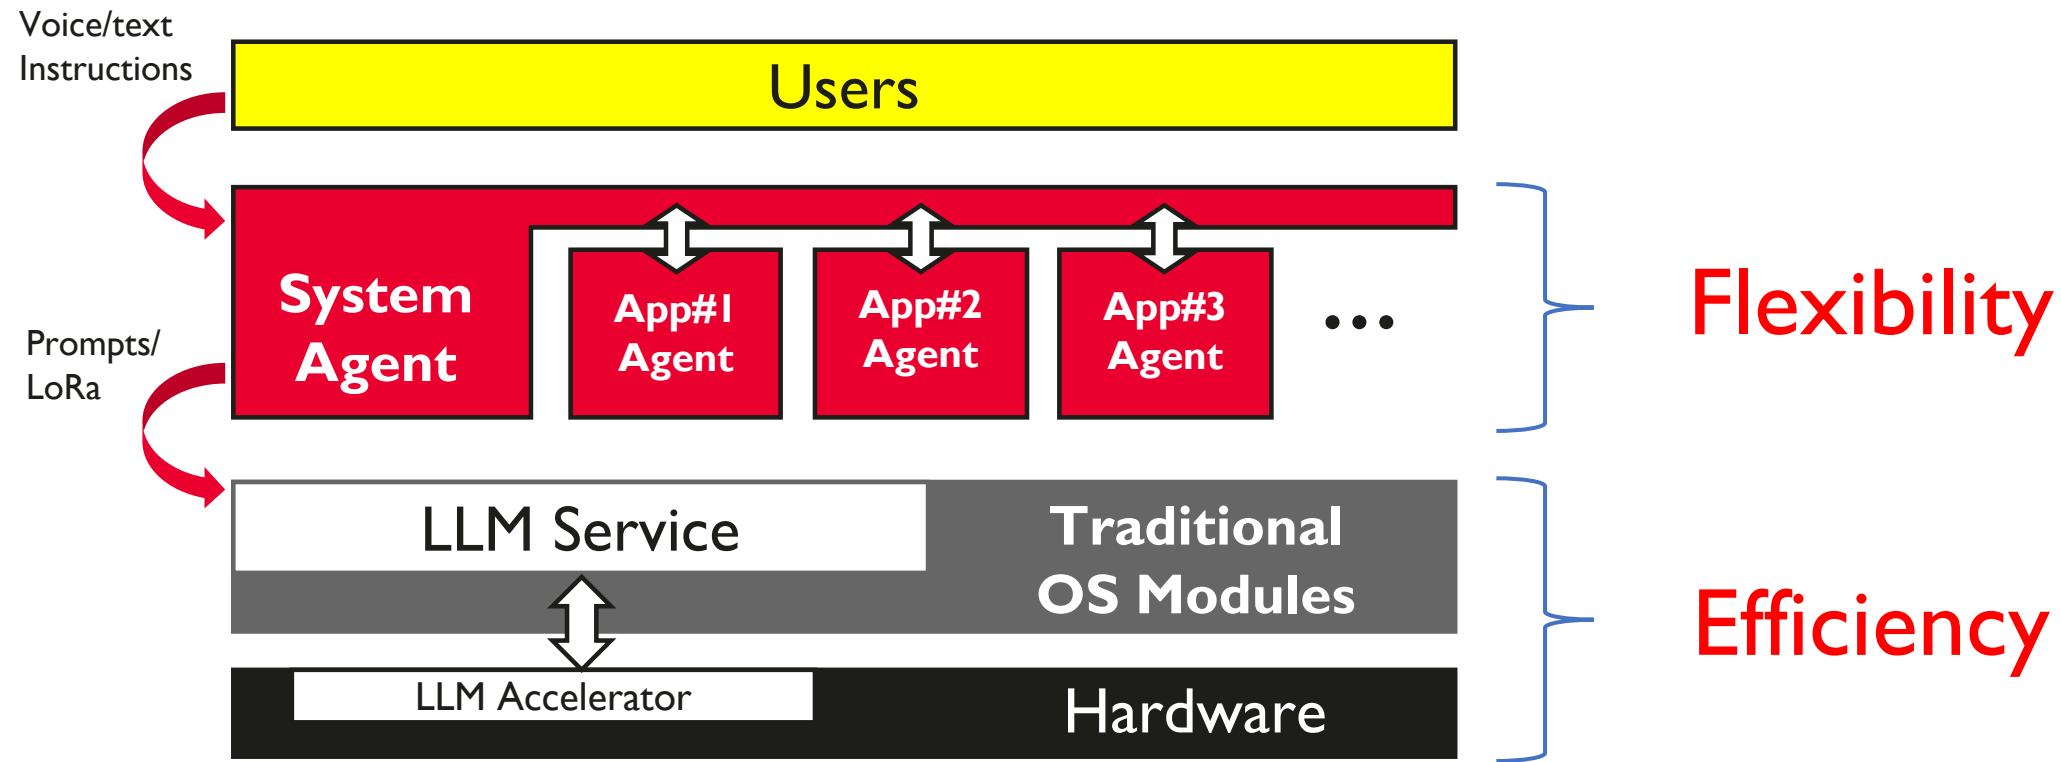


Time to sacrifice flexibility for efficiency!  
(we still have flexibility at agent workflow)

<https://innogyan.in/2024/10/28/die-shot-of-snapdragon-8-elite-reveals-component-space-allocation/>

# The Future: full-stack design!

- One LLM, Many Agents



# Takeaways

- On-device LLM is reinventing the mobile devices
  - A total paradigm shift of mobile AI ecosystem
- It calls for full-stack LLM research
  - OS, runtime, model, and application (agent)