



Proceedings of the 33rd ACM Web Conference 2024

# Towards Energy-efficient Federated Learning via INT8-based Training on Mobile DSPs

*Jinliang Yuan, Shangguang Wang, Hongyu Li, Daliang Xu, yuanchun Li, Mengwei Xu, Xuanzhe Liu*

Peking University



Beijing University of Posts and  
Telecommunications




Tsinghua University



AI-driven

Web Applications



History suggestions




Page recommend



Input prediction




Vicious prediction



Private user data



# Federated Learning Algorithm



- ① Keep private data locally
- ② Train the model
- ③ Upload  $\Delta W$
- ④ Download global model




Strict protect



High accuracy




# Is Federated Learning Algorithm enough to protect data?

No! The **energy** is the main fence.




Such huge **energy** is unacceptable.


$$\mathbf{E}(\text{Training VGG16 on CIFAR-10}) = \mathbf{E}(\text{Watch TikTok/YouTube for 12/24 hours}) = \mathbf{E}(\text{Play video game for 9 hours})$$



# Why energy consumption is so huge?




*DNN models are trained on **GPUs***




- floating computation
- matrix manipulation

*For phones, they need executing on **CPUs***



- integer computation
- general computing



We observe there is **DSP** on the phone

It's energy friendly and can handle matrix manipulation

Models	CPU, FP32		CPU, INT8		DSP, INT8	
	T	E	T	E	T	E
MobileNet-V1	11.4	88.2	4.3	22.7	2.5	5.4
MobileNet-V2	8.6	64.8	4.6	22.9	3.1	5.4
ResNet-50	78.7	597.6	27.8	131.4	9.2	28.8
Inception-V4	266.3	1,980	81.5	399.6	17.2	59.4
EfficientNet-V2	33.0	187.2	13.4	59.4	8.44	12.6

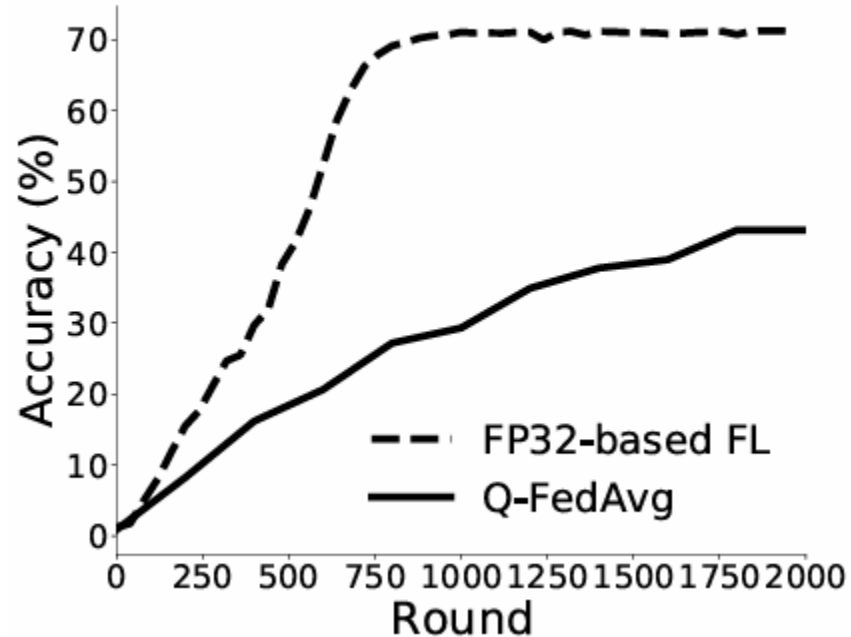
	FP32 Acc. (%)	Octo Acc. (%)	Acc. Degradation (%)
GoogLeNet, FM	99.1 – 99.5	97.9 – 98.6	0.9 – 1.2
GoogLeNet, CF	97.8 – 99.2	97.6 – 98.8	0.2 – 0.4
AlexNet, FM	95.6 – 98.4	92.8 – 94.3	2.8 – 4.1
AlexNet, CF	91.8 – 95.2	86.1 – 87.3	5.7 – 7.9
VGG11, FM	97.5 – 98.8	94.4 – 96.5	2.3 – 3.1
VGG11, CF	97.2 – 99.5	96.5 – 98.6	0.7 – 0.9

Save energy 32X, speeds up 9X

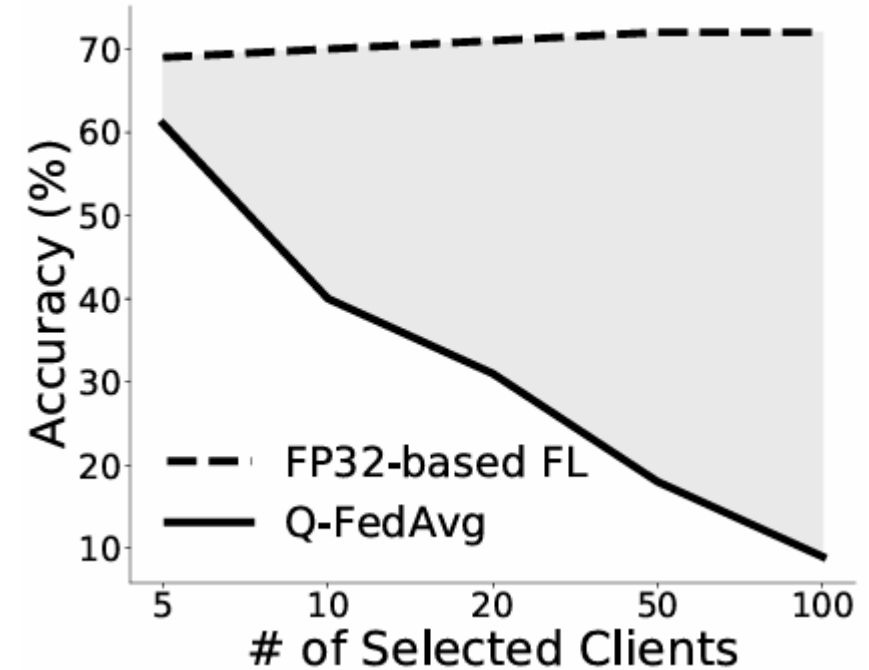
Accuracy losses only 2%

# Challenge: Directly using DSP with FL is impractical

*Low and slow convergence*

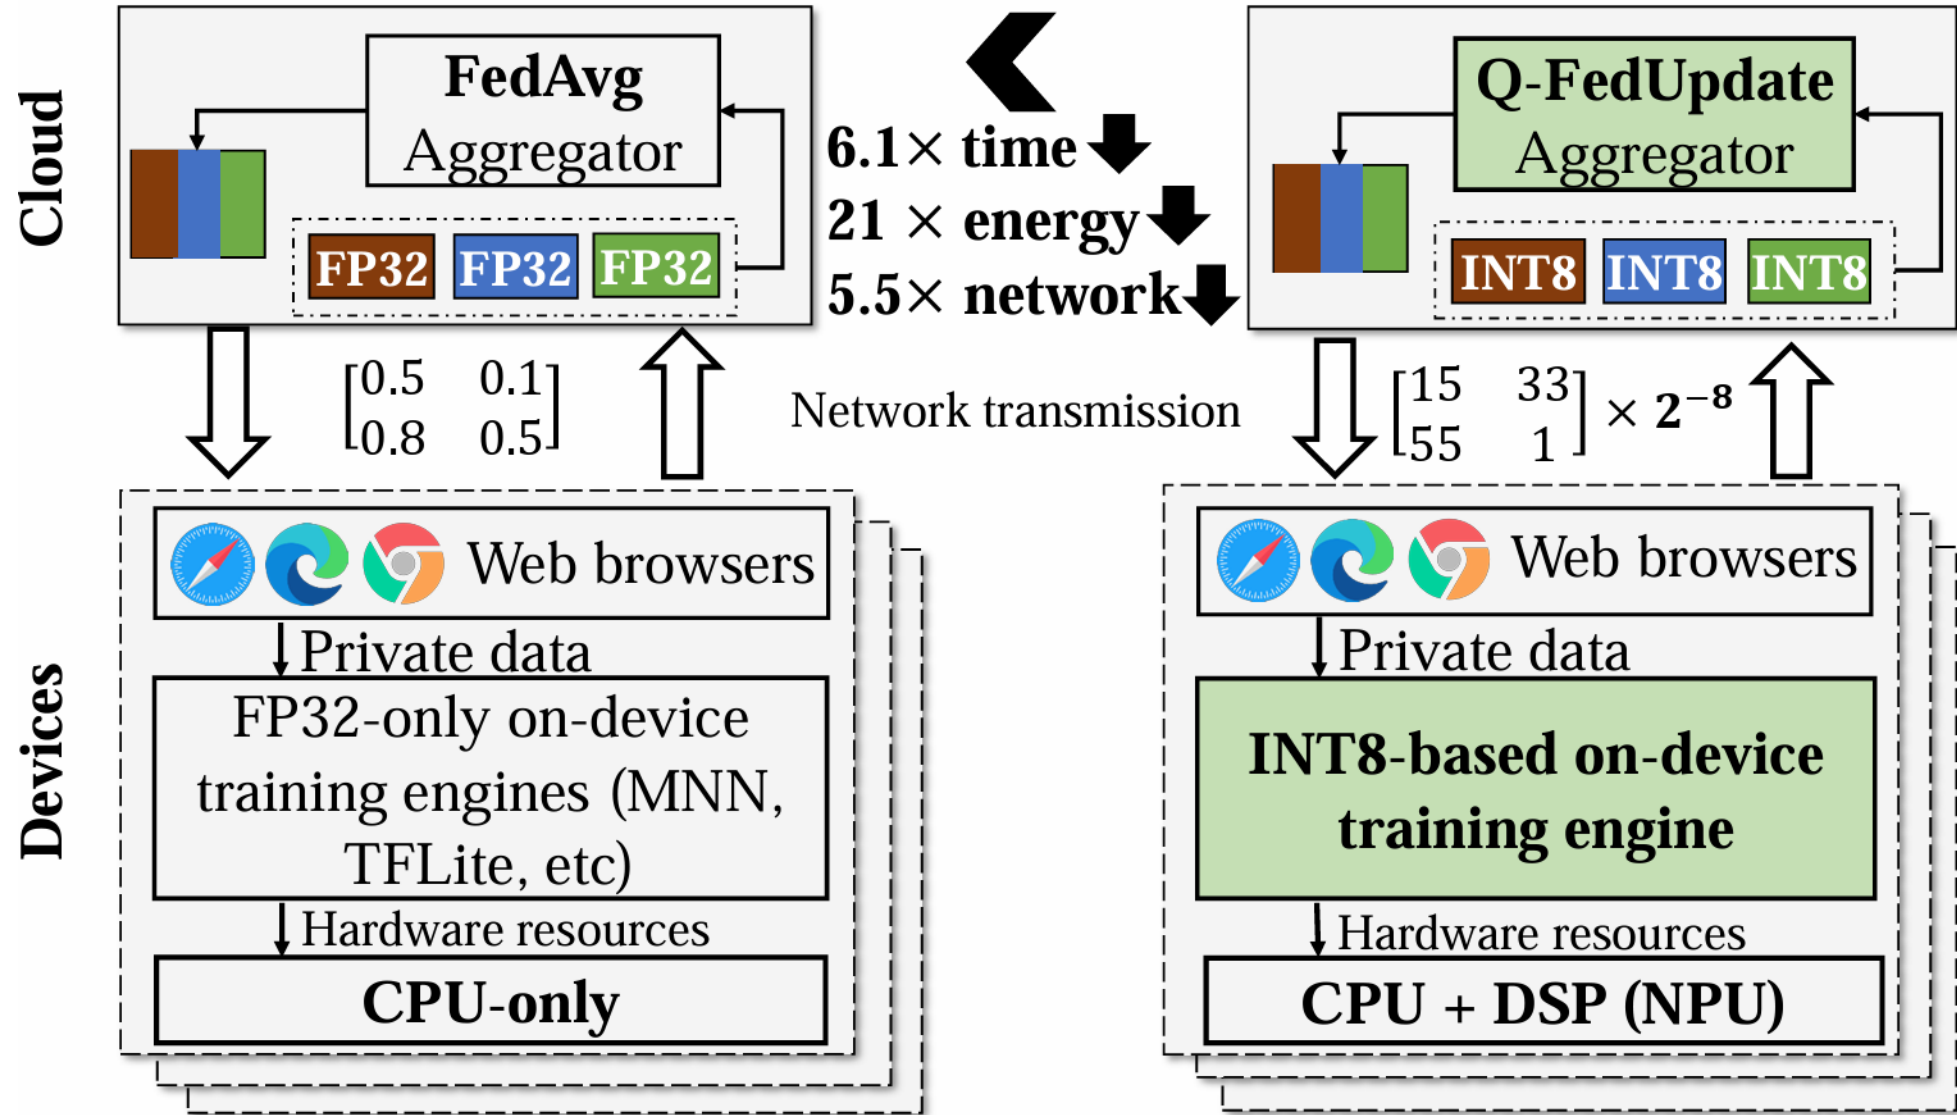


*Low scalability*



**GOAL: How to design an energy efficient and high accuracy algorithm with DSP?**

# Key Idea: FP32-INT8



(a) Traditional FP32-based FL

(b) Proposed INT8-based FL


# Design1:

## Error-Compensated Aggregation


$$\begin{aligned}
 w(t+1) &= w(t) - \sum_{k=1}^K \frac{n_k}{n} \Delta w_d^k(t+1) \\
 &= w(t) - \sum_{k=1}^K \frac{n_k}{n} (w_d(t) - w_d^k(t+1)) \\
 &= \underbrace{w(t) - w_d(t)}_{\text{Quantization Error}} + \underbrace{\sum_{k=1}^K \frac{n_k}{n} w_d^k(t+1)}_{\text{Q-FedAvg}}.
 \end{aligned}$$

# Design2:

## Pipelined Batch Quantization



(a) Traditional four procedures of local training on device



(b) Pipeline procedures: parallelization of  $P_2$  on CPU and  $P_4$  on DSP




# Evaluation Results


Algorithms	FEMNIST			CIFAR-10			CIFAR-100		
	Acc (%)	T (hours)	E (kJ)	Acc (%)	T (hours)	E (kJ)	Acc (%)	T (hours)	E (kJ)
FloatFL (#1)	<b>82 (1↓)</b>	0.57	0.56	<b>84 (3↓)</b>	8.94	42.5	<b>71 (2↓)</b>	7.58	2.2
QuanFL (#1)	80	0.40	0.62	82	6.17	53.1	67	3.09	2.6
Q-FedUpdate (#2)	81	0.14	0.18	81	2.80	18.9	69	1.71	0.9
Q-FedUpdate (L)	81	<b>0.12 (4.8×</b>	<b>0.04 (14×</b>	81	<b>2.40 (3.7×</b>	<b>4.6 (9×</b>	69	<b>1.61 (4.7×</b>	<b>0.21 (10×</b>
Q-FedUpdate (M)	81	<b>0.08 (7.1×</b>	<b>0.02 (28×</b>	81	<b>1.77 (5.1×</b>	<b>3.0 (14×</b>	69	<b>1.47 (5.2×</b>	<b>0.14 (16×</b>
Q-FedUpdate (H)	81	<b>0.07 (8.1×</b>	<b>0.02 (28×</b>	81	<b>1.62 (5.5×</b>	<b>2.7 (16×</b>	69	<b>1.44 (5.3×</b>	<b>0.13 (17×</b>

Time speeds up 6x

Energy saves 21X



Accuracy losses only 2%



# Thanks

Q&A